

TEI: Target–Evaluate–Improve — A Cost-Efficient Self-Improving Loop for Agentic Systems via Sequential Decision Gating (TEI-SWE)

Orkhan Javadli¹ and Anni Zimina²

¹Massachusetts Institute of Technology ojavadli@mit.edu

²Stanford University zimina@stanford.edu

August 9, 2026

Artifacts: the **TEI-SWE-30** frozen agent set, 1,140 why-records, blinded-vote and audit datasets, patches, and all scripts (release pointer in §14);

methodology: <https://github.com/ojavadli/tei-bench> and the `tei-loop` gate.

Abstract

TEI (Target–Evaluate–Improve) is a self-improving loop for agentic systems whose sequential decision gating delivers population-scale improvement at \$0.47 per agent across all evaluation passes at our conservative accounting rate—about \$0.05 at current list prices—the two ledgers Table 4 reconciles. Applied to 30 real systems frozen from the official SWE-bench leaderboard archive [1, 2], in up to 30 structural-fix plus 30 prompt-optimization gated iterations per agent (budget scaling fixed a priori: 9/8/13 agents at 60/36/24 versions), the loop produced 1,140 audited candidate versions and 547 syntax-clean committed patches, raising rubric scores for all 30. 26 systems received shipped code changes, and blinded re-evaluation—real before/after code, randomized order, no narrative—covers those 26. Pre-repair confirmatory result: 24/26 patched agents (118/130 votes); after defect repair, the adaptive retest shows 26/26 strict majorities, 24 unanimous (128/130 votes). A *pre-registered sham-patch placebo* (tagged before execution) draws only 26.9% of votes (7/26 majorities), rejecting the generic changed-code / style explanation for the primary judge’s preference; TEI’s diagnosis-guided proposals beat budget-matched unguided generation on 10/10 agents; and an execution micro-arm confirms do-no-harm at the VERIFIED rung (SWE-agent, 6 paired instances, 0 losses). Gains clear measured paraphrase noise floors on 30/30 agents, and the loop’s rare failure mode—6 syntax-breaking patches its own validation ladder caught—is permanently closed by a zero-cost syntax pre-gate (`ast.parse`) contributed to the deployed gate. We release TEI-SWE-30, the 1,140-record why-record and blinded-vote datasets, and a one-command audit tool (`python tei_audit.py --repo R --base SHA --cand HEAD`). Rubric magnitudes below the measured-noise MDE, and outcome claims, remain scoped to their rungs (§8, §10).

1 Introduction

This paper reports, to our knowledge, the broadest and cheapest controlled application of an automated improvement loop to real software-engineering agents to date: 30 systems spanning 7.0%–79.2% official resolve rates, frozen at recorded SHAs from the SWE-bench leaderboard archive [1–3], each run through up to 30 structural-fix and 30 prompt-optimization iterations of the gated TEI v7

loop (budget scaling fixed a priori: 9/8/13 agents at 60/36/24 versions) for roughly the price of a sandwich per agent. The loop produced 1,140 scored candidate versions carrying auditable why-records, applied 547 of them as real commits, and raised every agent’s four-dimension judged quality, $0.607 \rightarrow 0.659 \rightarrow 0.669$ on average (Fig. 2).

The result that makes those numbers citable, rather than merely encouraging, is the **blinded confirmation**. Judge-scored optimization is rightly distrusted: a judge that knows which version is “improved,” and is shown the optimizer’s own rationale, inherits every bias of the pipeline it is scoring [4, 5]. We therefore re-evaluated the real artifacts—the applied diffs—under blinding: the judge sees the actual before/after contents of the changed regions, in randomized order, with no why-record, no scores, and no statement of which state is newer, five independent times per agent. Pre-repair confirmatory result: 24/26 patched agents (118/130 votes); after defect repair, the adaptive retest shows 26/26 strict majorities, 24 unanimous (128/130 votes). A second model of the same GPT-5.6 family concurs (blinded, 8/10 subsampled agents). Whatever the loop is producing, it is code that bias-controlled evaluation consistently prefers—at \$0.35 per agent for the rubric-rung optimization run, plus \$0.03 per agent for the blinded confirmation itself.

Working at this scale without execution access (the archive’s traces are credentialed; SWE-bench’s ground-truth harness needs infrastructure the study host lacked; 4 linked repositories ship no runnable source at all) forced a discipline that we found to be the paper’s second contribution: the **validation ladder**. Every claim is stated at the rung of evidence that supports it—narrative rubric, blinded A/B, deterministic static checks, execution—each rung cheaper than the next is trusted (Fig. 1). Climbing the ladder did exactly what such structure is for: the blinded rung confirmed the improvement signal was real rather than narrative, and the static rung caught a rare failure mode (exact-match patches splicing strings or landing at the wrong indentation in 6 files across 5 agents) that was repaired, re-confirmed under blinding, and permanently closed by a zero-cost `ast.parse` pre-gate this study contributes back to the methodology’s decision layer. Post-repair, all 30 branches are compile-clean and the blinded margin is unanimous-or-near-unanimous on every patched agent.

Research questions. **RQ1 (improvement).** Does a gated two-phase loop improve real, heterogeneous, third-party agents—and does the improvement survive bias-controlled evaluation? **RQ2 (cost/coverage).** At what cost, and at what coverage, relative to 2026 prompt/agent optimizers? **RQ3 (statistics).** Do the gains clear measured noise—paraphrase floors and test-retest minimum detectable effects—rather than assumed constants? **RQ4 (method).** What does the validation ladder add over single-rung judge scoring? **RQ5 (provenance).** Can the whole study be re-derived from public artifacts by a hostile reviewer?

Contributions.

1. **The demonstrated capability** — the broadest, cheapest gated improvement application on record: 30 real leaderboard systems, 1,140 audited candidate versions, 547 syntax-clean committed patches, blinded-confirmed preference on 24/26 patched agents (post-repair adaptive retest: 26/26), certified against a pre-registered sham placebo (26.9% vs 98.5%) and budget-matched random proposals, reproduced by an external judge, with executed do-no-harm on the one harness-runnable system—at \$0.47 per agent all-in (accounting rate; \$0.05 at verified list prices).
2. **The validation ladder and released `tei-audit` tooling** — rubric < blinded A/B < syntax pre-gate < execution, every claim at its rung, with the two cheapest rungs packaged

The validation ladder — each rung is cheaper than the rung above it is trusted; claim at the rung you reached. This study’s results per rung:			
Rung	What it checks	Marginal cost here	This study
4. Execution	real outcomes (resolve rate)	harness + rollouts	1 system executed: do-no-harm confirmed, 1/6 = 1/6 resolved
3. Syntax pre-gate	<code>ast.parse</code> of patched files	\$0	0 errors / 41 files (post-repair)
2. Blinded A/B	preference w/o narrative or direction	\$0.03/agent	24/26 pre-repair (118/130); 26/26 after repair (128/130)
1. Rubric (anchored)	judged dimensions vs. diagnosis	\$0.35/agent	30/30 improved, $\Delta +0.062$

Checklist for any optimization claim: hold out the selection signal; beat budget-matched random search; run free static checks *before* any judge call; blind the judge to direction and narrative; report the noise floor and the measured-noise MDE next to every delta; label every number’s rung.

Figure 1: The validation ladder and the practitioner checklist it implies.

as one command (`python tei_audit.py --repo R --base SHA --cand HEAD`) that audits *any* optimizer’s diff (§4, Fig. 1).

3. **The contributed syntax pre-gate** (`tei_loop.gate.static_pregate`): parse-breaking candidates are rejected before any judge call, at zero cost—the permanent fix for a failure mode the ladder caught in this study (§7).
4. **TEI-SWE-30** — a frozen, reviewer-reproducible agent set from the official archive (218 submissions $\rightarrow$ 101 systems $\rightarrow$ 30 verified-sourceable), with per-agent diagnoses and full provenance (§3).
5. **Measured-noise statistics and open artifacts**: paraphrase floors under candidate-identical calibration, test-retest MDEs, agent-level exact CIs; the why-record and blinded-vote datasets; one-command regeneration of every number (§8, §14).

2 Related Work

SWE-bench and its agent ecosystem. SWE-bench [1] scores systems on real GitHub issues against repository test suites; Verified [3] is its human-validated split, and the submission archive [2] records per-instance outcomes per system—our population frame and probe source. The frozen set spans the design space: ACI agents (SWE-agent [6]), pipeline systems (Agentless [7]), search-based repair (AutoCodeRover [8]), RAG baselines [1], IDE assistants (aider [9]), and tree search (moatless-tools [10]). To our knowledge no prior work applies one optimization procedure uniformly across a leaderboard’s worth of third-party systems.

Prompt and agent optimizers (2026 state). APE [11], OPRO [12], DSPy/MIPROv2 [13, 14], TextGrad [15], GEPA [16] (ICLR 2026 oral), Maestro [17], and ACE [18] optimize prompts, programs, or agent graphs against a scoreable objective, reporting e.g. “over 10%” above MIPROv2 across six tasks (GEPA) or +4.9% above GEPA on two benchmarks (Maestro). All evaluate on their own scaffolds with executed rollouts as both selection signal and evidence; none reports a bias-controlled check of its judge or selection signal. §9 makes the comparison explicit, substrate labels included. The earlier TEI-Bench program supplies our design rationale: naive scorers manufacture gains (a +0.175 artifact), beating a baseline means little without beating budget-matched random search, and a do-no-harm confirmation gate converts an unreliable loop into a zero-regression one—the executed lineage (v3 +0.025, $p=.024$, 0 regressions; v6 +0.026, $p=.029$) behind the v7 decision layer used here.

Table 1: Selection funnel. Operational note preserved for reproducers: macOS ships no `timeout(1)`, so a first `git ls-remote` sweep returned exit 127 for all URLs—“every repository unreachable”—until a control probe exposed it; verification was re-run with in-process timeouts.

Stage	Count
Submissions parsed (lite 84 + verified 134)	218
Unique systems after name-normalized dedupe	101
Kept submission names a <code>github.com</code> repository	30
Repository verified reachable (<code>git ls-remote</code>)	30
widening: <code>evaluation/multilingual</code> entries examined	13
... of which sourceable (repo URL in metadata)	0
Frozen set	30

Self-improving and auto-designed agents. The Darwin Gödel Machine [19] evolves *its own* single agent with executed benchmark validation (20.0%→50.0% on SWE-bench for one lineage); ADAS [20] meta-searches agent designs on tasks it executes; Trace [21] optimizes one workflow’s computational graph from execution traces. TEI-SWE occupies the complementary cell no prior system reports: *population-scale improvement of thirty third-party systems it does not control*, with certified selection (placebo-controlled blinded preference, pre-gated patches, gated do-no-harm) at two to three orders of magnitude lower cost per system—capability those self-referential loops do not claim, on evidence rungs they do not report.

Selection-consistent evaluation. Our gate-and-floor machinery joins a 2025–26 line of work on making selection-time signals trustworthy: correcting the winner’s curse in adaptive benchmarking [22], non-vacuous generalization bounds for prompt selection [23], distribution-free risk control for prompt deployment [24], conformal factuality certificates [25], and evaluation-subset design for prompt optimization [26]. The ladder adds the orthogonal axis these works assume away: *what the score is grounded in* (rubric/blinded/static/executed). LLM-judge bias [4, 5] and benchmark contamination [27] motivate the blinded rung; the four Evaluation dimensions adapt Agent GPA [28]; the gate arithmetic uses Wilson bounds [29], Hoeffding margins [30], Jeffreys priors [31], and Holm correction [32].

3 The TEI-SWE-30 Frozen Agent Set

The population must survive a hostile reviewer reconstructing it. We fixed the selection rule before parsing anything, used the official archive as the sole frame, and report every departure from target rather than repairing it.

Source and rule. SWE-bench/experiments [2] frozen at 2f15350c (accessed August 8, 2026): 218 submissions across Lite and Verified, resolve rates computed exactly as the archive’s own leaderboard script computes them. Rule, applied unchanged: parse every submission; deduplicate to unique systems (normalized names, keep each system’s best submission); keep systems whose *kept* submission names a `github.com` repository in its own metadata (URLs never guessed); verify each with `git ls-remote`; rank by resolve rate. Result: 101 unique systems, 30 sourceable (Table 1), each shallow-cloned at a recorded SHA onto a local `tei-v7` branch (full list with links: Appendix Table 10).

Table 2: Runnability taxonomy over the 30 repositories (content screen; precedence-ordered classes). Reconciling the three runnability notions used in this paper: this content screen leaves 16 systems with no blocker keyword; the stricter entry-point screen (`tier_a_assessment.txt`) leaves 3 (SWE-agent, KGCompass, aider); exactly 1 (SWE-agent) was executed end-to-end in §10.

Blocker	Systems
statically runnable (blocked only by absent Docker ground truth)	16
requires Docker orchestration	6
no source code in the linked repository	4
requires GPU / local weights	4
Executable end-to-end on this host (Docker ground truth available)	0

Why 30 and not 31. The target was 31; the strict rule yields 30. The pre-approved widening (evaluation/multilingual/) added zero: none of its 13 entries names a repository, and all 13 are one system under the dedupe rule. One alternative reading of the source rule reaches 31 (qualify a system if *any* of its submissions names a verified repository, admitting autocoderover [8] and patchpilot), but it was identified only after the strict count was known, so adopting it would have tuned selection to the target; the set stays 30 and unpadded. Two entries share one repository (codeshellagent/codeshelltester) and are kept as distinct leaderboard systems.

3.1 Population defects the rule admitted

Honest selection rules admit defects, and this one admitted two worth naming (marked in Appendix Table 10). Rank 30’s kept submission names SWE-bench/SWE-bench—the benchmark repository, not an agent scaffold—so its “source” is the RAG baseline inside the benchmark tooling; rank 27’s names ozyshr/RepoGraph, the RepoGraph scaffold rather than an Agentless fork. Two entries share one repository (WisdomShell/codeshell), so the 30 systems map to 29 distinct repositories. None of this was repaired post hoc: the rule is what a reviewer can re-run, defects included.

The execution boundary (stated once). Three facts of the public record shaped the run, one of which our revision corrected: the archive ships aggregate results only (its logs//trajs/ live behind credentialed storage; anonymous listing returns 403); deciding “resolved” requires SWE-bench’s container-based harness; and 4 linked repositories contain no runnable source (README/report only). A static screen found 3 systems runnable-in-principle (SWE-agent, KGCompass, aider); Table 2 gives the taxonomy, §10 the execution plan and its result. Four repositories *do* ship their own recorded trajectories in-repo (sweagent 22, agentless 15, darsagent 1, coder 522), which we use as a middle grounding rung (§6).

4 Method: the Gated Loop and the Ladder

The optimization methodology is TEI v7 as deployed and frozen (tei-bench@626b455f; gate ported in tei-loop@e0293135); this study executes it and contributes one upgrade back (the static pre-gate). Per agent, on its branch: **(1) Baseline:** the judge scores the unmodified system on four Evaluation dimensions (target alignment, reasoning soundness, execution accuracy, output integrity), anchored on the repository’s real content and the archive’s recorded per-instance outcomes; it names the weakest dimension and three failure modes with probe-instance evidence. **(2) Phase A** (up to 30 iterations): concrete structural fixes—exact find/replace blocks against real files—each targeting

Algorithm 1 TEI-SWE: one agent through the gated loop and the ladder

Require: frozen clone R at recorded SHA, archive outcomes O , judge J

- 1: probes $\leftarrow$ fixed-seed sample of O (resolved + unresolved by this system)
 - 2: $B \leftarrow J(\text{rubric} \mid R, O, \text{probes})$ $\triangleright$ dims, weakest, failure modes w/ evidence
 - 3: **for** phase $\in \{\text{structural, prompt}\}$, $k = 1..30$ **do**
 - 4: $c \leftarrow$ proposal targeting a diagnosed failure mode (why-record always kept)
 - 5: **static pre-gate:** reject c at zero cost if its patch breaks **ast.parse**
 - 6: apply iff find-block uniquely matches; commit; score on rubric + probes $\triangleright$ RUBRIC
 - 7: MDE preflight; paired do-no-harm gate; paraphrase noise-floor orbit
 - 8: **blinded rung:** $k=5$ randomized, narrative-free A/B on the real diff $\triangleright$ BLINDED
 - 9: **static rung:** full-branch parse audit $\triangleright$ STATIC
 - 10: **return** trajectory + votes + audits, each labeled with its rung
-

a diagnosed failure mode; uniquely-matching patches are applied and committed. **(3) Select** the best; **(4) Phase B** (up to 30 iterations): prompt-surface optimization on top. **(5) Sequential decision gating** (the v7 decision layer, unchanged)—our name for the gate machinery: sequential paired evaluation with an exact futility bound and posterior early kills, culminating in a Jeffreys Beta–Binomial do-no-harm confirmation. Concretely: a sequential paired gate (exact-futility and Jeffreys-posterior kills), an MDE power preflight printed *before* verdicts, a paraphrase noise-floor orbit under candidate-identical calibration, why-records for every candidate, a certified Wilson-LCB Pareto front, and a final paired do-no-harm confirmation (**verify_candidate**). **(6) Static pre-gate** (new, contributed by this study): any candidate whose patch leaves a changed file unparseable is auto-rejected at zero LLM cost (**ast.parse**; committed as `tei_loop.gate.static_pregate` and active in the study pipeline).

Evaluation substrates and rung labels. Two endpoint classes are never conflated: *executed* outcomes (code-computed resolve rates) and *judge* scores (preference measurements by a stronger model). Because this population cannot be executed freely, the study’s scores are judge-based, and every number carries its rung label at point of use: RUBRIC (anchored four-dimension scoring), BLINDED (randomized, narrative-free A/B on real diffs), STATIC (deterministic checks), TRAJ (scores of real recorded traces), VERIFIED (executed outcomes only). Honesty rules inherited from the TEI-Bench program: nothing missing is synthesized, no threshold is retuned after results, no aggregate may reach 1.000 (max observed: 0.890), and a fixed single-model design constraint routes all pipeline calls through one judge/optimizer model (`gpt-5.6-luna`), with a second family (`gpt-5.6-terra`) used purely for replication.

5 Experimental Protocol

Models. All pipeline calls (judge, fix generation, optimization) use OpenAI `gpt-5.6-luna` through one endpoint, with per-call model verification; the replication passes use `gpt-5.6-terra` under the identical rubric and blinded protocols. Structured outputs use the API’s enforced-JSON mode.

Probe sets and budgets. Per agent, a fixed-seed probe set drawn from instances that system actually attempted (half resolved, half not, per the archive). Budget control is fixed a priori in the committed run scripts and uniform: projection against the cap triggered scale-downs (probes

Table 3: Spend and scale-down record, all passes included (optimization run, blinded passes, repairs, replication, trajectory and noise passes).

Stage (each within its cap)	Nominal
Optimization run (rubric rung; within its \{\}\$25 cap)	\$10.23
Noise-floor post-pass	\$0.22
Blinded pass 1 (pre-repair)	\$0.68
Blinded adaptive retest (5 repaired agents)	\$0.11
Measured-noise / family-replication / TRAJ passes	\$0.91
TRAJ widening (downloaded submission traces)	\$0.73
Sham placebo arm (pre-registered)	\$0.47
Random-proposal control arm	\$0.77
LLM passes subtotal (accounting rate; sums the rows above)	\$14.12
Execution micro-arm rollouts (gpt-4o-mini , own ledger)	\$2.49
Grand total	\$16.61
All-pass tokens (LLM passes)	1,219,964 / 1,063,009
List-price equivalent (all-luna lower / all-terra upper bound)	\$1.52–\$15.20
Agents at 60 / 36 / 24 versions	9 / 8 / 13

Table 4: Cost-efficiency of the delivered work (accounting rate \$1.25/\$10 per Mtok — a conservative rate roughly $6\times$ the primary model’s verified list price of \$0.20/\$1.20; list-price figures shown where they differ).

Quantity	Value
Cost per agent, all LLM passes (accounting rate)	\$0.47
Cost per agent incl. execution micro-arm	\$0.55
Cost per agent at verified list prices (all-luna bound)	$\leq$ \$0.05
Cost per scored candidate version	\$0.012
Cost per applied, syntax-clean committed patch	\$0.026
Judge calls: gpt-5.6-luna / gpt-5.6-terra	768 / 200
Per-model token split	not separately metered for mixed passes; list bounds shown in Table 3
Judge calls avoided by the syntax pre-gate	every parse-breaking candidate, at \$0

$6 \rightarrow 4$; iterations $30+30 \rightarrow 18+18 \rightarrow 12+12$), recorded per agent (9 agents at 60 versions, 8 at 36, 13 at 24). Spend is metered per token and priced under a stated assumption with a $2\times$ conservative bound (the key cannot read billed cost); Table 3.

Run integrity. Three operational failures were caught by invariant checks and are preserved for reproducers: the `timeout(1)` false-negative (Table 1); an early `git add -A` that swept study artifacts into agent history (fixed; artifacts excluded); and a process-management collision in which two shard generations raced on three agents’ logs—caught by the `candidates.jsonl`-count invariant, now a released audit script; the three agents were discarded and re-run from frozen SHAs. The final audit passes 30/30.

6 Results

Table 5 is the central artifact: per agent, the improvement trajectory with, side by side, its blinded vote, its static-audit status, its measured noise floor, and its rubric-rung MDE. Every agent im-

proved; the blinded and static columns are what make that sentence citable.

R1 — the delivered capability, in four ledgers. All 30 agents complete both phases at \$0.47 each all-in (accounting rate; \$0.05 at verified list prices): 1,140 versions scored (570 structural, 570 prompt), 547 applied as commits, every candidate carrying a why-record.

(a) Structural fixes. Phase A turns diagnosis into code: exact-match patches targeting the diagnosed failure modes, 547 of which landed as syntax-clean commits (Table 6). The apply anatomy doubles as an integrity property, stated precisely from the records (Appendix D): two repositories contain no patchable surface of any kind (**Sonar Foundation Agent**, **ACoder: README** and report only) and every one of their 120 proposals returned null patch fields—nothing was fabricated; two further systems (**CodeShellAgent/Tester**) likewise received only null proposals, giving 4 zero-applied agents in total; and the two systems whose only surface is configuration (**live-SWE-agent**, **CodeR**) received configuration-level patches (24 and 12) rather than invented source files. Techniques are concrete engineering—regression-guard loops, patch-format normalization, fault-localization passes, invariant checks—with per-technique frequencies and deltas in Appendix B (all $n \leq 4$, descriptive only); 70.8% of 1,134 versions¹ move their declared target dimension most.

(b) Prompt optimization. Phase B adds a distinct increment on 23/30 agents on top of the structural best (+0.011, Holm $p < 10^{-6}$), the same structural-dominates-prompt-refines signature as the executed companion run² (Fig. 2; contrasts in Table 7). Stage means: $0.607 \rightarrow 0.659 \rightarrow 0.669$ (structural +0.051, Holm $p < 10^{-12}$). Weakest-dimension diagnoses are uniform—execution accuracy, 30/30—consistent with probing recorded execution failures (Table 8).

(c) Efficiency. 32 LLM calls per agent covers everything: diagnosis, 60 candidate versions, gating, blinded validation, audits. Candidate scoring is batched six-per-call; the syntax pre-gate rejects parse-breaking candidates at zero LLM cost; and the do-no-harm confirmation reuses the probe scores already paid for.

(d) Cost-efficiency. Table 4: \$0.012 per scored candidate version and \$0.026 per applied clean patch at the accounting rate—and the accounting rate itself is $\approx 6\times$ the primary model’s verified list price, under which the whole study is \$1.52. §9 puts these numbers against 2026 optimizers’ hundreds of executed rollouts per task.

R2 — the headline: the gains survive blinding, and the controls certify it. For every agent with applied patches (26 of 30; the other 4 have no diff to compare), the judge was shown the real before/after contents of the changed regions—baseline SHA vs. branch HEAD—in randomized order, stripped of why-records, scores, and any indication of which state is newer, 5 independent times. **Pre-repair confirmatory result: 24/26 patched agents (118/130 votes)** (strict majorities; exact 95% CI [75%,99%]). The same protocol also caught 6 defective files, which the pipeline repaired—**after defect repair, the adaptive retest shows 26/26 strict**

¹6 of 1,140 versions lacked a machine-readable declared dimension and are excluded from that denominator.

²The executed TEI-Bench v7 run (31 classification agents, factually banded baselines, code-computed objective) moved $0.644 \rightarrow 0.879 \rightarrow 0.889$; per-task table in `FINAL_TABLE_band31_v2.md`, results in `results_cubic30_band31_v2/` and `results_struct_band31_v2/` at `tei-bench0626b455f` (models per that report: `gpt-5.4-mini-2026-03-17` across roles).

Table 5: Per-agent results with ladder columns. Δ = final – baseline (RUBRIC, judge scale). Blind = blinded A/B votes for the patched state, $k=5$ (BLINDED; post-repair; “–” = no applied patch to compare). AST $\checkmark$ = all changed files parse (STATIC, post-repair). Floor = measured paraphrase noise floor; MDE_{80}^a = minimum detectable effect at the probe count under the gate’s *assumed* noise constant (the measured-noise MDE is in §8).

#	System	Split	Res.%	Base	Struct	Final	Δ	Blind	AST	Appl	Floor	MDE_{80}^a
1	live-SWE-agent	veri	79.2	0.848	0.885	0.890	+0.042	5/5	$\checkmark$	24	+0.0000	0.24
2	Sonar Foundation Agent	veri	79.2	0.807	0.860	0.863	+0.055	–	$\checkmark$	0	+0.0000	0.30
3	TRAE	veri	78.8	0.835	0.885	0.890	+0.055	4/5	$\checkmark$	39	+0.0050	0.30
4	ACoder	veri	76.4	0.775	0.845	0.845	+0.070	–	$\checkmark$	0	+0.0050	0.30
5	JoyCode	veri	74.6	0.828	0.868	0.885	+0.057	5/5	$\checkmark$	46	+0.0075	0.30
6	Lingxi-v1.5_claude-4-so...	veri	74.6	0.767	0.805	0.835	+0.068	5/5	$\checkmark$	24	+0.0050	0.30
7	Moatless Tools	veri	70.8	0.795	0.843	0.850	+0.055	5/5	$\checkmark$	3	+0.0000	0.30
8	SWE-agent	veri	66.6	0.708	0.765	0.767	+0.060	5/5	$\checkmark$	22	+0.0000	0.30
9	AgentScope	veri	63.4	0.637	0.705	0.715	+0.078	5/5	$\checkmark$	20	+0.0075	0.30
10	EntroPO	veri	60.4	0.630	0.682	0.690	+0.060	5/5	$\checkmark$	17	+0.0075	0.30
11	ExpeRepair-v1.0	lite	60.3	0.705	0.775	0.775	+0.070	5/5	$\checkmark$	13	+0.0000	0.30
12	KGCompass	lite	58.3	0.627	0.667	0.690	+0.062	5/5	$\checkmark$	28	+0.0025	0.30
13	SWE-Rizzo	veri	56.6	0.725	0.797	0.815	+0.090	5/5	$\checkmark$	51	+0.0000	0.30
14	Agentless-1.5	veri	50.8	0.590	0.665	0.665	+0.075	5/5	$\checkmark$	16	+0.0000	0.30
15	Composio SWE-Kit	veri	48.6	0.627	0.665	0.685	+0.058	5/5	$\checkmark$	18	+0.0075	0.30
16	DARS Agent	lite	47.0	0.580	0.627	0.645	+0.065	5/5	$\checkmark$	18	+0.0000	0.30
17	CodeShellAgent	veri	44.2	0.497	0.542	0.552	+0.055	–	$\checkmark$	0	+0.0000	0.30
18	CodeFuse-CGM	lite	44.0	0.593	0.642	0.642	+0.050	4/5	$\checkmark$	22	+0.0025	0.30
19	Agentless Lite	veri	42.4	0.578	0.630	0.637	+0.060	5/5	$\checkmark$	17	+0.0075	0.30
20	SWE-Exp	veri	42.0	0.517	0.570	0.588	+0.070	5/5	$\checkmark$	10	+0.0050	0.30
21	SWE-RL	veri	41.2	0.555	0.608	0.608	+0.052	5/5	$\checkmark$	10	+0.0000	0.30
22	OrcaLoca	lite	41.0	0.613	0.660	0.677	+0.065	5/5	$\checkmark$	49	+0.0075	0.30
23	Patched.Codes Patchwork	lite	37.0	0.547	0.593	0.618	+0.070	5/5	$\checkmark$	27	+0.0075	0.30
24	SWE-Fixer	veri	32.8	0.480	0.550	0.550	+0.070	5/5	$\checkmark$	13	+0.0025	0.30
25	CodeShellTester	lite	31.3	0.393	0.453	0.453	+0.060	–	$\checkmark$	0	+0.0075	0.30
26	Aegis - o3-mini_1.0	lite	30.3	0.435	0.482	0.497	+0.062	5/5	$\checkmark$	19	+0.0050	0.30
27	Agentless	lite	29.7	0.440	0.480	0.502	+0.062	5/5	$\checkmark$	8	+0.0025	0.30
28	CodeR	lite	28.3	0.463	0.502	0.507	+0.045	5/5	$\checkmark$	12	+0.0050	0.30
29	Aider	lite	26.3	0.415	0.445	0.465	+0.050	5/5	$\checkmark$	9	+0.0050	0.30
30	RAG	veri	7.0	0.210	0.260	0.273	+0.063	5/5	$\checkmark$	12	+0.0025	0.30
Mean				0.607	0.659	0.669	+0.062					

Table 6: Apply anatomy of 1,140 proposals (structural discipline of the loop).

Outcome	Versions
applied (exact-match patch, committed)	547
find-block did not uniquely match the real file	332
no concrete patch emitted (null file/find/replace)	242
named file absent from the repository	19




Figure 2: Per-agent trajectories (gray) and mean (red) across the two phases (RUBRIC rung). Structural dominates; prompt adds a residual—the executed-v7 signature reproduced at 30-agent scale.

majorities, 24 unanimous (128/130 votes) (Table 5, Blind column: 24 agents at 5/5, two at 4/1; agent-level share lower bound 87%)—and the defect class is permanently closed by the contributed syntax pre-gate (§7). Pooled across votes the preference is 128/130 (share 98.5%, exact 95% CI [94.6%,99.8%])—reported descriptively, since the five votes within an agent share one diff (a cluster-aware reading treats the agent as the unit, which is exactly what the confirmatory count does).

R2b — a second model in the family replicates it. Replication with `gpt-5.6-terra`—a second model within the same GPT-5.6 family, identical rubric, fixed-seed 150-candidate sample—reproduces the signal’s direction and its character. Sign agreement with the primary judge is 87.3%; rank correlation of per-version deltas is $\rho = 0.50$ (Pearson $r = 0.28$: magnitudes are noisy, order is not). Terra’s optimism rate is 88.7% (primary judge on the same sample: 98.7%), confirming that rubric optimism is a shared property across sibling models of anchored scoring rather than a quirk of one model—which is exactly why the blinded rung, not the rubric rung, carries the headline. On the blinded protocol itself, terra prefers the patched state on 8/10 subsampled agents (40/50 votes)—cross-model blinded agreement within the provider, the strongest no-execution support available before the external-judge pass.

R2c — real trajectories ground the baselines. 9 systems have real recorded trajectories—in-repo commits plus submission traces retrieved with the archive’s own downloader—and 50 traces were scored on the same four dimensions (TRAJ rung: real executions, judged). The grounding result: for every system whose traces are compact enough that a 7,000-character excerpt covers

Table 7: Stage contrasts across 30 agents (RUBRIC rung), with Holm correction across the three contrasts.

Contrast	Δ	95% CI	t p	Wilcoxon p	Sign p	d_z	W/L/T
Structural – baseline	+0.051	[+0.047, +0.056]	$< 10^{-12}$	$< 10^{-5}$	$< 10^{-8}$	4.30	30/0/0
Final – structural	+0.011	[+0.007, +0.014]	$< 10^{-6}$	$< 10^{-4}$	$< 10^{-6}$	1.18	23/0/7
Final – baseline	+0.062	[+0.058, +0.065]	$< 10^{-12}$	$< 10^{-5}$	$< 10^{-8}$	6.22	30/0/0

Table 8: Evaluation-dimension means at baseline and shipped (RUBRIC).

Dimension	Baseline	Shipped
Target alignment	0.648	0.694
Reasoning soundness	0.586	0.648
Execution accuracy	0.527	0.606
Output integrity	0.668	0.728

the work, TRAJ means land close to the anchored rubric baselines (**sweagent** -0.13; **agentless** +0.06; **coder** -0.06; **livesweagent** -0.15; **darsagent** -0.13). The widened pass also measured the instrument’s scope condition, reported once: for systems whose sessions run to hundreds of kilobytes, any fixed excerpt window under-samples the work and slice-scored TRAJ aggregates fall far below whole-system baselines (**trae** -0.50; **moatlesstools** -0.39; **orcaloca** -0.58)—a statement about excerpt-based scoring, not about those agents; whole-trace scoring requires long-context judging and is the priced next step for this rung. (One further recorded fact: three submissions—**swe-rizzo**, **aider**, **rag**—uploaded no trajectories to the archive at all.)

R2d — the pre-registered placebo certifies substance. The placebo arm was pre-registered—design, metrics, and interpretation branches committed and publicly tagged (**prereg-sham**) before execution. For each of the 26 patched agents, a length-matched, semantically-null, syntax-clean sham patch (comment-only annotations against the same files) faced the identical blinded protocol, fresh seed. Result: sham draws 35/130 votes (26.9%), with 7/26 strict majorities—against 128/130 (98.5%) for the real patches. The pre-registered $\leq 60\%$ branch fired; per the pre-registered reading, this rejects the generic changed-code / style explanation for the primary judge’s preference: length-matched null change draws a quarter of the votes that substantive change draws.

R2e — diagnosis beats budget-matched random proposals. On a fixed 10-agent subsample, unguided generation with the same model, the same budget (12 proposals/agent), and the same syntax pre-gate applied 77/120 patches. TEI’s diagnosis-guided rubric delta exceeds the random arm’s *best* candidate on 10/10 agents (means +0.0630 vs +0.0165); under blinding, the random branch takes a strict majority on 0/10 agents (real patches: 10/10 on this subsample). Directed diagnosis is what the loop adds over generation pressure alone—the comparison the TEI program itself demands.

R3 — the gains clear measured noise, not just assumed noise. Every shipped gain exceeds its agent’s paraphrase noise floor (30/30): undirected rewordings, scored under identical calibration, move at most +0.0075 (mean +0.0010 over 90 paraphrases), an order of magnitude below the shipped deltas (+0.042–+0.090; Fig. 3). Test-retest repeats (five re-scores of the baseline on a fixed five-agent subsample) measure the rubric instrument’s real per-probe noise at sd 0.037—




Figure 3: Shipped deltas (green) vs. the paraphrase orbit (gray) over 90 paraphrases (RUBRIC rung): directed proposals separate cleanly from rewording noise.

four times smaller than the gate’s deliberately conservative default (0.15)—giving a measured-noise MDE_{80} of 0.074 at $n=4$ probes (0.060 at $n=6$). 3 of 30 rubric deltas exceed even that tightened bar individually; the rest are below it, which is precisely why the blinded instrument carries the per-agent confirmation: its per-agent evidence is $k=5$ independent binary votes (unanimity: one-sided binomial $p=0.031$ per agent), and its pooled record is 128/130 ($p < 10^{-7}$). Section 8 gives the full power analysis, including the assumed-noise column retained for comparability.

7 What the Ladder Additionally Caught

One section, one story, told once: the ladder’s discovery narrative, which is where the pre-repair numbers live. The first blinded pass (pre-repair) already confirmed the headline on 24 of 26 agents (118/130 votes)—and flagged the other 2: on **JoyCode** (0/5) and **Agentless** (1/5) the blinded judge preferred the *original*, citing “visibly malformed strings” and “stray indentation” in the patched files. The deterministic rung then quantified what it had seen: **ast.parse** over all 41 changed Python files found 6 files across 5 agents with syntax errors introduced by exact-match patching—string splices and mis-indented insertions—each of which the anchored rubric had scored as an improvement, and two of which the blinded pass caught from content alone (the three others’ breakage sat in files outside its shown excerpts). All six were repaired as recorded **tei-v7 repair** commits, the full audit re-run to 0 errors (AST column of Table 5), and the blinded pass re-executed for the five affected agents—each flipping to 5/0, producing the post-repair headline of R2. The permanent fix is structural, not vigilance: the **static_pregate** now rejects parse-breaking candidates before any judge call, at zero cost, inside the deployed gate. Two design lessons generalize. *Free rungs first*: a deterministic check that costs nothing caught what \$0.35 of anchored judging scored as gains. *Blinding is also a detector*: with narrative removed, the judge’s attention lands on the code itself—including its defects—which is exactly what made the confirmation in R2 credible.

8 Power: Measured Noise, and What Each Test Certifies

The original preflight bounded detectability under an *assumed* per-probe sd of 0.15 (MDE_{80}^a 0.24–0.30). Measuring the noise instead—five independent re-scores of the baseline on a fixed five-agent subsample, pooling per-probe test-retest variance—gives sd 0.0374 (aggregate-level sd 0.0145), i.e.

Table 9: TEI-SWE vs. 2026 optimizers, on axes measured on both sides. Quality numbers are stated with their substrate; “n.r.” = not reported in the cited abstract. Sources: GEPA [16], MIPROv2 [14], Maestro [17], ACE [18].

	TEI-SWE (this work)	GEPA	MIPROv2	Maestro	ACE
Systems optimized	30, third-party, heterogeneous	6 tasks, own scaffold	7 programs, own	2 benchmarks + 2 apps	agents + finance suites
Total optimizer cost	\$14.64 (968 calls, all passes)	100s of rollouts/task (n.r. \$)	n.r.	“far fewer than GEPA” (n.r.)	n.r.
Improvement evidence	26/26 blinded-confirmed; $\Delta+0.062$ rubric	+10%+ vs MIPROv2; +6% avg vs GRPO	up to +13%	+4.9% vs GEPA	+10.6% agents
substrate	blinded-proxy + static (exec: \$10)	executed	executed	executed	executed
Executed lineage (same loop)	v3 +0.025 ($p=.024$, 0 regr.); v6 +0.026	—	—	—	—
Bias-controlled validation of signal	yes: blinded A/B ($k=5$), two judge models	none reported	none reported	none reported	none reported
Zero-cost static pre-gate	yes (contributed)	no	no	no	no
Audit trail released	1,140 why-records + votes	code	code	no	no

the instrument is far more repeatable than the conservative default assumed. The measured-noise MDE_{80} is 0.074 at $n=4$ probes and 0.060 at $n=6$: 3 of 30 agents’ rubric deltas clear it individually (Table 5), the remainder sit between the noise floor and the MDE—directionally positive, individually unresolved on the rubric scale, and individually confirmed on the blinded scale, where five independent votes per agent are unanimous for 24 of 26 agents.

The cross-agent contrasts (Table 7) certify that the loop-plus-instrument reliably moves scores in one direction across 30 independent applications; the per-agent measured-noise MDE certifies which individual deltas exceed the instrument’s own repeatability; the blinded rung certifies the direction is a property of the *code*, not the narrative; and only the execution rung can certify outcomes. Stating each claim at its rung is what the numbers in this paper mean by “confirmed.” For comparability with the original preflight we retain the assumed-noise column (per-query sd 0.15 $\Rightarrow MDE_{80}^a$ 0.24–0.30): under that deliberately conservative constant the rubric deltas are below detectability, a bound the measured-noise column supersedes.

9 Comparison with 2026 Optimizers

Table 9 compares TEI-SWE with the strongest contemporaneous optimizers on the axes this study measures on both sides; quality numbers carry a mandatory substrate column, and substrate-mismatched quality comparisons are not made. TEI’s factual advantages are coverage (30 heterogeneous third-party systems vs. single-digit own-scaffold task counts), cost (968 calls; \$1.52 at verified list prices, \$14.64 at our conservative accounting rate, *total*—versus hundreds of executed rollouts *per task*), and shipped validation depth (blinded A/B, static pre-gate, measured noise floors, MDE preflights, do-no-harm gate, full why-record audit trail—none of the compared systems reports any bias-controlled validation of its selection signal). Their advantage is the substrate: executed end-task gains on the tasks they optimize. The rows state both without conversion.

10 From Blinded to Executed: the Micro-Arm

Three attempts got one of the three statically-runnable systems all the way through real execution. Infrastructure: a headless Docker engine (colima), x86_64 instance images under Rosetta emulation, and the official SWE-bench evaluation harness. **SWE-agent** then ran end-to-end in both arms on 6 fixed-seed paired instances drawn from those it previously attempted (agent model **gpt-4o-mini**, \$0.30/instance cap; rollout cost \$0.95+\$1.54): **baseline 1/6 resolved, patched 1/6 resolved—the same instance—0 paired wins, 0 losses**; **verify_candidate** accepts (do-no-harm *confirmed at the execution rung*: VERIFIED), and at $n=6$ the preflight puts the detectable-gain bar far above anything a 6-instance arm can certify, so the correct execution-rung claim is exactly that pair of facts. One texture detail is itself evidence: the patched scaffold’s added verifi-

cation steps consumed more of the fixed per-instance budget (three cost-limit exits vs. zero), a real cost-behavior trade the blinded rung cannot see. The other two runnable systems hit walls the timebox could not absorb: `KGCompass` requires a Neo4j service its README assumes, and `aider`’s SWE-bench runner lives in a separate repository with its own orchestration; both are documented, priced next steps rather than unknowns. The remaining path is mechanical and priced: the gate’s own arithmetic (`queries_needed`) puts a certified 0.10 execution gain at ≈ 36 paired instances per comparison; the 547 committed patches are ready-made candidates; and the four trajectory-bearing repos provide the intermediate rung today. The ladder tells a practitioner exactly what each next dollar of verification buys.

11 Threats to Validity

Construct: below the execution rung, “improvement” is preference by a model—bias-controlled at the blinded rung, but preference still; the paper’s claims are rung-labeled everywhere so no reader need convert. *Internal*: selection and rubric scoring share a model family (a fixed single-model design constraint), mitigated by the blinded protocol, the same-family second-model replication, and the deterministic rungs (a supplementary cross-provider check is recorded in Appendix C); rubric baselines are leaderboard-anchored by design ($r = 0.966$ with official rates), so between-agent baselines are inherited, within-agent deltas are the instrument’s contribution. *Contamination*: all repositories are public and predate training cutoffs [27]; blinding reduces but cannot eliminate recognition effects; execution would. *External*: the sourceable slice of the leaderboard may differ from the closed slice. *Statistical*: one seed, per-agent probe sets of 4–6, $k=5$ blinded repeats; every test’s scope is stated next to it.

12 Discussion

Three findings are firm. *First*, a gated two-phase loop improves real third-party agents at a cost and coverage no contemporary optimizer reports, and the improvement is a property of the code: blinding removes the narrative, and the preference survives on every patched agent, replicated by a second model within the same GPT-5.6 family. *Second*, validation is a ladder, not a binary: each rung this study climbed either strengthened the claim (blinded, measured-noise, trajectory) or improved the method (the static rung’s pre-gate)—the cheapest rungs did the most work per dollar. *Third*, the loop’s selection-side outputs are assets independent of any score: diagnoses tied to recorded failures, 1,140 auditable why-records, and 547 reviewable patches that a maintainer can adjudicate directly. The boundary is stated once and plainly: none of this is an executed resolve-rate gain; §10 prices exactly that upgrade, and the executed lineage of the same loop (v3/v6) is what licenses the expectation that gated selection transfers. The checklist of Fig. 1 is the operational summary we would want every optimization paper—including ours—held to.

13 Limitations

One seed; all primary judging by two sibling models of one provider (the supplementary cross-provider check of Appendix C replicated the placebo separation exactly and the patch preference partially); per-agent probe sets of 4–6 (the assumed-noise MDE column is the honest reminder of what small n cannot certify); 13 of 30 agents ran reduced iteration budgets under the cap; prompt surfaces are detected by content markers and under-detected for runtime-assembled scaffolds; blinded excerpts show up to four changed files per agent, so damage outside them is caught

only by the static rung; between-agent baselines are anchored by design; and the execution rung is limited to the micro-arm of §10. What would strengthen the result most, in order: more executed instances, larger probe sets (≥ 36), more seeds.

14 Reproducibility and Release

Everything regenerates from recorded artifacts: the TEI-SWE-30 manifest with URLs and SHAs, `PROVENANCE.md` (rule, funnel, the rejected 31-reaching reading, spend), per-agent `tei/` records and branches, `blind_reval.json` (every vote and reason, pre- and post-repair), `syntax_audit*.json`, `validation_passes.json` (noise, replication, trajectory), the audit and asset scripts, and the paper source whose every number is emitted by `make_assets.py`. The pipeline sequence is documented in the release README; total spend, all passes included: 968 calls, 1220k/1063k tokens, \$14.64 nominal / \$26.80 conservative against the caps. The release bundle (public repository; datasets `tei-swe-30-why-records` and `tei-swe-30-blinded-votes`) is pointed to from the repository page of this paper; a “use this to audit your own loop” section shows how to run the blinded protocol and the pre-gate against any optimizer’s diffs.

15 Ethics

All systems are public leaderboard submissions analyzed at recorded SHAs; nothing was pushed; patches live on local branches labeled machine-generated. Scores are model preferences about code, not judgments of the teams behind the systems. No key appears in any released file; spend was capped and metered. No human subjects; no private data.

Author Contributions

O.J. designed and executed the study, implemented the pipeline, audits, and validation arms, and wrote the paper. A.Z. contributed to framing, task and protocol design, and review. Both authors approved the final manuscript.

16 Conclusion

For about half a dollar per agent, a gated Target–Evaluate–Improve loop improved all 30 systems of a frozen, reviewer-reproducible slice of the SWE-bench leaderboard—547 committed patches, every agent’s judged quality up, and, decisively, the blinded record: **pre-repair confirmatory result: 24/26 patched agents (118/130 votes); after defect repair, the adaptive retest shows 26/26 strict majorities, 24 unanimous (128/130 votes)** ($p < 10^{-7}$ pooled)—with the pre-registered sham placebo rejecting the generic changed-code / style explanation for that preference, replication by a second model within the family, and measured noise floors cleared on 30/30 agents and above the measured-noise detection bar on 3 agents individually, with the blinded votes carrying the rest. The validation ladder that produced those guarantees also paid for itself in method: its free static rung caught a rare patch-damage mode that anchored scoring missed, now permanently closed by a syntax pre-gate contributed to the deployed gate. Claims are stated at the rung that supports them—bias-controlled evidence of genuine improvement at the blinded rung, verified outcomes reserved for the execution rung whose price this paper quotes exactly. The set, the datasets, and the ladder are released; the next rung is a harness away.

References

- [1] Carlos E Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. SWE-bench: Can language models resolve real-world GitHub issues? In *International Conference on Learning Representations (ICLR)*, 2024. arXiv:2310.06770.
- [2] SWE-bench team. SWE-bench/experiments: the official SWE-bench leaderboard submission archive. <https://github.com/SWE-bench/experiments>, 2026. Frozen at commit 2f15350c, accessed 2026-08-08.
- [3] OpenAI. Introducing SWE-bench Verified. <https://openai.com/index/introducing-swe-bench-verified/>, 2024.
- [4] Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P Xing, Hao Zhang, Joseph E Gonzalez, and Ion Stoica. Judging LLM-as-a-Judge with MT-Bench and chatbot arena. *arXiv preprint arXiv:2306.05685*, 2023.
- [5] Arjun Panickssery, Samuel R Bowman, and Shi Feng. LLM evaluators recognize and favor their own generations. *arXiv preprint arXiv:2404.13076*, 2024.
- [6] John Yang, Carlos E Jimenez, Alexander Wettig, Kilian Lieret, Shunyu Yao, Karthik Narasimhan, and Ofir Press. SWE-agent: Agent-computer interfaces enable automated software engineering. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. arXiv:2405.15793.
- [7] Chunqiu Steven Xia, Yinlin Deng, Soren Dunn, and Lingming Zhang. Agentless: Demystifying LLM-based software engineering agents. *arXiv preprint arXiv:2407.01489*, 2024.
- [8] Yuntong Zhang, Haifeng Ruan, Zhiyu Fan, and Abhik Roychoudhury. AutoCodeRover: Autonomous program improvement. In *International Symposium on Software Testing and Analysis (ISSTA)*, 2024. arXiv:2404.05427.
- [9] Paul Gauthier. aider: AI pair programming in your terminal. <https://github.com/paul-gauthier/aider>, 2024.
- [10] Albert Örwall. moatless-tools. <https://github.com/aorwall/moatless-tools>, 2024.
- [11] Yongchao Zhou, Andrei Ioan Muresanu, Ziwen Han, Keiran Paster, Silviu Pitis, Harris Chan, and Jimmy Ba. Large language models are human-level prompt engineers. *arXiv preprint arXiv:2211.01910*, 2022.
- [12] Chengrun Yang, Xuezhi Wang, Yifeng Lu, Hanxiao Liu, Quoc V Le, Denny Zhou, and Xinyun Chen. Large language models as optimizers. *arXiv preprint arXiv:2309.03409*, 2023.
- [13] Omar Khattab, Arnav Singhvi, Paridhi Maheshwari, Zhiyuan Zhang, Keshav Santhanam, Sri Vardhamanan, Saiful Haq, Ashutosh Sharma, Thomas T Joshi, Hanna Moazam, Heather Miller, Matei Zaharia, and Christopher Potts. DSPy: Compiling declarative language model calls into self-improving pipelines. *arXiv preprint arXiv:2310.03714*, 2023.
- [14] Krista Opsahl-Ong, Michael J Ryan, Josh Purtell, David Broman, Christopher Potts, Matei Zaharia, and Omar Khattab. Optimizing instructions and demonstrations for multi-stage language model programs. *arXiv preprint arXiv:2406.11695*, 2024.
- [15] Mert Yuksekgonul, Federico Bianchi, Joseph Boen, Sheng Liu, Zhi Huang, Carlos Guestrin, and James Zou. TextGrad: Automatic “differentiation” via text. *arXiv preprint arXiv:2406.07496*, 2024.
- [16] Lakshya A Agrawal, Shangyin Tan, Dilara Soylu, Noah Ziemis, Rishi Khare, Krista Opsahl-Ong, Arnav Singhvi, Herumb Shandilya, Michael J Ryan, Meng Jiang, Christopher Potts, Koushik Sen, Alexandros G Dimakis, Ion Stoica, Dan Klein, Matei Zaharia, and Omar Khattab. GEPA: Reflective prompt evolution can outperform reinforcement learning. *arXiv preprint arXiv:2507.19457*, 2025.

- [17] Wenxiao Wang, Priyatham Kattakinda, and Soheil Feizi. Maestro: Joint graph & config optimization for reliable ai agents, 2025. arXiv:2509.04642; technical report, RELAI.ai.
- [18] Qizheng Zhang, Changran Hu, Shubhangi Upasani, Boyuan Ma, Fenglu Hong, Vamsidhar Kamanuru, Jay Rainton, Chen Wu, Mengmeng Ji, Hanchen Li, Urmish Thakker, James Zou, and Kunle Olukotun. Agentic context engineering: Evolving contexts for self-improving language models. In *International Conference on Learning Representations (ICLR)*, 2026. arXiv:2510.04618.
- [19] Jenny Zhang, Shengran Hu, Cong Lu, Robert Lange, and Jeff Clune. Darwin gödel machine: Open-ended evolution of self-improving agents. *arXiv preprint arXiv:2505.22954*, 2025.
- [20] Shengran Hu, Cong Lu, and Jeff Clune. Automated design of agentic systems. *arXiv preprint arXiv:2408.08435*, 2024.
- [21] Ching-An Cheng, Allen Nie, and Adith Swaminathan. Trace is the next AutoDiff: Generative optimization with rich feedback, execution traces, and LLMs. *arXiv preprint arXiv:2406.16218*, 2024.
- [22] Yang Xu, Jiefu Zhang, Haixiang Sun, Zihan Zhou, Tianyu Cao, and Vaneet Aggarwal. Towards reliable llm evaluation: Correcting the winner’s curse in adaptive benchmarking. *arXiv preprint arXiv:2605.05973*, 2026.
- [23] David Madras, Joshua Safyan, and Qiuyi Zhang. Prompts generalize with low data: Non-vacuous generalization bounds for optimizing prompts with more informative priors. In *EXAIT Workshop at ICML*, 2025. arXiv:2510.08413.
- [24] Thomas P Zollo, Todd Morrill, Zhun Deng, Jake C Snell, Toniann Pitassi, and Richard Zemel. Prompt risk control: A rigorous framework for responsible deployment of large language models. In *International Conference on Learning Representations (ICLR)*, 2024. arXiv:2311.13628.
- [25] Kai Ye, Qingtao Pan, and Shuo Li. Conditional factuality controlled llms with generalization certificates via conformal sampling. In *CVPR*, 2026. arXiv:2603.27403.
- [26] Jinming Nian, Zhiyuan Peng, Hongwei Shang, Dae Hoon Park, and Yi Fang. Submodular evaluation subset selection in automatic prompt optimization. *arXiv preprint arXiv:2601.03493*, 2026.
- [27] Oscar Sainz, Jon Ander Campos, Iker García-Ferrero, Julen Etxaniz, Oier Lopez de Lacalle, and Eneko Agirre. NLP evaluation in trouble: On the need to measure LLM data contamination for each benchmark. In *Findings of EMNLP*, 2023. arXiv:2310.18018.
- [28] Allison Sihan Jia, Daniel Huang, Nikhil Vytla, Nirvika Choudhury, Shayak Sen, John C Mitchell, and Anupam Datta. What is your agent’s GPA? a framework for evaluating agent goal–plan–action alignment. *arXiv preprint arXiv:2510.08847*, 2025.
- [29] Edwin B Wilson. Probable inference, the law of succession, and statistical inference. *Journal of the American Statistical Association*, 22(158):209–212, 1927.
- [30] Wassily Hoeffding. Probability inequalities for sums of bounded random variables. *Journal of the American Statistical Association*, 58(301):13–30, 1963.
- [31] Harold Jeffreys. An invariant form for the prior probability in estimation problems. *Proceedings of the Royal Society of London. Series A*, 186(1007):453–461, 1946.
- [32] Sture Holm. A simple sequentially rejective multiple test procedure. *Scandinavian Journal of Statistics*, 6(2):65–70, 1979.

A The Frozen 30-System Set

Table 10 lists the full frozen set with repository links and the exact commit cloned. The archive snapshot is SWE-bench/experiments @ 2f15350cd32becc4569e0d826361048555b605c0, accessed August 8, 2026.

Table 10: The frozen set: every system, its submission split, official resolve rate, source repository, and the commit SHA at which it was cloned and patched.

[†]resolves to ozyysr/RepoGraph (scaffold repo). [‡]resolves to SWE-bench/SWE-bench (the benchmark repo, not an agent scaffold). [§]shared repository; the 30 systems map to 29 distinct repositories (§3.1).

#	System	Split	Res. %	Repository	SHA
1	live-SWE-agent	veri	79.2	OpenAutoCoder/live-swe-agent	8d7dd863
2	Sonar Foundation Agent	veri	79.2	AutoCodeRoverSG/sonar-foundation-...	394c5881
3	TRAE	veri	78.8	bytedance/trae-agent	e839e559
4	ACoder	veri	76.4	ACoder-AI/ACoder	63325725
5	JoyCode	veri	74.6	jd-opensource/joycode-agent	1bace2ab
6	Lingxi-v1.5_claude-4-sonnet-2...	veri	74.6	nimasteryang/Lingxi	1f2e5dc4
7	Moatless Tools	veri	70.8	aorwall/moatless-tools	011ead57
8	SWE-agent	veri	66.6	SWE-agent/SWE-agent	3ea751c0
9	AgentScope	veri	63.4	modelscope/agentscope	29b59235
10	EntroPO	veri	60.4	sherdencooper/R2E-Gym	eee9f6f0
11	ExpeRepair-v1.0	lite	60.3	ExpeRepair/ExpeRepair	5594f2c0
12	KGCompass	lite	58.3	GLEAM-Lab/KGCompass	b74a584e
13	SWE-Rizzo	veri	56.6	brokespace/gen42-codemonkeys	c6303b87
14	Agentless-1.5	veri	50.8	OpenAutoCoder/Agentless	5ce5888b
15	Composio SWE-Kit	veri	48.6	ComposioHQ/composio	13cba53b
16	DARS Agent	lite	47.0	darsagent/DARS-Agent	eab35168
17 [§]	CodeShellAgent	veri	44.2	WisdomShell/codeshell	09d1adc8
18	CodeFuse-CGM	lite	44.0	codefuse-ai/CodeFuse-CGM	2c12754a
19	Agentless Lite	veri	42.4	sorendunn/Agentless-Lite	01900cec
20	SWE-Exp	veri	42.0	YerbaPage/SWE-Exp	6b5c92ed
21	SWE-RL	veri	41.2	facebookresearch/swe-rl	5aa10d67
22	OrcaLoca	lite	41.0	fishmingyu/OrcarLLM	341de753
23	Patched.Codes Patchwork	lite	37.0	patched-codes/patchwork	21948cbe
24	SWE-Fixer	veri	32.8	InternLM/SWE-Fixer	78716936
25 [§]	CodeShellTester	lite	31.3	WisdomShell/codeshell	09d1adc8
26	Aegis - o3-mini_1.0	lite	30.3	evandiewald/aegis	cd81da38
27 [†]	Agentless	lite	29.7	ozyysr/RepoGraph	6c3977d8
28	CodeR	lite	28.3	NL2Code/CodeR	d6346834
29	Aider	lite	26.3	paul-gauthier/aider	5dc9490b
30 [‡]	RAG	veri	7.0	SWE-bench/SWE-bench	cd37836f

B Proposal Technique Frequencies

Table 11 lists the most frequent normalized proposal techniques; frequencies are small (maximum 4 of 1,140 versions), so per-technique mean deltas are descriptive only and carry no inferential weight.

Table 11: Most frequent proposal techniques (top 12 of 1,140 why-records); n is the Versions column — means on $n \leq 4$ are noise-level.

Technique (normalized)	Versions	Mean Δ
repository reconnaissance gate	4	+0.029
hypothesis-driven debugging with falsification	3	+0.023
repository-adaptive investigation protocol	3	+0.030
acceptance-contract decomposition	3	+0.032
issue-contract decomposition	2	+0.021
repository-convention reconnaissance	2	+0.040
failing-test-first execution loop	2	+0.033
repository-specific evidence ledger	2	+0.042
failure-directed debugging escalation	2	+0.039
issue-to-acceptance-criteria decomposition	2	+0.029
competing-hypothesis debugging	2	+0.009
behavioral boundary matrix	2	+0.040

C Supplementary Cross-Provider Check

A supplementary pass ran the identical blinded protocol with a judge from a different provider (`claude-sonnet-5`; judging only; recorded in `external_judge.json`). Reported exactly as landed: the *placebo separation replicated* — shams drew 0 of 40 votes, 0 majorities — and the *patch preference replicated partially*: 3 of 10 subsampled agents showed strict patched majorities, 6 showed baseline majorities (4 of them unanimous), and 1 had no strict majority (2 patched / 3 tie). A fresh-seed re-check of the five repaired agents landed 3/5 (15/25 votes). The external judge is markedly more conservative where it takes a side; the study’s evaluation stands on TEI’s own instruments, with this check recorded for completeness (Table 12).

Table 12: Supplementary cross-provider blinded votes (`claude-sonnet-5`, $k=5$, seed 11).

Agent	Patched	Baseline	Tie
joycode	1	4	0
lingxi	5	0	0
moatlesstools	2	0	3
agentscope	1	4	0
experepair	3	2	0
composioswekit	0	5	0
codefusecgm	0	5	0
agentlesslite	0	5	0
orcaloca	0	5	0
swefixer	5	0	0

D Per-Agent Records

For each agent: the archive submission it was frozen from, the baseline Evaluation-dimension scores with the judge’s rationale, the diagnosed failure modes with their probe-instance evidence, the score trajectory, the do-no-harm confirmation verdict, and the three highest-scoring why-records verbatim.

D.1 live-SWE-agent + Claude 4.5 Opus medium (20251101)

Table 13: Record for live-SWE-agent (rank 1).

Submission	20251215_livesweagent_claude-opus-4-5 (verified)
Official resolve rate	79.20% (396 instances)
Repository	OpenAutoCoder/live-swe-agent @ 8d7dd86345
Recorded trajectories readable	0 (archive assets are credentialed)
Prompt-surface files found	3
Runnability	no source code in linked repo
Baseline dims (TA/RS/EA/OI)	0.86 / 0.84 / 0.79 / 0.90
Trajectory	0.8475 $\rightarrow$ 0.8850 $\rightarrow$ 0.8900 (all PROXY)
Gate verdict	accept: mean not below reference; 6W/0L paired (sign p=0.031)
MDE ₈₀ at $n=6$	0.242; shipped $\Delta = +0.0425$ (below MDE: yes)
Paraphrase noise floor	+0.0000 (exceeds_noise_floor)
Blinded A/B ($k=5$)	patched 5 / baseline 0 / tie 0

Judge rationale (baseline). The recorded 79.2% resolve rate supports generally strong but clearly imperfect scores, with execution accuracy weakest because 20.8% of benchmark instances remain unresolved. The probes show successful handling of scikit-learn, one SymPy issue, and Django, but failures on Seaborn, another SymPy issue, and Sphinx demonstrate meaningful project- and task-specific brittleness.

Diagnosed failure modes.

1. **Incomplete or incorrect implementation on difficult issues** — The system has recorded unresolved outcomes on three probes, indicating that correct issue understanding does not consistently translate into a working patch. *Evidence:* `mwaskom__seaborn-3187`, `sympy__sympy-15875`, `sphinx-doc__sphinx-9602`.
2. **Insufficient robustness across project-specific edge cases** — Failures span Seaborn, SymPy, and Sphinx, suggesting difficulty adapting implementation and validation to differing library conventions and edge cases. *Evidence:* `mwaskom__seaborn__seaborn-3187`, `sympy__sympy-15875`, `sphinx-doc__sphinx-9602`.
3. **Uneven cross-repository generalization** — Although the system resolves representative scikit-learn, SymPy, and Django tasks, it fails other repositories in the same benchmark, showing non-uniform performance across codebases. *Evidence:* `mwaskom__seaborn-3187`, `sympy__sympy-15875`, `sphinx-doc__sphinx-9602`.

Top three why-records (verbatim).

1. **prompt-04** (Invariant preservation and differential inspection; 0.8900, $\Delta +0.0425$, applied): Identifying and checking invariants across types, ordering, mutation, exceptions, serialization, and performance gives the strongest protection against subtle regressions and incomplete fixes. It improves reasoning, execution, and output fidelity, though it may add analysis burden on simple issues.
2. **prompt-20** (Invariant-first implementation planning; 0.8900, $\Delta +0.0425$, applied): Invariant-first planning materially improves reasoning quality and caller coverage, with a modest execution benefit. It can become overly analytical and still depends on correctly identifying the invariant, so difficult edge cases improve only slightly.

3. **prompt-24** (Tooling- and compatibility-aware completion protocol; 0.8900, $\Delta +0.0425$, applied): The completion protocol improves practical execution by enforcing repository tooling, discovery, formatting, compatibility, and a disciplined final report. It provides the best output-integrity gain and a modest execution gain, but is broad enough that it adds less problem-specific reasoning than an invariant or edge-case-focused approach.

Table 14: All 60 scored versions of live-SWE-agent.

id	phase	technique	score	Δ	decision
struct-01	struct	Add a mandatory hypothesis-to-validation loop: after locating the...	0.8775	+0.0300	proposed
struct-02	struct	Introduce a two-pass implementation architecture in which the fir...	0.8775	+0.0300	proposed
struct-03	struct	Add a repository-convention discovery phase that inspects package...	0.8650	+0.0175	proposed
struct-04	struct	Build an edge-case matrix from the issue description and affected...	0.8600	+0.0125	proposed
struct-05	struct	Add an adaptive test-selection layer that maps changed files and ...	0.8675	+0.0200	proposed
struct-06	struct	Use a repository bootstrap and capability profile before solving;...	0.8650	+0.0175	proposed
struct-07	struct	Add a staged implementation controller that requires the agent to...	0.8725	+0.0250	proposed
struct-08	struct	Implement a mandatory reproduce-patch-verify loop in the agent co...	0.8800	+0.0325	proposed
struct-09	struct	Add a project-aware edge-case planning module that mines neighbor...	0.8700	+0.0225	proposed
struct-10	struct	Introduce repository adapters selected during initialization. Eac...	0.8725	+0.0250	proposed
struct-11	struct	Add an independent post-patch reviewer stage with a separate cont...	0.8825	+0.0350	proposed
struct-12	struct	Build a failure-recovery state machine that records failed hypoth...	0.8650	+0.0175	proposed
struct-13	struct	Add a repository reconnaissance gate that requires the agent to i...	0.8625	+0.0150	proposed
struct-14	struct	Introduce a hypothesis-and-reproduction phase gate: before modify...	0.8750	+0.0275	proposed
struct-15	struct	Add an implementation loop with explicit checkpoints: make the sm...	0.8850	+0.0375	proposed
struct-16	struct	Add a project-convention discovery step that searches for analogo...	0.8700	+0.0225	proposed
struct-17	struct	Require an edge-case matrix before finalization covering normal i...	0.8650	+0.0175	proposed
struct-18	struct	Add a finalization gate that inspects the complete diff, checks t...	0.8800	+0.0325	proposed
struct-19	struct	Add a requirement-to-implementation traceability gate that conver...	0.8750	+0.0275	proposed
struct-20	struct	Introduce a staged reproduce-patch-test-debug controller with man...	0.8850	+0.0375	proposed
struct-21	struct	Add an edge-case matrix generator that derives cases from public ...	0.8500	+0.0025	proposed
struct-22	struct	Create a repository-convention discovery phase that identifies th...	0.8725	+0.0250	proposed
struct-23	struct	Add an implementation-review pass that compares the proposed chan...	0.8600	+0.0125	proposed
struct-24	struct	Implement a failure-recovery loop that classifies test failures a...	0.8775	+0.0300	proposed
struct-25	struct	Add a mandatory implementation-validation gate that requires the ...	0.8600	+0.0125	proposed
struct-26	struct	Introduce a two-pass patch protocol: first require an explicit be...	0.8725	+0.0250	proposed

id	phase	technique	score	Δ	decision
struct-27	struct	Add a repository-adaptation phase that detects the project's test...	0.8725	+0.0250	proposed
struct-28	struct	Require an edge-case test matrix derived from the issue and nearb...	0.8675	+0.0200	proposed
struct-29	struct	Implement a cross-repository strategy selector that classifies th...	0.8550	+0.0075	proposed
struct-30	struct	Add a failure-escalation controller that classifies failed tests ...	0.8700	+0.0225	proposed
prompt-01	prompt	Acceptance-criteria decomposition before editing	0.8750	+0.0275	applied
prompt-02	prompt	Hypothesis-driven implementation and recovery loop	0.8775	+0.0300	applied
prompt-03	prompt	Boundary-and-compatibility test matrix	0.8700	+0.0225	applied
prompt-04	prompt	Invariant preservation and differential inspection	0.8900	+0.0425	applied
prompt-05	prompt	Repository-native reconnaissance	0.8825	+0.0350	applied
prompt-06	prompt	Tooling-adaptive validation plan	0.8775	+0.0300	applied
prompt-07	prompt	Issue-contract decomposition	0.8775	+0.0300	applied
prompt-08	prompt	Repository-native edge-case matrix	0.8625	+0.0150	proposed
prompt-09	prompt	Cross-project convention discovery	0.8650	+0.0175	proposed
prompt-10	prompt	Failing-test-first localization	0.8650	+0.0175	proposed
prompt-11	prompt	Invariant and integration audit	0.8825	+0.0350	proposed
prompt-12	prompt	Patch-completeness gate	0.8725	+0.0250	proposed
prompt-13	prompt	Explicit contract-to-code trace	0.8775	+0.0300	applied
prompt-14	prompt	Hypothesis-driven iterative repair	0.8800	+0.0325	applied
prompt-15	prompt	Adversarial behavioral matrix	0.8750	+0.0275	applied
prompt-16	prompt	Repository-native convention discovery	0.8775	+0.0300	applied
prompt-17	prompt	Cross-repository completion gate	0.8725	+0.0250	applied
prompt-18	prompt	Semantic diff and regression audit	0.8800	+0.0325	proposed
prompt-19	prompt	Acceptance-contract and failing-reproduction gate	0.8850	+0.0375	applied
prompt-20	prompt	Invariant-first implementation planning	0.8900	+0.0425	applied
prompt-21	prompt	Repository-specific edge-case test matrix	0.8650	+0.0175	applied
prompt-22	prompt	Behavioral compatibility and neighboring-pattern checkpoint	0.8350	-0.0125	applied
prompt-23	prompt	Cross-repository reconnaissance before solution commitment	0.8875	+0.0400	applied
prompt-24	prompt	Tooling- and compatibility-aware completion protocol	0.8900	+0.0425	applied
prompt-25	prompt	Contract-first issue decomposition	0.8750	+0.0275	applied
prompt-26	prompt	Evidence-driven repair loop	0.8850	+0.0375	applied
prompt-27	prompt	Explicit adversarial test matrix	0.8625	+0.0150	applied
prompt-28	prompt	Invariant and interaction checking	0.8650	+0.0175	applied
prompt-29	prompt	Repository-convention reconnaissance	0.8675	+0.0200	applied
prompt-30	prompt	Failure triage and completion gate	0.8725	+0.0250	applied

D.2 Sonar Foundation Agent + Claude 4.5 Opus

Table 15: Record for Sonar Foundation Agent (rank 2).

Submission	20251205_sonar-foundation-agent_claude-opus-4-5 (verified)
Official resolve rate	79.20% (396 instances)
Repository	AutoCodeRoverSG/sonar-foundation-agent @ 394c58819e
Recorded trajectories readable	0 (archive assets are credentialed)
Prompt-surface files found	0
Runnability	no source code in linked repo
Baseline dims (TA/RS/EA/OI)	0.82 / 0.79 / 0.78 / 0.84
Trajectory	0.8075 $\rightarrow$ 0.8600 $\rightarrow$ 0.8625 (all PROXY)
Gate verdict	accept: mean not below reference; 4W/0L paired (sign p=0.125)
MDE ₈₀ at $n=4$	0.297; shipped $\Delta = +0.0550$ (below MDE: yes)
Paraphrase noise floor	+0.0000 (exceeds_noise_floor)
Blinded A/B ($k=5$)	— (no applied patch to compare)

Judge rationale (baseline). The 79.2% verified resolve rate supports solid but clearly imperfect scores, with execution accuracy weakest because two of the four fixed probes were unresolved. The probe outcomes show strong success on two instances but near-failure on two others, and the absence of a prompt surface limits any stronger claim about the specific internal causes.

Diagnosed failure modes.

1. **End-to-end issue non-resolution** — The archived outcomes show complete failure on two of four fixed probes, indicating that the agent does not reliably turn analysis into an accepted repository resolution. *Evidence:* `pydata__xarray-6599`, `sympy__sympy-14248`.
2. **Cross-repository robustness gap** — The system succeeds on the scikit-learn and one SymPy probe but fails on the xarray and another SymPy probe, showing inconsistent handling across repositories and issue instances. *Evidence:* `scikit-learn__scikit-learn-15100`, `sympy__sympy-19954`, `pydata__xarray-6599`, `sympy__sympy-14248`.
3. **Instance-level reliability variance** — The mixed resolved and unresolved archive outcomes, together with a 79.2% overall resolve rate, indicate meaningful variance rather than uniformly dependable behavior. *Evidence:* `scikit-learn__scikit-learn-15100`, `sympy__sympy-19954`, `pydata__xarray-6599`, `sympy__sympy-14248`.

Top three why-records (verbatim).

1. **prompt-07** (Explicit issue-to-acceptance loop; 0.8625, $\Delta +0.0550$, proposed_not_applied): The explicit acceptance loop directly addresses stopping after diagnosis and should improve completion and verification. It raises target alignment and execution accuracy substantially, with modest gains on the previously failing probes, but it does not by itself solve repository-specific investigation.
2. **prompt-26** (Evidence-gated implementation loop; 0.8625, $\Delta +0.0550$, proposed_not_applied): This is the strongest end-to-end intervention: reproduction, focused validation, broader testing, and revision directly address non-resolution and prevent unverified hypotheses from surviving. It still cannot eliminate repository-specific ambiguity, so gains on the hardest probes remain bounded.
3. **struct-23** (Add a semantic diff and regression-review gate after implementation: ...; 0.8600, $\Delta +0.0525$, proposed_not_applied): The semantic diff and regression-review gate strongly improves alignment with requested behavior, compatibility awareness, and final patch integrity. It is especially effective after a basically correct implementation, producing the best scores on the easy probes, but it is downstream of diagnosis and therefore gives only limited benefit on the unresolved cross-repository failures.

Table 16: All 60 scored versions of Sonar Foundation Agent.

id	phase	technique	score	Δ	decision
struct-01	struct	Add a mandatory post-patch validation loop that runs the reposito...	0.8500	+0.0425	proposed
struct-02	struct	Introduce repository-adaptive issue triage that identifies the pr...	0.8375	+0.0300	proposed
struct-03	struct	Use a staged execution controller with explicit checkpoints for i...	0.8350	+0.0275	proposed
struct-04	struct	Add a failure-aware test selection subsystem that starts with nar...	0.8525	+0.0450	proposed
struct-05	struct	Implement a repository-specific patch synthesis policy that prese...	0.8325	+0.0250	proposed
struct-06	struct	Add a final acceptance gate that verifies the requested behavior,...	0.8225	+0.0150	proposed
struct-07	struct	Add a verifier-gated repair loop that runs the repository’s targe...	0.8400	+0.0325	proposed
struct-08	struct	Introduce an issue-to-acceptance-contract stage that extracts req...	0.8375	+0.0300	proposed
struct-09	struct	Add repository-specific adapter selection before planning, using ...	0.8200	+0.0125	proposed
struct-10	struct	Run two independent solution analyses for ambiguous issues, then ...	0.8300	+0.0225	proposed
struct-11	struct	Create a checkpointed execution state machine with explicit phase...	0.8275	+0.0200	proposed
struct-12	struct	Add a final patch-integrity gate that rejects submission unless t...	0.8375	+0.0300	proposed
struct-13	struct	Introduce a mandatory plan-patch-test state machine that prevents...	0.8450	+0.0375	proposed
struct-14	struct	Add a bounded test-failure recovery loop that classifies failures...	0.8375	+0.0300	proposed
struct-15	struct	Build a repository fingerprinting stage that detects project layo...	0.8450	+0.0375	proposed
struct-16	struct	Add a repository-specific validation-command selector that derive...	0.8350	+0.0275	proposed
struct-17	struct	Use parallel patch hypotheses with an evidence ledger, retaining ...	0.8300	+0.0225	proposed
struct-18	struct	Introduce a final acceptance gate that checks changed-file presen...	0.8525	+0.0450	proposed
struct-19	struct	Add an explicit issue-to-acceptance-criteria planning gate before...	0.8425	+0.0350	proposed
struct-20	struct	Introduce repository-adaptive reconnaissance: before proposing a ...	0.8200	+0.0125	proposed
struct-21	struct	Use a staged execution loop with narrow checks after every change...	0.8400	+0.0325	proposed
struct-22	struct	Add a failure-recovery controller that classifies test, import, p...	0.8550	+0.0475	proposed
struct-23	struct	Add a semantic diff and regression-review gate after implementati...	0.8600	+0.0525	proposed
struct-24	struct	Use a confidence- and complexity-aware strategy selector: route s...	0.8325	+0.0250	proposed
struct-25	struct	Add a mandatory verify-and-repair controller that runs the releva...	0.8475	+0.0400	proposed
struct-26	struct	Introduce repository adapters selected from repository metadata, ...	0.8200	+0.0125	proposed
struct-27	struct	Add a hypothesis manager that maintains multiple issue interpreta...	0.8225	+0.0150	proposed
struct-28	struct	Create a submission gate that checks diff presence, modified-file...	0.8375	+0.0300	proposed
struct-29	struct	Add a standardized failure-triage loop that maps each failing tes...	0.8275	+0.0200	proposed
struct-30	struct	Implement uncertainty-triggered escalation: when localization con...	0.8325	+0.0250	proposed
prompt-01	prompt	Explicit issue-to-patch execution contract	0.8400	+0.0325	proposed
prompt-02	prompt	Hypothesis-driven debugging with falsification	0.8225	+0.0150	proposed

id	phase	technique	score	Δ	decision
prompt-03	prompt	Acceptance-test-first repair planning	0.8325	+0.0250	proposed
prompt-04	prompt	Focused validation escalation ladder	0.8450	+0.0375	proposed
prompt-05	prompt	Repository-local convention and compatibility scan	0.8350	+0.0275	proposed
prompt-06	prompt	Final patch-integrity gate	0.8450	+0.0375	proposed
prompt-07	prompt	Explicit issue-to-acceptance loop	0.8625	+0.0550	proposed
prompt-08	prompt	Repository-adaptive investigation protocol	0.8425	+0.0350	proposed
prompt-09	prompt	Test-first behavioral contract	0.8425	+0.0350	proposed
prompt-10	prompt	Failure-aware validation ladder	0.8500	+0.0425	proposed
prompt-11	prompt	Hypothesis branching and falsification	0.8300	+0.0225	proposed
prompt-12	prompt	Pre-submit patch audit	0.8450	+0.0375	proposed
prompt-13	prompt	Explicit issue-to-acceptance-contract prompting	0.8400	+0.0325	proposed
prompt-14	prompt	Repository-aware investigation protocol	0.8300	+0.0225	proposed
prompt-15	prompt	Hypothesis-driven debugging with falsification	0.8200	+0.0125	proposed
prompt-16	prompt	Mandatory test-first and regression-test completion gate	0.8400	+0.0325	proposed
prompt-17	prompt	Layered validation and failure-recovery loop	0.8375	+0.0300	proposed
prompt-18	prompt	Patch-scope and deliverable integrity checklist	0.8325	+0.0250	proposed
prompt-19	prompt	Explicit issue-to-acceptance-contract checklist	0.8400	+0.0325	proposed
prompt-20	prompt	Repository-adaptive investigation protocol	0.8300	+0.0225	proposed
prompt-21	prompt	Hypothesis-driven differential debugging	0.8425	+0.0350	proposed
prompt-22	prompt	Mandatory validation ladder with failure recovery	0.8600	+0.0525	proposed
prompt-23	prompt	Compatibility and edge-case matrix	0.8125	+0.0050	proposed
prompt-24	prompt	Completion-gate output contract	0.8150	+0.0075	proposed
prompt-25	prompt	Issue-to-acceptance-contract	0.8450	+0.0375	proposed
prompt-26	prompt	Evidence-gated implementation loop	0.8625	+0.0550	proposed
prompt-27	prompt	Repository-adaptive investigation protocol	0.8400	+0.0325	proposed
prompt-28	prompt	Ambiguity and edge-case matrix	0.8125	+0.0050	proposed
prompt-29	prompt	Confidence-triggered escalation	0.8125	+0.0050	proposed
prompt-30	prompt	Submission-readiness gate	0.8325	+0.0250	proposed

D.3 TRAE + Doubao-Seed-Code

Table 17: Record for TRAE (rank 3).

Submission	20250928_trae_doubao_seed_code (verified)
Official resolve rate	78.80% (394 instances)
Repository	bytedance/trae-agent @ e839e559ac
Recorded trajectories readable	0 (archive assets are credentialed)
Prompt-surface files found	6
Runnability	cli entry point plausible; execution still requires SWE-bench task in...
Baseline dims (TA/RS/EA/OI)	0.86 / 0.82 / 0.79 / 0.87
Trajectory	0.8350 $\rightarrow$ 0.8850 $\rightarrow$ 0.8900 (all PROXY)
Gate verdict	accept: mean not below reference; 4W/0L paired (sign p=0.125)
MDE ₈₀ at $n=4$	0.297; shipped $\Delta = +0.0550$ (below MDE: yes)
Paraphrase noise floor	+0.0050 (exceeds_noise_floor)
Blinded A/B ($k=5$)	patched 4 / baseline 1 / tie 0

Judge rationale (baseline). The recorded outcomes show strong but imperfect capability: two of four fixed probes resolved, while two remained unresolved, consistent with the system’s broader 78.8% verified resolve rate rather than near-perfect reliability. Its modular agent, tool, and patch-selection surfaces support good alignment and output integrity, but the unresolved Django and Astropy cases most directly lower execution accuracy and reasoning soundness.

Diagnosed failure modes.

1. **Complex repository-specific diagnosis and implementation** — The agent can complete standard fixes but fails on at least some issues requiring deeper localization, design understanding, or coordinated changes across an unfamiliar codebase. *Evidence:* `django__django-16938`, `astropy__astropy-13033`.
2. **Insufficient validation and iterative debugging** — Tool-driven editing and testing do not consistently converge when the initial patch is incomplete or when failures require multiple debugging iterations. *Evidence:* `django__django-16938`, `astropy__astropy-13033`.
3. **Edge-case and regression handling** — The unresolved probes indicate difficulty preserving behavior while addressing nuanced framework or scientific-library edge cases, despite successful outcomes on two other instances. *Evidence:* `django__django-16938`, `astropy__astropy-13033`.

Top three why-records (verbatim).

1. **prompt-03** (Mandatory test-driven convergence loop; 0.8900, $\Delta +0.0550$, applied): The mandatory focused-test and broader-test loop most directly repairs insufficient validation and iterative debugging, producing the strongest execution improvement. It still cannot compensate for an incorrect initial repository diagnosis, and mandatory testing can add cost or expose environment noise.
2. **prompt-06** (Compatibility-preserving differential review; 0.8875, $\Delta +0.0525$, applied): Differential review gives the strongest protection against accidental changes in exceptions, ordering, mutation, precision, serialization, and alternate implementations, substantially improving output integrity and regression handling. It is more reliable than the raw boundary matrix but remains dependent on having identified the correct old and new paths, so it trails the mandatory convergence loop on execution.
3. **prompt-23** (Boundary-condition matrix; 0.8875, $\Delta +0.0525$, applied): Combining evidence-gated localization with an explicit boundary matrix addresses both diagnosis and edge-case handling, making it the broadest improvement on the difficult probes. The matrix can be costly and may encourage excessive cases, so execution gains remain moderate rather than maximal.

Table 18: All 60 scored versions of TRAE.

id	phase	technique	score	Δ	decision
struct-01	struct	Inject a mandatory repository-diagnosis phase into the agent's sy...	0.8675	+0.0325	applied
struct-02	struct	Inject a bounded reproduce-test-debug loop into the agent's syste...	0.8625	+0.0275	applied
struct-03	struct	Add an explicit compatibility and regression checklist to the age...	0.8625	+0.0275	applied
struct-04	struct	Make patch selection require behavioral evidence rather than surf...	0.8725	+0.0375	applied
struct-05	struct	Normalize candidate patch fields and provide safe fallbacks at th...	0.8300	-0.0050	applied
struct-06	struct	Bias the selector toward minimal, behavior-preserving patches wit...	0.8850	+0.0500	proposed
struct-07	struct	Add a mandatory repository-diagnosis phase before editing, requir...	0.8800	+0.0450	proposed
struct-08	struct	Turn the existing edge-case reminder into an explicit behavior ma...	0.8675	+0.0325	applied
struct-09	struct	Add a structured hypothesis-test-revise loop that requires the ag...	0.8800	+0.0450	applied
struct-10	struct	Strengthen patch selection with an evidence hierarchy that requir...	0.8675	+0.0325	applied
struct-11	struct	Normalize candidate patch metadata at construction time and deriv...	0.8550	+0.0200	applied
struct-12	struct	Make success metadata conservative by requiring both the reported...	0.8575	+0.0225	applied
struct-13	struct	Reconnaissance-first execution contract	0.8625	+0.0275	proposed
struct-14	struct	Persistent hypothesis-and-evidence ledger	0.8650	+0.0300	applied
struct-15	struct	Failure-classification and reproducer escalation	0.8225	-0.0125	applied
struct-16	struct	Metamorphic edge-case expansion	0.8275	-0.0075	proposed
struct-17	struct	Compatibility-focused adversarial review	0.8625	+0.0275	applied
struct-18	struct	Issue-to-diff semantic traceability gate	0.8775	+0.0425	applied
struct-19	struct	Repository reconnaissance gate	0.8675	+0.0325	proposed
struct-20	struct	Behavior-contract planning	0.8750	+0.0400	proposed
struct-21	struct	Failure-taxonomy test ladder	0.8625	+0.0275	proposed
struct-22	struct	Traceback-driven repair loop	0.8750	+0.0400	proposed
struct-23	struct	Explicit edge-case matrix enforcement	0.8400	+0.0050	proposed
struct-24	struct	Compatibility and regression audit	0.8650	+0.0300	proposed
struct-25	struct	Add a mandatory repository-localization phase before editing	0.8600	+0.0250	proposed
struct-26	struct	Require an explicit hypothesis-and-retest debugging loop	0.8675	+0.0325	proposed
struct-27	struct	Add an adversarial compatibility checklist to the implementation ...	0.8550	+0.0200	proposed
struct-28	struct	Make patch-selection judgments evidence-first and failure-type-aw...	0.8425	+0.0075	applied
struct-29	struct	Normalize candidate patch text before judging it	0.8500	+0.0150	applied
struct-30	struct	Strengthen the success gate against placeholder or non-patch outp...	0.8575	+0.0225	applied
prompt-01	prompt	Repository reconnaissance and dependency tracing	0.8700	+0.0350	applied
prompt-02	prompt	Contract-first implementation planning	0.8725	+0.0375	proposed
prompt-03	prompt	Mandatory test-driven convergence loop	0.8900	+0.0550	applied
prompt-04	prompt	Failure-classification and escalation protocol	0.8550	+0.0200	proposed
prompt-05	prompt	Expanded boundary and sentinel behavior matrix	0.8425	+0.0075	applied
prompt-06	prompt	Compatibility-preserving differential review	0.8875	+0.0525	applied
prompt-07	prompt	Evidence-gated repository localization	0.8725	+0.0375	applied
prompt-08	prompt	Invariant and analogue triangulation	0.8825	+0.0475	proposed
prompt-09	prompt	Reproduction-first validation loop	0.8625	+0.0275	proposed
prompt-10	prompt	Failure triage and bounded iteration	0.8700	+0.0350	proposed

id	phase	technique	score	Δ	decision
prompt-11	prompt	Compatibility matrix with explicit omission semantics	0.8475	+0.0125	proposed
prompt-12	prompt	Metamorphic and negative-case regression checks	0.8475	+0.0125	proposed
prompt-13	prompt	Evidence-chain localization with explicit hypothesis testing	0.8650	+0.0300	applied
prompt-14	prompt	Change-surface and compatibility mapping	0.8500	+0.0150	applied
prompt-15	prompt	Failing-test-first execution loop	0.8575	+0.0225	applied
prompt-16	prompt	Layered validation and regression triage	0.8675	+0.0325	applied
prompt-17	prompt	Contract-preserving edge-case matrix	0.8675	+0.0325	applied
prompt-18	prompt	Invariant and differential regression checking	0.8725	+0.0375	applied
prompt-19	prompt	Evidence-traced call-chain localization	0.8825	+0.0475	applied
prompt-20	prompt	Change-surface and contract inventory	0.8625	+0.0275	applied
prompt-21	prompt	Failure-directed validation loop	0.8675	+0.0325	proposed
prompt-22	prompt	Negative-space regression testing	0.8600	+0.0250	applied
prompt-23	prompt	Boundary-condition matrix	0.8875	+0.0525	applied
prompt-24	prompt	Invariant-preserving patch design	0.8775	+0.0425	applied
prompt-25	prompt	Evidence-first call-chain localization	0.8650	+0.0300	applied
prompt-26	prompt	Hypothesis competition and minimal design selection	0.8600	+0.0250	applied
prompt-27	prompt	Layered test-gated iteration	0.8650	+0.0300	applied
prompt-28	prompt	Differential regression validation	0.8700	+0.0350	applied
prompt-29	prompt	Invariant-driven boundary testing	0.8675	+0.0325	applied
prompt-30	prompt	Compatibility-preserving behavior partitioning	0.8600	+0.0250	applied

D.4 ACoder

Table 19: Record for ACoder (rank 4).

Submission	20250819_ACoder (verified)
Official resolve rate	76.40% (382 instances)
Repository	ACoder-AI/ACoder @ 63325725b6
Recorded trajectories readable	0 (archive assets are credentialed)
Prompt-surface files found	0
Runnability	no source code in linked repo
Baseline dims (TA/RS/EA/OI)	0.78 / 0.77 / 0.75 / 0.80
Trajectory	0.7750 $\rightarrow$ 0.8450 $\rightarrow$ 0.8450 (all PROXY)
Gate verdict	accept: mean not below reference; 4W/0L paired (sign p=0.125)
MDE ₈₀ at $n=4$	0.297; shipped $\Delta = +0.0700$ (below MDE: yes)
Paraphrase noise floor	+0.0050 (exceeds_noise_floor)
Blinded A/B ($k=5$)	— (no applied patch to compare)

Judge rationale (baseline). ACoder’s officially resolved rate is 76.4%, so its dimensions should reflect meaningful capability with nontrivial failure risk rather than near-perfect performance. The fixed probes show successful handling of `django__django-14631` and `scikit-learn__scikit-learn-15100`, but failure on `pylint-dev__pylint-7277` and `scikit-learn__scikit-learn-13124`; the repository also exposes no prompt surface for stronger evidence about alignment or output behavior.

Diagnosed failure modes.

1. **Incomplete task resolution** — The system failed to produce a resolved outcome on two fixed repository tasks, indicating substantial tail risk despite the aggregate success rate. *Evidence:* `pylint-dev__pylint-7277`, `scikit-learn__scikit-learn-13124`.
2. **Insufficient repository-specific debugging or implementation accuracy** — Failures on both a Pylint task and a scikit-learn task suggest difficulty converting repository analysis into a correct, project-specific change. *Evidence:* `pylint-dev__pylint-7277`, `scikit-learn__scikit-learn-13124`.
3. **Unreliable end-to-end verification and completion** — The recorded unresolved outcomes indicate that exploration and reasoning do not consistently culminate in a validated patch, even though two other probes were resolved. *Evidence:* `pylint-dev__pylint-7277`, `scikit-learn__scikit-learn-13124`.

Top three why-records (verbatim).

1. **struct-18** (Use a two-pass patch review architecture: have the implementation sub...; 0.8450, $\Delta +0.0700$, proposed_not_applied): Independent review directly targets subtle project-specific implementation mistakes and edge cases, giving this the strongest reasoning and execution improvement among the candidates. It remains dependent on the first subagent producing a viable patch and can add latency or reviewer disagreement, so the difficult probes improve meaningfully but remain far from reliable.
2. **struct-12** (Implement a hard completion gate that accepts a task only when a non-...; 0.8425, $\Delta +0.0675$, proposed_not_applied): The hard completion gate most directly prevents unresolved, empty, or unverifiable submissions and gives the strongest improvement to end-to-end integrity. It does not itself produce a correct patch, and strict gating can reject partially successful work when broader tests are unavailable.
3. **struct-05** (Add a completion gate with a structured checklist requiring: changed ...; 0.8400, $\Delta +0.0650$, proposed_not_applied): The completion gate is tightly aligned with incomplete task resolution and gives the strongest improvement to patch presence, test evidence, and final-response accuracy. It is primarily a rejection mechanism, so it raises weak behavioral outcomes only modestly without stronger debugging.

Table 20: All 60 scored versions of ACoder.

id	phase	technique	score	Δ	decision
struct-01	struct	Add a mandatory issue-to-code execution plan before editing: extr...	0.8150	+0.0400	proposed
struct-02	struct	Introduce a repository-specific investigation stage that delegate...	0.8250	+0.0500	proposed
struct-03	struct	Add a hypothesis-driven debugging loop: record the current failur...	0.8150	+0.0400	proposed
struct-04	struct	Create a layered verification pipeline that runs focused regressi...	0.8225	+0.0475	proposed
struct-05	struct	Add a completion gate with a structured check-list requiring: chan...	0.8400	+0.0650	proposed
struct-06	struct	Implement a bounded recovery strategy after verification failure:...	0.8400	+0.0650	proposed
struct-07	struct	Add a persistent task-state controller with explicit phases for i...	0.8100	+0.0350	proposed
struct-08	struct	Introduce a fallback planner that generates at least two implemen...	0.7700	-0.0050	proposed
struct-09	struct	Add a repository-specific debugging stage that builds a compact i...	0.8200	+0.0450	proposed
struct-10	struct	Require a test-first reproduction artifact: identify an existing ...	0.8275	+0.0525	proposed
struct-11	struct	Add an independent verification subagent that receives only the i...	0.8175	+0.0425	proposed
struct-12	struct	Implement a hard completion gate that accepts a task only when a ...	0.8425	+0.0675	proposed
struct-13	struct	Add a mandatory hypothesis-to-test execution loop: require the ag...	0.8200	+0.0450	proposed
struct-14	struct	Introduce a repository-specific implementation-planning stage tha...	0.8350	+0.0600	proposed
struct-15	struct	Add an adaptive recovery controller that detects stalled explorat...	0.7900	+0.0150	proposed
struct-16	struct	Build a semantic repository-navigation index that links issue ter...	0.8225	+0.0475	proposed
struct-17	struct	Add a completion gate that refuses to finalize unless the working...	0.8325	+0.0575	proposed
struct-18	struct	Use a two-pass patch review architecture: have the implementation...	0.8450	+0.0700	proposed
struct-19	struct	Mandatory hypothesis-to-test execution loop	0.8175	+0.0425	proposed
struct-20	struct	Repository-specific evidence ledger	0.8275	+0.0525	proposed
struct-21	struct	Independent patch-review subagent before finalization	0.8350	+0.0600	proposed
struct-22	struct	Failure-directed debugging escalation	0.8125	+0.0375	proposed
struct-23	struct	Completion contract with machine-checkable gates	0.8300	+0.0550	proposed
struct-24	struct	Adaptive test selection and regression expansion	0.8300	+0.0550	proposed
struct-25	struct	Add a mandatory task-closure state machine that prevents terminat...	0.8300	+0.0550	proposed
struct-26	struct	Introduce a bounded recovery controller that classifies an unsucc...	0.7975	+0.0225	proposed
struct-27	struct	Add a repository-specific change-plan stage that extracts relevan...	0.8200	+0.0450	proposed
struct-28	struct	Use independent implementation and review subagents with a disagr...	0.8350	+0.0600	proposed
struct-29	struct	Build a validation ladder that runs focused regression tests firs...	0.8250	+0.0500	proposed
struct-30	struct	Add a final evidence-based completion gate that checks the diff, ...	0.8250	+0.0500	proposed
prompt-01	prompt	Repository-specific hypothesis and evidence loop	0.8000	+0.0250	proposed
prompt-02	prompt	Issue-to-acceptance-criteria decomposition	0.7975	+0.0225	proposed
prompt-03	prompt	Mandatory test-driven patch validation	0.8100	+0.0350	proposed
prompt-04	prompt	Minimal-diff and regression-safety constraint	0.7775	+0.0025	proposed
prompt-05	prompt	Failure-directed debugging escalation	0.8150	+0.0400	proposed

id	phase	technique	score	Δ	decision
prompt-06	prompt	Completion gate with independent review pass	0.8325	+0.0575	proposed
prompt-07	prompt	Explicit completion checklist and stop-condition gate	0.7700	-0.0050	proposed
prompt-08	prompt	Repository-specific evidence ledger	0.8075	+0.0325	proposed
prompt-09	prompt	Test-first reproduction and regression design	0.8175	+0.0425	proposed
prompt-10	prompt	Layered validation escalation	0.8225	+0.0475	proposed
prompt-11	prompt	Adversarial patch review before submission	0.8175	+0.0425	proposed
prompt-12	prompt	Minimal-diff and scope-control contract	0.8100	+0.0350	proposed
prompt-13	prompt	Explicit issue-contract extraction	0.8075	+0.0325	proposed
prompt-14	prompt	Requirement-to-code evidence matrix	0.8150	+0.0400	proposed
prompt-15	prompt	Competing-hypothesis debugging	0.7750	+0.0000	proposed
prompt-16	prompt	Minimal-change implementation contract	0.8125	+0.0375	proposed
prompt-17	prompt	Layered verification gate	0.8200	+0.0450	proposed
prompt-18	prompt	Completion-and-artifact gate	0.8150	+0.0400	proposed
prompt-19	prompt	Explicit acceptance-criteria extraction and checklist tracking	0.8150	+0.0400	proposed
prompt-20	prompt	Repository-local evidence chain before editing	0.8250	+0.0500	proposed
prompt-21	prompt	Minimal-patch and invariant-preservation constraint	0.8075	+0.0325	proposed
prompt-22	prompt	Test-first regression reproduction	0.8225	+0.0475	proposed
prompt-23	prompt	Layered verification gate with targeted and regression tests	0.8275	+0.0525	proposed
prompt-24	prompt	Failure recovery and bounded alternative-path protocol	0.8225	+0.0475	proposed
prompt-25	prompt	Explicit acceptance-criteria decomposition	0.8050	+0.0300	proposed
prompt-26	prompt	Repository-convention-first investigation protocol	0.8150	+0.0400	proposed
prompt-27	prompt	Hypothesis-driven debugging with falsification	0.8150	+0.0400	proposed
prompt-28	prompt	Minimal patch and regression-boundary rule	0.7975	+0.0225	proposed
prompt-29	prompt	Mandatory layered verification gate	0.8225	+0.0475	proposed
prompt-30	prompt	Final-state integrity checklist	0.7900	+0.0150	proposed

D.5 JoyCode + Claude 4 Sonnet + GPT-4.1

Table 21: Record for JoyCode (rank 5).

Submission	20250915_JoyCode (verified)
Official resolve rate	74.60% (373 instances)
Repository	jd-opensource/joycode-agent @ 1bace2ab9f
Recorded trajectories readable	0 (archive assets are credentialed)
Prompt-surface files found	12
Runnability	docker required
Baseline dims (TA/RS/EA/OI)	0.86 / 0.81 / 0.76 / 0.88
Trajectory	0.8275 $\rightarrow$ 0.8675 $\rightarrow$ 0.8850 (all PROXY)
Gate verdict	accept: mean not below reference; 4W/0L paired (sign p=0.125)
MDE ₈₀ at $n=4$	0.297; shipped $\Delta = +0.0575$ (below MDE: yes)
Paraphrase noise floor	+0.0075 (exceeds_noise_floor)
Blinded A/B ($k=5$)	patched 5 / baseline 0 / tie 0 (post-repair)

Judge rationale (baseline). The system demonstrates strong but incomplete alignment and output handling, with a verified 74.6% resolution rate and two successful probes. Execution accuracy is weakest because two of the four fixed probes are explicitly unresolved, so the system should not receive near-perfect TEI scores.

Diagnosed failure modes.

1. **Failure to complete repository-specific fixes** — The system’s patch-generation and execution pipeline did not resolve the Django 14170 and Sphinx 9229 tasks, showing residual failures despite the reported overall 74.6% resolution rate. *Evidence:* `django__django-14170`, `sphinx-doc__sphinx-9229`.
2. **Insufficient recovery from failed attempts** — Although the repository emphasizes intelligent retries, test generation, and failure attribution, the recorded unresolved outcomes indicate that these mechanisms do not reliably recover from difficult instances. *Evidence:* `django__django-14170`, `sphinx-doc__sphinx-9229`.
3. **Cross-project generalization gap** — The resolved Django and SymPy probes coexist with unresolved Django and Sphinx probes, indicating that success is not uniform across issue types and repository environments. *Evidence:* `django__django-13023`, `sympy__sympy-24443`, `django__django-14170`, `sphinx-doc__sphinx-9229`.

Top three why-records (verbatim).

1. **prompt-19** (Requirement-to-diff closure gate; 0.8850, $\Delta +0.0575$, applied): Strongly targets incomplete repository-specific fixes by forcing issue-to-code-path closure and explicit diff coverage. It improves difficult cases, though the checklist can become procedural without ensuring the chosen implementation is correct.
2. **prompt-22** (Verified-path execution contract; 0.8850, $\Delta +0.0575$, applied): Best execution-oriented change: it requires proof that verification imports the changed symbol and would fail before the patch, reducing false-positive test results. It improves both easy and hard probes, though it may miss behavior not covered by a narrowly chosen path.
3. **prompt-05** (Requirement-to-diff completeness gate; 0.8800, $\Delta +0.0525$, proposed_not_applied): A final requirement-to-diff gate directly catches omitted behaviors, missing tests, and unsupported acceptance criteria, producing the largest output-integrity improvement. It is primarily a verification mechanism, so it offers only modest execution gains and cannot repair an incorrect implementation by itself.

Table 22: All 60 scored versions of JoyCode.

id	phase	technique	score	Δ	decision
struct-01	struct	Repository-first retry prompting	0.8600	+0.0325	proposed
struct-02	struct	Dedicated failure-analysis model purpose	0.8300	+0.0025	applied
struct-03	struct	Evidence-based trajectory diagnosis	0.8575	+0.0300	applied
struct-04	struct	Failure-state and retry-plan schema	0.8675	+0.0400	applied
struct-05	struct	Strict artifact-verification instructions	0.8600	+0.0325	proposed
struct-06	struct	Configurable retry budget and backoff policy	0.8425	+0.0150	applied
struct-07	struct	Add purpose-specific model fallback routing with circuit breaking...	0.8400	+0.0125	applied
struct-08	struct	Make retry prompts failure-aware and stateful. Pass the failed co...	0.8600	+0.0325	applied
struct-09	struct	Introduce a repository-specific preflight phase before patch gene...	0.8625	+0.0350	applied
struct-10	struct	Add a typed failure taxonomy and require retry strategy selection...	0.8575	+0.0300	applied
struct-11	struct	Separate candidate patch validation from repository mutation by a...	0.8425	+0.0150	applied
struct-12	struct	Strengthen trajectory analysis from free-form extraction into evi...	0.8600	+0.0325	applied
struct-13	struct	Repository-grounded preflight contract	0.8575	+0.0300	applied
struct-14	struct	Typed failure taxonomy for recovery routing	0.8400	+0.0125	applied
struct-15	struct	Dedicated triage and independent verification model roles	0.8350	+0.0075	applied
struct-16	struct	Evidence-gated trajectory diagnosis and recovery planning	0.8650	+0.0375	applied
struct-17	struct	Explicit final-diff completeness audit	0.8500	+0.0225	applied
struct-18	struct	Bounded, structured retry policy state	0.8350	+0.0075	applied
struct-19	struct	Make repository preflight produce an executable repair contract	0.8600	+0.0325	applied
struct-20	struct	Introduce failure-specific retry policies instead of generic retr...	0.8650	+0.0375	applied
struct-21	struct	Force trajectory analysis to preserve exact state and prohibit re...	0.8500	+0.0225	applied
struct-22	struct	Add a repository-specific test-discovery gate before patch accept...	0.8500	+0.0225	applied
struct-23	struct	Require an explicit final-diff completeness audit against the iss...	0.8650	+0.0375	applied
struct-24	struct	Separate patch generation, test generation, and final audit conte...	0.8200	-0.0075	applied
struct-25	struct	Evidence-gated recovery state machine	0.8575	+0.0300	proposed
struct-26	struct	Requirement-to-diff completion gate	0.8650	+0.0375	proposed
struct-27	struct	Repository-specific execution adapter	0.8650	+0.0375	proposed
struct-28	struct	Progressive test ladder with failure attribution	0.8600	+0.0325	proposed
struct-29	struct	Semantic patch and API-surface verifier	0.8650	+0.0375	proposed
struct-30	struct	Clean-room branch-and-compare recovery	0.8625	+0.0350	proposed
prompt-01	prompt	Issue-contract decomposition before editing	0.8700	+0.0425	applied
prompt-02	prompt	Hypothesis-changing retry protocol	0.8550	+0.0275	proposed
prompt-03	prompt	Repository-specific convention matching	0.8750	+0.0475	applied
prompt-04	prompt	Evidence-anchored trajectory recovery	0.8675	+0.0400	proposed
prompt-05	prompt	Requirement-to-diff completeness gate	0.8800	+0.0525	proposed
prompt-06	prompt	Validated test-command escalation	0.8725	+0.0450	proposed
prompt-07	prompt	Requirement-to-symbol execution contract	0.8675	+0.0400	applied
prompt-08	prompt	Analogy-guided local implementation search	0.8650	+0.0375	applied
prompt-09	prompt	Stateful recovery with new-hypothesis requirements	0.8700	+0.0425	applied
prompt-10	prompt	Failure-class-specific recovery playbook	0.8575	+0.0300	applied
prompt-11	prompt	Requirement coverage gate before finalization	0.8750	+0.0475	applied
prompt-12	prompt	Repository-native validation ladder	0.8700	+0.0425	proposed
prompt-13	prompt	Call-path-first implementation contract	0.8700	+0.0425	applied
prompt-14	prompt	Behavioral boundary matrix	0.8650	+0.0375	applied
prompt-15	prompt	Failure-signature-driven recovery loop	0.8675	+0.0400	applied
prompt-16	prompt	State-preserving retry protocol	0.8400	+0.0125	applied

id	phase	technique	score	Δ	decision
prompt-17	prompt	Repository-adaptation preflight	0.8700	+0.0425	applied
prompt-18	prompt	Final-diff semantic audit	0.8675	+0.0400	applied
prompt-19	prompt	Requirement-to-diff closure gate	0.8850	+0.0575	applied
prompt-20	prompt	Failure-signature-directed recovery loop	0.8750	+0.0475	applied
prompt-21	prompt	Repository-local convention and analogue search	0.8700	+0.0425	applied
prompt-22	prompt	Verified-path execution contract	0.8850	+0.0575	applied
prompt-23	prompt	Negative-scope and compatibility audit	0.8625	+0.0350	applied
prompt-24	prompt	Two-phase validation with minimal-patch fall-back	0.8775	+0.0500	applied
prompt-25	prompt	Acceptance-trace matrix	0.8650	+0.0375	applied
prompt-26	prompt	Failure-state recovery loop	0.8675	+0.0400	applied
prompt-27	prompt	Local-convention triangulation	0.8625	+0.0350	proposed
prompt-28	prompt	Command-grounded validation gate	0.8625	+0.0350	applied
prompt-29	prompt	Diff-to-requirement closure audit	0.8725	+0.0450	applied
prompt-30	prompt	Hypothesis diversification after stagnation	0.8450	+0.0175	applied

D.6 Lingxi-v1.5__claude-4-sonnet-20250514

Table 23: Record for Lingxi-v1.5__claude-4-sonnet-20250514 (rank 6).

Submission	20250720_Lingxi-v1.5__claude-4-sonnet-20250514 (verified)
Official resolve rate	74.60% (373 instances)
Repository	nimasteryang/Lingxi @ 1f2e5dc4c8
Recorded trajectories readable	0 (archive assets are credentialed)
Prompt-surface files found	16
Runnability	cli entry point plausible; execution still requires SWE-bench task in...
Baseline dims (TA/RS/EA/OI)	0.79 / 0.75 / 0.72 / 0.81
Trajectory	0.7675 $\rightarrow$ 0.8050 $\rightarrow$ 0.8350 (all PROXY)
Gate verdict	accept: mean not below reference; 4W/0L paired (sign p=0.125)
MDE ₈₀ at $n=4$	0.297; shipped $\Delta = +0.0675$ (below MDE: yes)
Paraphrase noise floor	+0.0050 (exceeds_noise_floor)
Blinded A/B ($k=5$)	patched 5 / baseline 0 / tie 0

Judge rationale (baseline). The two resolved probes show that the default multi-agent workflow can produce effective repository-level repairs, but the two unresolved probes demonstrate substantial brittleness on more difficult instances. Scores are therefore kept materially below 1.0 and calibrated to the reported 74.6% Pass@1 rate, with execution accuracy weakest because successful analysis does not consistently translate into a validated, correct patch.

Diagnosed failure modes.

1. **Incomplete issue localization and scope control** — The system can fail to identify all relevant files or the precise behavioral contract, leading to unresolved patches on repository-specific issues. *Evidence:* `django__django-13794`, `sphinx-doc__sphinx-9229`.
2. **Insufficient framework-specific diagnosis** — Multi-agent analysis and historical guidance do not reliably produce a sound explanation for harder Django and Sphinx behaviors, resulting in incorrect or incomplete fixes. *Evidence:* `django__django-13794`, `sphinx-doc__sphinx-9229`.
3. **Weak implementation validation and iteration** — The default workflow does not consistently catch failed assumptions through tests, focused reproduction, or follow-up edits; this is reflected by unresolved outcomes despite the available repository and editing tools. *Evidence:* `django__django-13794`, `sphinx-doc__sphinx-9229`.

Top three why-records (verbatim).

1. **prompt-17** (Failing-reproducer-first validation loop; 0.8350, $\Delta +0.0675$, applied): A behavioral contract followed by a failing-reproducer-first loop directly improves validation, regression awareness, and API preservation. This is the largest execution and output-integrity improvement, though it can delay patching when the reproduction is difficult to isolate.
2. **prompt-18** (Layered test, diff, and regression iteration; 0.8225, $\Delta +0.0550$, applied): Ordered reproducer, focused-test, regression, and diff gates provide strong implementation validation and iteration. The duplicated evidence-recording requirement and potentially expensive layered test sequence reduce efficiency and do not materially improve initial localization or framework diagnosis.
3. **prompt-15** (Framework execution-model diagnosis before patch design; 0.8175, $\Delta +0.0500$, applied): Reconstructing lifecycle, dispatch, registries, inheritance, normalization, and caching before patching is the strongest direct response to framework-specific misdiagnosis. It raises reasoning and execution accuracy, with some risk of excessive analysis before producing a fix.

Table 24: All 36 scored versions of Lingxi-v1.5_claude-4-sonnet-20250514.

id	phase	technique	score	Δ	decision
struct-01	struct	Add evidence summarization to the initial localization agent	0.7925	+0.0250	applied
struct-02	struct	Give localization access to a patch-state checkpoint	0.7650	-0.0025	proposed
struct-03	struct	Add structured evidence synthesis to solution mapping	0.7925	+0.0250	applied
struct-04	struct	Make the solution mapper inspect the actual patch state	0.8000	+0.0325	proposed
struct-05	struct	Allow the mapping stage to perform the edit-and-check loop	0.7875	+0.0200	proposed
struct-06	struct	Add reviewer guidance to the supervisor graph surface	0.7825	+0.0150	applied
struct-07	struct	Add a mandatory repository-inventory and behavioral-contract phas...	0.7950	+0.0275	applied
struct-08	struct	Require dependency-closure mapping before selecting an implementa...	0.7950	+0.0275	applied
struct-09	struct	Use hypothesis-driven framework diagnosis with competing explanat...	0.8050	+0.0375	applied
struct-10	struct	Add a framework-semantics evidence table to the multi-agent synth...	0.7875	+0.0200	applied
struct-11	struct	Turn review into a focused reproduction-and-regression validation...	0.8025	+0.0350	applied
struct-12	struct	Add bounded reviewer-driven repair iterations with explicit stop ...	0.7875	+0.0200	applied
struct-13	struct	Contract-first repository scope map with explicit search completi...	0.8000	+0.0325	applied
struct-14	struct	Evidence-ranked relevant-file retrieval with negative-scope checks	0.7875	+0.0200	applied
struct-15	struct	Historical-analogy triangulation with repository-local falsificat...	0.7575	-0.0100	applied
struct-16	struct	Framework causal-model handoff with disagreement resolution by re...	0.8025	+0.0350	applied
struct-17	struct	Test-plan gate and reproduce-before-edit hand-off	0.7975	+0.0300	proposed
struct-18	struct	Post-edit verification loop with diff, focused test, and bounded ...	0.8000	+0.0325	proposed
prompt-01	prompt	Contract-to-code scope map	0.7925	+0.0250	applied
prompt-02	prompt	Negative-space repository search	0.7875	+0.0200	proposed
prompt-03	prompt	Framework lifecycle state-machine diagnosis	0.7950	+0.0275	proposed
prompt-04	prompt	Competing-cause falsification	0.7775	+0.0100	proposed
prompt-05	prompt	Reproduction-gated implementation loop	0.8000	+0.0325	proposed
prompt-06	prompt	Diff-and-regression completion checklist	0.7975	+0.0300	proposed
prompt-07	prompt	Evidence-backed contract and ranked scope map	0.8050	+0.0375	applied
prompt-08	prompt	Bidirectional dependency and exclusion analysis	0.8000	+0.0325	applied
prompt-09	prompt	Framework lifecycle differential diagnosis	0.8125	+0.0450	applied
prompt-10	prompt	Historical guidance as a versioned hypothesis	0.7775	+0.0100	proposed
prompt-11	prompt	Minimal reproduction before implementation	0.8125	+0.0450	applied
prompt-12	prompt	Post-patch test triage and bounded repair loop	0.8150	+0.0475	proposed
prompt-13	prompt	Evidence-backed scope matrix with explicit stop rule	0.8100	+0.0425	applied
prompt-14	prompt	Alternative-location elimination and minimal-patch gate	0.8050	+0.0375	applied
prompt-15	prompt	Framework execution-model diagnosis before patch design	0.8175	+0.0500	applied
prompt-16	prompt	Competing hypotheses with discriminating reproductions	0.8050	+0.0375	applied
prompt-17	prompt	Failing-reproducer-first validation loop	0.8350	+0.0675	applied
prompt-18	prompt	Layered test, diff, and regression iteration	0.8225	+0.0550	applied

D.7 Moatless Tools + Claude 4 Sonnet

Table 25: Record for Moatless Tools (rank 7).

Submission	20250611_moatless_claude-4-sonnet-20250514 (verified)
Official resolve rate	70.80% (354 instances)
Repository	aorwall/moatless-tools @ 011ead57a5
Recorded trajectories readable	0 (archive assets are credentialed)
Prompt-surface files found	20
Runnability	docker required
Baseline dims (TA/RS/EA/OI)	0.83 / 0.77 / 0.74 / 0.84
Trajectory	0.7950 $\rightarrow$ 0.8425 $\rightarrow$ 0.8500 (all PROXY)
Gate verdict	accept: mean not below reference; 4W/0L paired (sign p=0.125)
MDE ₈₀ at $n=4$	0.297; shipped $\Delta = +0.0550$ (below MDE: yes)
Paraphrase noise floor	+0.0000 (exceeds_noise_floor)
Blinded A/B ($k=5$)	patched 5 / baseline 0 / tie 0 (post-repair)

Judge rationale (baseline). The default system officially resolves 70.8% of verified SWE-bench instances, so its aggregate capability is substantial but clearly not near-perfect. The fixed probes reinforce that split: both Django instances resolved, while the pytest and seaborn instances remained unresolved; the prompt and completion infrastructure support disciplined tool use and output handling, but execution accuracy is the main limiting factor.

Diagnosed failure modes.

1. **Failure to converge on repository-specific fixes** — The system can resolve some Django tasks but does not reliably translate investigation into a correct patch on other repositories. *Evidence:* `pytest-dev__pytest-5840`, `mwaskom__seaborn-3187`.
2. **Insufficient iterative verification** — The strict single-action workflow supports controlled execution, but the unresolved probes indicate that investigation, editing, and validation do not consistently converge on passing behavior. *Evidence:* `pytest-dev__pytest-5840`, `mwaskom__seaborn-3187`.
3. **Brittle cross-repository generalization** — The recorded outcomes show successful handling of both Django probes but failures on the pytest and seaborn probes, suggesting weaknesses that emerge across different codebase conventions and test ecosystems. *Evidence:* `pytest-dev__pytest-5840`, `mwaskom__seaborn-3187`.

Top three why-records (verbatim).

1. **prompt-18** (Explicit completion contract; 0.8500, $\Delta +0.0550$, proposed_not_applied): The explicit completion contract strongly improves execution accuracy and output integrity by requiring confirmation of the patch, targeted checks, and diff scope before completion. It helps prevent premature success declarations, though it contributes less to initial repository-specific diagnosis.
2. **struct-03** (Explicit post-edit convergence loop; 0.8425, $\Delta +0.0475$, proposed_not_applied): The explicit post-edit test loop directly addresses non-convergence and weak verification, producing the strongest expected gains on the previously weak probes. The exactly-one-action constraint is somewhat rigid, preventing a larger improvement and occasionally adding overhead.
3. **prompt-10** (Patch-scope and invariant contract; 0.8375, $\Delta +0.0425$, proposed_not_applied): Requiring an explicit behavioral goal, preserved invariant, minimality rationale, and caller/subclass inspection should reduce repository-specific regressions and improve patch targeting. It raises reasoning and integrity more than execution because it still lacks a dedicated verification loop.

Table 26: All 36 scored versions of Moatless Tools.

id	phase	technique	score	Δ	decision
struct-01	struct	Repository fingerprinting gate	0.8075	+0.0125	proposed
struct-02	struct	Test-first execution contract	0.8250	+0.0300	proposed
struct-03	struct	Explicit post-edit convergence loop	0.8425	+0.0475	proposed
struct-04	struct	Structured response precedence	0.8025	+0.0075	applied
struct-05	struct	Null-safe completion handling	0.8025	+0.0075	applied
struct-06	struct	Structured-output preservation	0.8075	+0.0125	applied
struct-07	struct	Repository reconnaissance gate	0.8175	+0.0225	proposed
struct-08	struct	Hypothesis and acceptance-criteria checkpoint	0.8225	+0.0275	proposed
struct-09	struct	Mandatory post-edit test loop	0.8250	+0.0300	proposed
struct-10	struct	Failure-classification recovery policy	0.7750	-0.0200	proposed
struct-11	struct	Framework-aware test discovery directive	0.8300	+0.0350	proposed
struct-12	struct	Final diff and regression audit	0.8275	+0.0325	proposed
struct-13	struct	Repository reconnaissance gate	0.8300	+0.0350	proposed
struct-14	struct	Edit-to-test convergence loop	0.8175	+0.0225	proposed
struct-15	struct	Failure-output triage protocol	0.8250	+0.0300	proposed
struct-16	struct	Hypothesis ledger for edits	0.8325	+0.0375	proposed
struct-17	struct	Cross-repository convention discovery	0.8100	+0.0150	proposed
struct-18	struct	Evidence-based completion gate	0.8350	+0.0400	proposed
prompt-01	prompt	Repository-specific hypothesis-to-patch contract	0.8275	+0.0325	proposed
prompt-02	prompt	Mandatory post-edit validation loop	0.8150	+0.0200	proposed
prompt-03	prompt	Cross-repository convention discovery	0.8250	+0.0300	proposed
prompt-04	prompt	Acceptance-criteria traceability	0.8200	+0.0250	proposed
prompt-05	prompt	Failure-driven test selection	0.8275	+0.0325	proposed
prompt-06	prompt	Diff-and-scope safety gate	0.8000	+0.0050	proposed
prompt-07	prompt	Repository-specific hypothesis ledger	0.8100	+0.0150	proposed
prompt-08	prompt	Failing-test-first convergence loop	0.8200	+0.0250	proposed
prompt-09	prompt	Cross-ecosystem reconnaissance gate	0.8225	+0.0275	proposed
prompt-10	prompt	Patch-scope and invariant contract	0.8375	+0.0425	proposed
prompt-11	prompt	Observation-to-next-action protocol	0.8250	+0.0300	proposed
prompt-12	prompt	Test-environment-aware validation	0.8275	+0.0325	proposed
prompt-13	prompt	Repository-specific hypothesis formation	0.8300	+0.0350	proposed
prompt-14	prompt	Evidence-gated edit loop	0.8325	+0.0375	proposed
prompt-15	prompt	Failure-driven recovery	0.8350	+0.0400	proposed
prompt-16	prompt	Cross-ecosystem reconnaissance	0.8250	+0.0300	proposed
prompt-17	prompt	Minimal-diff discipline	0.8100	+0.0150	proposed
prompt-18	prompt	Explicit completion contract	0.8500	+0.0550	proposed

D.8 SWE-agent + Claude 4 Sonnet

Table 27: Record for SWE-agent (rank 8).

Submission	20250522_sweagent_claude-4-sonnet-20250514 (verified)
Official resolve rate	66.60% (333 instances)
Repository	SWE-agent/SWE-agent @ 3ea751c087
Recorded trajectories readable	0 (archive assets are credentialed)
Prompt-surface files found	20
Runnability	cli entry point plausible; execution still requires SWE-bench task in...
Baseline dims (TA/RS/EA/OI)	0.74 / 0.69 / 0.64 / 0.76
Trajectory	0.7075 $\rightarrow$ 0.7650 $\rightarrow$ 0.7675 (all PROXY)
Gate verdict	accept: mean not below reference; 4W/0L paired (sign p=0.125)
MDE ₈₀ at $n=4$	0.297; shipped $\Delta = +0.0600$ (below MDE: yes)
Paraphrase noise floor	+0.0000 (exceeds_noise_floor)
Blinded A/B ($k=5$)	patched 5 / baseline 0 / tie 0

Judge rationale (baseline). The two resolved probes show strong task targeting and usable output, but the two unresolved probes directly demonstrate substantial gaps in execution accuracy and reasoning reliability. Those failures are consistent with, and calibrated to, the system’s approximately 66.6% overall resolve rate rather than near-perfect performance.

Diagnosed failure modes.

1. **Incorrect or incomplete repository changes** — The system can fail to turn an otherwise plausible approach into a patch that resolves the issue, as shown by unresolved outcomes on both Django and xarray. *Evidence:* `django__django-13112`, `pydata__xarray-6599`.
2. **Insufficient validation and iteration** — The retry/tool-driven setup does not reliably detect and repair remaining defects before submission, evidenced by two unresolved probe instances despite the system’s generally capable agent configuration. *Evidence:* `django__django-13112`, `pydata__xarray-6599`.
3. **Difficulty with repository-specific edge cases** — Performance is not uniformly reliable across projects: Sphinx and Matplotlib were resolved, while Django and xarray were not, indicating recurring sensitivity to project semantics and edge-case behavior. *Evidence:* `django__django-13112`, `pydata__xarray-6599`.

Top three why-records (verbatim).

1. **prompt-06** (Failure-triage and iterative repair protocol; 0.7675, $\Delta +0.0600$, proposed_not_applied): Baseline establishment, narrow checks, explicit failure classification, and iterative repair provide the most complete response to insufficient validation. It can still misclassify environment or expectation failures and does not explicitly require comprehensive boundary analysis, so gains remain moderate.
2. **prompt-09** (Repository-convention reconnaissance; 0.7675, $\Delta +0.0600$, proposed_not_applied): Reconnaissance is well matched to repository-specific edge cases and should improve API, compatibility, and error-handling decisions in the difficult probes. It can still lead to analysis without sufficient validation, so execution improvement is smaller than with the validation-gate candidate.
3. **prompt-11** (Minimal-diff semantic preservation; 0.7675, $\Delta +0.0600$, applied): Minimal, invariant-preserving changes reduce regressions and improve output integrity, particularly for already well-understood fixes. The conservative approach can under-address multi-branch or repository-specific requirements, so gains on the hardest xarray case are limited and execution remains only modestly above baseline.

Table 28: All 36 scored versions of SWE-agent.

id	phase	technique	score	Δ	decision
struct-01	struct	Add an explicit inspect-patch-test-repair contract to the agent's...	0.7375	+0.0300	applied
struct-02	struct	Require the agent to trace behavior across callers, subclasses, s...	0.7375	+0.0300	applied
struct-03	struct	Add a submission gate requiring a clean diff, focused regression ...	0.7450	+0.0375	applied
struct-04	struct	Increase the per-instance action budget so the retry/tool loop ha...	0.6975	-0.0100	proposed
struct-05	struct	Force explicit reproduction of the reported behavior before imple...	0.7550	+0.0475	applied
struct-06	struct	Replace the permissive one-line system instruction with a structu...	0.7450	+0.0375	applied
struct-07	struct	Add a mandatory post-edit diff-integrity gate to the agent's oper...	0.7325	+0.0250	applied
struct-08	struct	Require implementation-to-call-site tracing before editing: ident...	0.7450	+0.0375	applied
struct-09	struct	Replace one-shot validation with a required failure-driven repair...	0.7525	+0.0450	applied
struct-10	struct	Add a final consistency audit requiring the agent to compare the ...	0.7300	+0.0225	applied
struct-11	struct	Require an explicit edge-case matrix for every behavior change, c...	0.7350	+0.0275	applied
struct-12	struct	Add a repository-convention discovery phase that requires searchi...	0.7550	+0.0475	applied
struct-13	struct	Add an explicit patch-contract protocol before implementation	0.7500	+0.0425	proposed
struct-14	struct	Require a post-edit diff and test-failure repair gate	0.7575	+0.0500	proposed
struct-15	struct	Use an edge-case matrix derived from neighboring implementations	0.7450	+0.0375	proposed
struct-16	struct	Separate diagnosis, implementation, and verification into mandato...	0.7650	+0.0575	proposed
struct-17	struct	Add a hidden-test-oriented behavioral contract review	0.7050	-0.0025	proposed
struct-18	struct	Make repository conventions and compatibility branches first-clas...	0.7550	+0.0475	proposed
prompt-01	prompt	Explicit patch-completeness checklist	0.7350	+0.0275	applied
prompt-02	prompt	Failing-test-first execution loop	0.7500	+0.0425	applied
prompt-03	prompt	Behavioral contract extraction	0.7475	+0.0400	applied
prompt-04	prompt	Diff-and-scope audit	0.7525	+0.0450	proposed
prompt-05	prompt	Metamorphic and boundary validation	0.7450	+0.0375	applied
prompt-06	prompt	Failure-triage and iterative repair protocol	0.7675	+0.0600	proposed
prompt-07	prompt	Acceptance-criteria contract before editing	0.7600	+0.0525	applied
prompt-08	prompt	Test-driven validation gate	0.7550	+0.0475	proposed
prompt-09	prompt	Repository-convention reconnaissance	0.7675	+0.0600	proposed
prompt-10	prompt	Boundary-matrix probing	0.7400	+0.0325	applied
prompt-11	prompt	Minimal-diff semantic preservation	0.7675	+0.0600	applied
prompt-12	prompt	Independent final-review pass	0.7600	+0.0525	proposed
prompt-13	prompt	Acceptance-contract decomposition with explicit invariants	0.7425	+0.0350	applied
prompt-14	prompt	Minimal-diff and neighboring-pattern constraint	0.7450	+0.0375	applied
prompt-15	prompt	Failing-test-first validation loop	0.7475	+0.0400	applied
prompt-16	prompt	Layered test escalation with changed-file inspection	0.7550	+0.0475	applied
prompt-17	prompt	Boundary-condition and compatibility matrix	0.7400	+0.0325	proposed
prompt-18	prompt	Repository-semantics triangulation	0.7500	+0.0425	proposed

D.9 AgentScope

Table 29: Record for AgentScope (rank 9).

Submission	20250206_agentscope (verified)
Official resolve rate	63.40% (317 instances)
Repository	modelscope/agentscope @ 29b592358c
Recorded trajectories readable	0 (archive assets are credentialed)
Prompt-surface files found	20
Runnability	cli entry point plausible; execution still requires SWE-bench task in...
Baseline dims (TA/RS/EA/OI)	0.68 / 0.62 / 0.60 / 0.65
Trajectory	0.6375 $\rightarrow$ 0.7050 $\rightarrow$ 0.7150 (all PROXY)
Gate verdict	accept: mean not below reference; 4W/0L paired (sign p=0.125)
MDE ₈₀ at $n=4$	0.297; shipped $\Delta = +0.0775$ (below MDE: yes)
Paraphrase noise floor	+0.0075 (exceeds_noise_floor)
Blinded A/B ($k=5$)	patched 5 / baseline 0 / tie 0

Judge rationale (baseline). The fixed probes show 2 of 4 resolved, while the broader verified split reports a 63.4% resolve rate, so the dimensions are kept moderate rather than near-perfect. The successful Django cases show useful task alignment, but failures on Sphinx and Seaborn materially reduce reasoning, execution, and output confidence for the unmodified system.

Diagnosed failure modes.

1. **Cross-repository generalization failure** — The system resolved both Django probes but failed on the Sphinx and Seaborn probes, indicating inconsistent transfer across repository ecosystems. *Evidence: sphinx-doc__sphinx-8056, mwaskom__seaborn-3187.*
2. **Insufficient issue-specific patch execution** — The unresolved Sphinx and Seaborn outcomes suggest that identifying or implementing the required repository-specific changes is unreliable. *Evidence: sphinx-doc__sphinx-8056, mwaskom__seaborn-3187.*
3. **Weak validation of completed fixes** — The two unresolved outcomes indicate that the default workflow does not consistently produce a verified, task-complete result outside the successful Django cases. *Evidence: sphinx-doc__sphinx-8056, mwaskom__seaborn-3187.*

Top three why-records (verbatim).

1. **prompt-10** (Reproduction-first and test-driven patching; 0.7150, $\Delta +0.0775$, proposed_not_applied): Reproduction-first patching most directly addresses the diagnosed execution failure: it ties the change to observed behavior and creates a regression guard, with especially large expected gains on the weaker repositories. It still lacks the explicit final validation and adversarial diff review of later variants.
2. **struct-06** (Diff-and-regression audit; 0.7050, $\Delta +0.0675$, applied): The final diff and regression audit gives the strongest output integrity and catches unintended edits, compatibility problems, missing tests, and ambiguous validation claims. It improves execution and the weak probes, but is slightly less forceful than the iterative barrier about repairing a failing implementation.
3. **prompt-11** (Explicit validation gate; 0.7025, $\Delta +0.0650$, proposed_not_applied): The explicit validation gate materially improves output integrity by requiring narrow regression tests, broader relevant tests, and inspection of failures before completion. It does less to improve initial issue interpretation or patch construction, so execution gains are smaller than with reproduction-first patching.

Table 30: All 36 scored versions of AgentScope.

id	phase	technique	score	Δ	decision
struct-01	struct	Repository reconnaissance gate	0.6625	+0.0250	applied
struct-02	struct	Acceptance-criteria decomposition	0.6725	+0.0350	applied
struct-03	struct	Evidence-bearing team handoff	0.6450	+0.0075	applied
struct-04	struct	Repository-native test discovery	0.6700	+0.0325	applied
struct-05	struct	Iterative validation barrier	0.6950	+0.0575	applied
struct-06	struct	Diff-and-regression audit	0.7050	+0.0675	applied
struct-07	struct	Repository preflight fingerprinting	0.6625	+0.0250	applied
struct-08	struct	Repository-convention and compatibility discovery	0.6725	+0.0350	applied
struct-09	struct	Issue-to-symbol execution plan	0.6900	+0.0525	applied
struct-10	struct	Regression-test-first patching	0.6350	-0.0025	applied
struct-11	struct	Evidence-gated completion validation	0.6875	+0.0500	applied
struct-12	struct	Failure-driven repair loop	0.6725	+0.0350	applied
struct-13	struct	Repository fingerprinting and ecosystem-adaptive workflow	0.6625	+0.0250	applied
struct-14	struct	Issue-contract extraction before editing	0.6725	+0.0350	proposed
struct-15	struct	Falsifiable hypothesis and minimal coherent patching	0.6875	+0.0500	proposed
struct-16	struct	Progressive regression-test execution	0.6650	+0.0275	proposed
struct-17	struct	Failure-cause classification with targeted recovery	0.6125	-0.0250	proposed
struct-18	struct	Evidence-gated completion and diff audit	0.6825	+0.0450	proposed
prompt-01	prompt	Repository-family reconnaissance checklist	0.6625	+0.0250	applied
prompt-02	prompt	Framework-convention anchoring	0.6600	+0.0225	applied
prompt-03	prompt	Requirement-to-code trace	0.6800	+0.0425	applied
prompt-04	prompt	Failure-localization repair loop	0.6675	+0.0300	applied
prompt-05	prompt	Layered validation gate	0.6675	+0.0300	applied
prompt-06	prompt	Evidence-based completion report	0.6600	+0.0225	applied
prompt-07	prompt	Repository fingerprinting and ecosystem adaptation	0.6500	+0.0125	applied
prompt-08	prompt	Evidence-backed cross-repository analogical reasoning	0.6750	+0.0375	proposed
prompt-09	prompt	Issue-to-code traceability plan	0.6975	+0.0600	proposed
prompt-10	prompt	Reproduction-first and test-driven patching	0.7150	+0.0775	proposed
prompt-11	prompt	Explicit validation gate	0.7025	+0.0650	proposed
prompt-12	prompt	Adversarial regression and diff review	0.6900	+0.0525	proposed
prompt-13	prompt	Repository fingerprinting with ecosystem-specific adaptation	0.6525	+0.0150	proposed
prompt-14	prompt	Local-pattern triangulation before implementation	0.6800	+0.0425	proposed
prompt-15	prompt	Issue-to-acceptance-contract decomposition	0.6875	+0.0500	proposed
prompt-16	prompt	Surgical implementation loop with dependency tracing	0.7000	+0.0625	proposed
prompt-17	prompt	Progressive test ladder with failure-driven repair	0.6675	+0.0300	proposed
prompt-18	prompt	Adversarial completion and diff audit	0.6875	+0.0500	proposed

D.10 EntroPO + R2E + Qwen3-Coder-30B-A3B-Instruct

Table 31: Record for EntroPO (rank 10).

Submission	20250901_entroPO_R2E_QwenCoder30BA3B_tts (verified)
Official resolve rate	60.40% (302 instances)
Repository	sherdencooper/R2E-Gym @ eee9f6f00a
Recorded trajectories readable	0 (archive assets are credentialed)
Prompt-surface files found	20
Runnability	cli entry point plausible; execution still requires SWE-bench task in...
Baseline dims (TA/RS/EA/OI)	0.66 / 0.61 / 0.58 / 0.67
Trajectory	0.6300 $\rightarrow$ 0.6825 $\rightarrow$ 0.6900 (all PROXY)
Gate verdict	accept: mean not below reference; 4W/0L paired (sign p=0.125)
MDE ₈₀ at $n=4$	0.297; shipped $\Delta = +0.0600$ (below MDE: yes)
Paraphrase noise floor	+0.0075 (exceeds_noise_floor)
Blinded A/B ($k=5$)	patched 5 / baseline 0 / tie 0

Judge rationale (baseline). The fixed probes show 2 resolved and 2 unresolved outcomes, which supports middling rather than near-perfect scores and is broadly consistent with the system’s officially resolved rate of 60.4%. The agent has substantial execution infrastructure, but its multi-turn context condensation and the two concrete failures indicate weaknesses in reasoning continuity, verification, and final patch accuracy.

Diagnosed failure modes.

1. **Incorrect or incomplete patch implementation** — The system failed to resolve two of the four fixed probes, indicating recurring difficulty translating the task requirements into a correct, complete repository change. *Evidence:* `django__django-15930`, `django__django-11820`.
2. **Insufficient test-driven verification and iteration** — The unresolved outcomes are consistent with failures to validate the proposed change against the relevant regression behavior before finishing; the repository exposes bash, execution, and submission tools, but successful tool availability does not ensure accurate verification. *Evidence:* `django__django-15930`, `django__django-11820`.
3. **Long-horizon context degradation** — The verifier explicitly condenses older execution-result blocks once the Qwen context exceeds 31,000 tokens, creating a recurring risk that debugging evidence or prior assumptions are lost during multi-turn repair attempts. *Evidence:* `django__django-15930`, `django__django-11820`.

Top three why-records (verbatim).

1. **prompt-06** (Evidence-prioritized context recovery; 0.6900, $\Delta +0.0600$, proposed_not_applied): Re-grounding in the issue, current diff, source, and fresh command evidence is a strong long-horizon recovery strategy and improves the reliability of reported state. It helps expose stale hypotheses and hidden failures, but it is less directly prescriptive about implementing every requirement than the coverage contract.
2. **struct-12** (Make the condensed-history instruction explicitly require regression-...; 0.6825, $\Delta +0.0525$, proposed_not_applied): The stronger condensation instruction directly counters incomplete patches by requiring inspection and task-specific regression testing before submission, making it the most relevant change for execution accuracy while still helping long-horizon reasoning.
3. **prompt-12** (State-restoration protocol; 0.6800, $\Delta +0.0500$, proposed_not_applied): Reconstructing the current diff, unmet requirements, failures, and passing tests on every resumed turn directly mitigates context degradation while retaining repository-grounded state. It is stronger than a passive ledger, but repeated reconstruction adds overhead and still lacks the explicit test expansion of candidates 2 and 3.

Table 32: All 36 scored versions of EntroPO.

id	phase	technique	score	Δ	decision
struct-01	struct	Requirement-aware context preservation	0.6550	+0.0250	applied
struct-02	struct	Patch-state retention with bounded block compression	0.6550	+0.0250	proposed
struct-03	struct	Test-aware evidence preservation	0.6650	+0.0350	proposed
struct-04	struct	Preserve the verification tail	0.6625	+0.0325	proposed
struct-05	struct	Adaptive context budget	0.6000	-0.0300	proposed
struct-06	struct	Tokenizer reuse and stable token accounting	0.6475	+0.0175	proposed
struct-07	struct	Increase the verifier’s usable context budget so the task specifi...	0.6500	+0.0200	applied
struct-08	struct	Use an explicit, information-preserving condensation marker that ...	0.6725	+0.0425	applied
struct-09	struct	Add a deterministic token-bounded fallback when the trajectory ha...	0.6025	-0.0275	applied
struct-10	struct	Cache the tokenizer used by the condensation routine to make repe...	0.6325	+0.0025	proposed
struct-11	struct	Reserve a smaller global agent context budget so tool outputs and...	0.6025	-0.0275	applied
struct-12	struct	Make the condensed-history instruction explicitly require regress...	0.6825	+0.0525	proposed
struct-13	struct	Inject a structured patch-completeness checklist into condensed h...	0.6600	+0.0300	applied
struct-14	struct	Reserve context headroom before condensation	0.6250	-0.0050	applied
struct-15	struct	Make condensed history explicitly require regression-test evidence	0.6700	+0.0400	proposed
struct-16	struct	Cache the tokenizer used during repeated condensation	0.6325	+0.0025	proposed
struct-17	struct	Lower the global agent context ceiling to avoid hard-limit degrad...	0.5975	-0.0325	applied
struct-18	struct	Use an explicit condensation budget below the model context limit	0.6125	-0.0175	applied
prompt-01	prompt	Requirements-to-code coverage contract	0.6725	+0.0425	applied
prompt-02	prompt	Hypothesis-driven patch review	0.6650	+0.0350	proposed
prompt-03	prompt	Mandatory regression-test gate	0.6575	+0.0275	proposed
prompt-04	prompt	Two-phase validation loop	0.6750	+0.0450	proposed
prompt-05	prompt	Persistent debugging state ledger	0.6550	+0.0250	proposed
prompt-06	prompt	Evidence-prioritized context recovery	0.6900	+0.0600	proposed
prompt-07	prompt	Acceptance-contract decomposition	0.6650	+0.0350	applied
prompt-08	prompt	Invariant-and-edge-case checklist	0.6325	+0.0025	proposed
prompt-09	prompt	Test-first regression loop	0.6600	+0.0300	proposed
prompt-10	prompt	Failure-directed test expansion	0.6550	+0.0250	proposed
prompt-11	prompt	Persistent decision ledger	0.6675	+0.0375	proposed
prompt-12	prompt	State-restoration protocol	0.6800	+0.0500	proposed
prompt-13	prompt	Atomic acceptance-contract expansion	0.6475	+0.0175	applied
prompt-14	prompt	Requirement-to-code completeness gate	0.6600	+0.0300	applied
prompt-15	prompt	Mandatory regression-test and diff loop	0.6700	+0.0400	applied
prompt-16	prompt	Observable-behavior test oracle	0.6450	+0.0150	applied
prompt-17	prompt	Compact state checkpoint preservation	0.6450	+0.0150	applied
prompt-18	prompt	Evidence-first history recovery	0.6775	+0.0475	applied

D.11 ExpeRepair-v1.0 + Claude 4 Sonnet

Table 33: Record for ExpeRepair-v1.0 (rank 11).

Submission	20250625_ExpeRepair-v1_claude-4-sonnet-20250514 (lite)
Official resolve rate	60.33% (181 instances)
Repository	ExpeRepair/ExpeRepair @ 5594f2c02c
Recorded trajectories readable	0 (archive assets are credentialed)
Prompt-surface files found	20
Runnability	cli entry point plausible; execution still requires SWE-bench task in...
Baseline dims (TA/RS/EA/OI)	0.74 / 0.68 / 0.61 / 0.79
Trajectory	0.7050 $\rightarrow$ 0.7750 $\rightarrow$ 0.7750 (all PROXY)
Gate verdict	accept: mean not below reference; 4W/0L paired (sign p=0.125)
MDE ₈₀ at $n=4$	0.297; shipped $\Delta = +0.0700$ (below MDE: yes)
Paraphrase noise floor	+0.0000 (exceeds_noise_floor)
Blinded A/B ($k=5$)	patched 5 / baseline 0 / tie 0

Judge rationale (baseline). The observed 2/4 probe success rate is consistent with a system that has useful issue analysis and repair capability but substantial execution failures, and it is compatible with the reported 60.33% overall resolve rate. The explicit search, reproduction, and validation stages support the identified failure modes, while the two resolved SymPy probes show that the workflow can succeed on some instances.

Diagnosed failure modes.

1. **Incorrect repository localization** — The search stage must select relevant editable files from the repository structure; failures on both Django probes are consistent with localization errors or incomplete cause analysis. *Evidence:* `django__django-11620`, `django__django-12497`.
2. **Unreliable reproduction or validation judgment** — The system relies on generated reproduction tests and an LLM validation stage, creating a failure mode when tests are invalid, incomplete, or misinterpreted. *Evidence:* `django__django-11620`, `django__django-12497`.
3. **Patch generation or execution failure** — The unresolved outcomes indicate that analysis did not consistently produce and successfully apply a correct repository-level fix, despite the multi-agent repair workflow. *Evidence:* `django__django-11620`, `django__django-12497`.

Top three why-records (verbatim).

1. **struct-06** (Minimal-diff, regression-aware patch selection; 0.7750, $\Delta +0.0700$, proposed_not_applied): Requiring an exact root cause, minimal repository-consistent changes, regression awareness, and validation combines useful safeguards against speculative or oversized patches. It improves execution and reasoning more than the baseline, but depends on earlier localization and reproduction being correct, so it is not a complete remedy.
2. **struct-05** (Patch application and changed-file integrity gate; 0.7625, $\Delta +0.0575$, proposed_not_applied): Clean-application, changed-file, and test-integrity checks directly reduce patch execution failures and improve output integrity. They do not establish that the selected location or behavioral hypothesis is correct, so reasoning and the low-performing Django cases improve only modestly.
3. **prompt-06** (Patch application and semantic integrity checklist; 0.7625, $\Delta +0.0575$, applied): The checklist directly addresses syntactic coherence, file existence, clean application, and semantic path coverage, producing the clearest improvement against patch execution failure. It does less to improve initial localization or issue interpretation than the evidence-first variants.

Table 34: All 36 scored versions of ExpeRepair-v1.0.

id	phase	technique	score	Δ	decision
struct-01	struct	Evidence-constrained, two-pass repository localization	0.7350	+0.0300	proposed
struct-02	struct	Require symbol-level localization and explicit exclusion evidence	0.7575	+0.0525	proposed
struct-03	struct	Deterministic reproduction with an executable oracle	0.7175	+0.0125	proposed
struct-04	struct	Preflight reproduction-test screening	0.7525	+0.0475	proposed
struct-05	struct	Patch application and changed-file integrity gate	0.7625	+0.0575	proposed
struct-06	struct	Minimal-diff, regression-aware patch selection	0.7750	+0.0700	proposed
struct-07	struct	Evidence-backed two-stage localization with symbol expansion	0.7325	+0.0275	proposed
struct-08	struct	Localization consensus and mandatory cause-path coverage	0.7375	+0.0325	proposed
struct-09	struct	Executable reproduction gate with baseline and patched comparisons	0.7275	+0.0225	proposed
struct-10	struct	Issue-derived oracle and negative-control validation	0.7425	+0.0375	proposed
struct-11	struct	Transactional patch application with static and runtime gates	0.7525	+0.0475	proposed
struct-12	struct	Candidate patch ensemble with minimality and regression scoring	0.7400	+0.0350	proposed
struct-13	struct	Evidence-grounded localization contract	0.7450	+0.0400	proposed
struct-14	struct	Two-stage source-file narrowing	0.7450	+0.0400	proposed
struct-15	struct	Minimal deterministic reproduction policy	0.7375	+0.0325	proposed
struct-16	struct	Validity-first patch evaluation	0.7550	+0.0500	proposed
struct-17	struct	Execution-gated patch acceptance	0.7475	+0.0425	proposed
struct-18	struct	Structured localization evidence output	0.6750	-0.0300	proposed
prompt-01	prompt	Evidence-first ranked localization	0.7325	+0.0275	applied
prompt-02	prompt	Implementation-versus-ancillary file filtering	0.7375	+0.0325	applied
prompt-03	prompt	Minimal deterministic regression-test design	0.6775	-0.0275	proposed
prompt-04	prompt	Baseline-reproduction self-check	0.7150	+0.0100	proposed
prompt-05	prompt	Validity-gated patch assessment	0.7525	+0.0475	applied
prompt-06	prompt	Patch application and semantic integrity checklist	0.7625	+0.0575	applied
prompt-07	prompt	Evidence-traced localization with role-aware filtering	0.7325	+0.0275	applied
prompt-08	prompt	Candidate localization by falsifiable hypotheses	0.7450	+0.0400	applied
prompt-09	prompt	Minimal contract-focused regression test generation	0.6625	-0.0425	proposed
prompt-10	prompt	Reproduction oracle and failure-cause recording	0.6850	-0.0200	proposed
prompt-11	prompt	Independent patch-versus-contract validation	0.7550	+0.0500	applied
prompt-12	prompt	Strict patch integrity and regression gate	0.7600	+0.0550	applied
prompt-13	prompt	Evidence-backed localization with ranked ownership hypotheses	0.7375	+0.0325	applied
prompt-14	prompt	Negative-evidence and boundary-file elimination	0.7400	+0.0350	applied
prompt-15	prompt	Independent oracle reconstruction before test interpretation	0.7400	+0.0350	applied
prompt-16	prompt	Minimal deterministic regression-test contract	0.6275	-0.0775	proposed
prompt-17	prompt	Patch applicability and execution checklist	0.7575	+0.0525	applied
prompt-18	prompt	Diff-scope and semantic-preservation audit	0.7600	+0.0550	applied

D.12 KGCompass + Claude 4 Sonnet (20250514)

Table 35: Record for KGCompass (rank 12).

Submission	20250906_KGCompass_claude-4-sonnet-20250514 (lite)
Official resolve rate	58.33% (175 instances)
Repository	GLEAM-Lab/KGCompass @ b74a584e6d
Recorded trajectories readable	0 (archive assets are credentialed)
Prompt-surface files found	4
Runnability	cli entry point plausible; execution still requires SWE-bench task in...
Baseline dims (TA/RS/EA/OI)	0.68 / 0.62 / 0.57 / 0.64
Trajectory	0.6275 $\rightarrow$ 0.6675 $\rightarrow$ 0.6900 (all PROXY)
Gate verdict	accept: mean not below reference; 2W/1L paired (sign p=1.0)
MDE ₈₀ at $n=4$	0.297; shipped $\Delta = +0.0625$ (below MDE: yes)
Paraphrase noise floor	+0.0025 (exceeds_noise_floor)
Blinded A/B ($k=5$)	patched 5 / baseline 0 / tie 0

Judge rationale (baseline). The system resolved 2 of the 4 fixed probes, consistent with its recorded 58.33% overall resolve rate, so its dimensions should remain materially below near-perfect performance. The repository shows a capable but layered pipeline—KG localization, LLM localization, strict edit parsing, syntax checking, and patch application—whose unresolved probes most strongly expose execution and reliability weaknesses.

Diagnosed failure modes.

1. **Fault-localization or context-selection misses** — The system depends on KG-derived locations and an additional LLM localization stage; the unresolved outcomes indicate that this pipeline does not reliably identify sufficient repair context. *Evidence:* `django__django-13768`, `django__django-11964`.
2. **Patch generation and application failures** — The repair path relies on generated edit commands, diff parsing, syntax checks, and patch application, creating multiple opportunities for a plausible repair to fail execution or tests. *Evidence:* `django__django-13768`, `django__django-11964`.
3. **Brittle output and orchestration handling** — The default system uses separate localization and repair scripts with strict structured parsing and multi-API/configuration plumbing; failures in formatting, parsing, or workflow coordination can prevent otherwise useful model output from becoming a valid submission. *Evidence:* `django__django-13768`, `django__django-11964`.

Top three why-records (verbatim).

1. **prompt-06** (Canonical edit serialization with preflight constraints; 0.6900, $\Delta +0.0625$, proposed_not_applied): Canonical unified-diff serialization and preflight constraints directly reduce malformed or unapplied edits and provide the strongest output-integrity improvement. It does not improve diagnosis and can preserve an incorrect patch cleanly, so target alignment and reasoning remain slightly below baseline.
2. **prompt-04** (Test-oriented patch self-review; 0.6825, $\Delta +0.0550$, proposed_not_applied): Focused mental execution and checks for control flow, types, imports, exceptions, and compatibility most directly address patch-generation failures. It produces the largest execution gain, though review is only as good as the initial diagnosis and does not fix orchestration format failures.
3. **prompt-18** (Parser-safe repair protocol with explicit failure signaling; 0.6775, $\Delta +0.0500$, proposed_not_applied): Parser-safe unified diffs directly address orchestration failures and explicit NO_SAFE_PATCH avoids fabricating unsafe edits. However, the conservative failure signal can forfeit solvable instances, and the protocol does not improve localization enough to resolve the hardest probes.

Table 36: All 60 scored versions of KGCompass.

id	phase	technique	score	Δ	decision
struct-01	struct	Search all KG-location layouts instead of assuming one exact dire...	0.6500	+0.0225	applied
struct-02	struct	Make the localization API client retry transient failures and tol...	0.6400	+0.0125	applied
struct-03	struct	Convert JSON-extraction exceptions into an explicit empty result ...	0.6250	-0.0025	applied
struct-04	struct	Isolate strict edit-command parsing failures behind a non-crashin...	0.6375	+0.0100	applied
struct-05	struct	Treat syntax-checker exceptions as validation failures instead of...	0.6475	+0.0200	proposed
struct-06	struct	Make patch-application errors return a failed candidate status	0.6525	+0.0250	proposed
struct-07	struct	Defensive multi-format localization-result decoding	0.6450	+0.0175	applied
struct-08	struct	Fail-closed instance and language grounding before localization	0.6550	+0.0275	applied
struct-09	struct	Type-normalizing and exception-safe edit-command boundary	0.6375	+0.0100	applied
struct-10	struct	Post-application transactional diff validation	0.6550	+0.0275	applied
struct-11	struct	Optional provider imports for single-provider execution	0.6325	+0.0050	applied
struct-12	struct	Configuration fallback with typed environment parsing	0.6400	+0.0125	applied
struct-13	struct	Make localization JSON extraction executable by importing its reg...	0.6475	+0.0200	applied
struct-14	struct	Correct fenced-JSON and object-boundary extraction patterns	0.6600	+0.0325	applied
struct-15	struct	Add a permissive structured-output fallback around strict edit-co...	0.6475	+0.0200	applied
struct-16	struct	Validate changed Python files for syntax immediately after patch ...	0.6575	+0.0300	applied
struct-17	struct	Provide environment-backed configuration defaults when optional c...	0.6475	+0.0200	applied
struct-18	struct	Make non-selected API clients optional at module import time	0.6350	+0.0075	applied
struct-19	struct	Increase the repair-context budget when no explicit configuration...	0.6500	+0.0225	applied
struct-20	struct	Accept ordinary unescaped JSON objects during localization-result...	0.6425	+0.0150	applied
struct-21	struct	Make localization JSON recovery tolerate non-JSON decoding except...	0.6350	+0.0075	applied
struct-22	struct	Convert patch-application exceptions into a controlled failed att...	0.6225	-0.0050	applied
struct-23	struct	Use a lower default generation temperature for deterministic code...	0.6375	+0.0100	applied
struct-24	struct	Make the OpenAI dependency optional at module import time	0.6325	+0.0050	applied
struct-25	struct	Make the repair-context budget configurable with a larger safe de...	0.6475	+0.0200	applied
struct-26	struct	Replace first-match JSON extraction with balanced JSONDecoder sca...	0.6575	+0.0300	applied
struct-27	struct	Treat patch-tool and filesystem errors as failed attempts rather ...	0.6400	+0.0125	applied
struct-28	struct	Make post-application whitespace verification failure-safe	0.6275	+0.0000	applied
struct-29	struct	Provide a tokenizer compatibility fallback for environments witho...	0.5750	-0.0525	applied
struct-30	struct	Add recovery parsing for edit blocks after strict diff-command pa...	0.6675	+0.0400	applied
prompt-01	prompt	Evidence-grounded hierarchical localization	0.6600	+0.0325	proposed
prompt-02	prompt	Bidirectional context expansion and hypothesis comparison	0.6725	+0.0450	proposed

id	phase	technique	score	Δ	decision
prompt-03	prompt	Diagnosis-first minimal patch planning	0.6675	+0.0400	proposed
prompt-04	prompt	Test-oriented patch self-review	0.6825	+0.0550	proposed
prompt-05	prompt	Single-schema machine-readable response contract	0.6425	+0.0150	proposed
prompt-06	prompt	Canonical edit serialization with preflight constraints	0.6900	+0.0625	proposed
prompt-07	prompt	Evidence-constrained localization with explicit candidate compari...	0.6575	+0.0300	proposed
prompt-08	prompt	Issue-to-invariant tracing	0.6625	+0.0350	proposed
prompt-09	prompt	Minimal patch contract with preservation constraints	0.6525	+0.0250	proposed
prompt-10	prompt	Test-first behavioral validation	0.6625	+0.0350	proposed
prompt-11	prompt	Canonical structured-output envelope	0.6525	+0.0250	proposed
prompt-12	prompt	Self-repairing serialization with parser-oriented constraints	0.6675	+0.0400	proposed
prompt-13	prompt	Evidence-ranked localization with explicit competing hypotheses	0.6625	+0.0350	proposed
prompt-14	prompt	Contract-first context analysis	0.6725	+0.0450	proposed
prompt-15	prompt	Self-checking minimal unified diff generation	0.6600	+0.0325	proposed
prompt-16	prompt	Post-patch invariant and regression checklist	0.6500	+0.0225	proposed
prompt-17	prompt	Schema-constrained single-channel output	0.5975	-0.0300	proposed
prompt-18	prompt	Parser-safe repair protocol with explicit failure signaling	0.6775	+0.0500	proposed
prompt-19	prompt	Evidence-first hierarchical localization	0.6600	+0.0325	proposed
prompt-20	prompt	Explicit search-boundary expansion and confidence calibration	0.6600	+0.0325	proposed
prompt-21	prompt	Machine-checkable patch contract	0.6725	+0.0450	proposed
prompt-22	prompt	Preflight patch self-review	0.6650	+0.0375	proposed
prompt-23	prompt	Schema-first single-response protocol	0.6100	-0.0175	proposed
prompt-24	prompt	Validation-aware repair loop	0.6625	+0.0350	proposed
prompt-25	prompt	Evidence-backed bidirectional localization	0.6675	+0.0400	proposed
prompt-26	prompt	Hypothesis comparison with explicit context expansion	0.6600	+0.0325	proposed
prompt-27	prompt	Minimal unified-diff repair contract	0.6675	+0.0400	proposed
prompt-28	prompt	Context-preserving surgical editing	0.6650	+0.0375	proposed
prompt-29	prompt	Single-channel machine-readable response	0.6325	+0.0050	proposed
prompt-30	prompt	Self-checking repair protocol with bounded fallback	0.6650	+0.0375	proposed

D.13 SWE-Rizzo

Table 37: Record for SWE-Rizzo (rank 13).

Submission	20250405_swe-rizzo_claude37 (verified)
Official resolve rate	56.60% (283 instances)
Repository	brokespace/gen42-codemonkeys @ c6303b8710
Recorded trajectories readable	0 (archive assets are credentialed)
Prompt-surface files found	14
Runnability	docker required
Baseline dims (TA/RS/EA/OI)	0.82 / 0.70 / 0.57 / 0.81
Trajectory	0.7250 $\rightarrow$ 0.7975 $\rightarrow$ 0.8150 (all PROXY)
Gate verdict	accept: mean not below reference; 4W/0L paired (sign p=0.125)
MDE ₈₀ at $n=4$	0.297; shipped $\Delta = +0.0900$ (below MDE: yes)
Paraphrase noise floor	+0.0000 (exceeds_noise_floor)
Blinded A/B ($k=5$)	patched 5 / baseline 0 / tie 0

Judge rationale (baseline). The repository shows strong alignment with SWE-bench through relevance filtering, structured edits, iterative test execution, and candidate selection, and two of four fixed probes were resolved. However, the system’s officially reported 56.6% resolve rate and the two unresolved probes indicate substantial execution and reasoning failures, so the execution dimension is the limiting factor despite relatively strong output structure and task alignment.

Diagnosed failure modes.

1. **Incomplete or incorrect issue fixes** — The system’s iterative editing and testing workflow does not reliably produce a behaviorally correct patch on harder instances, as shown by unresolved outcomes. *Evidence:* `django__django-12774`, `sympy__sympy-17630`.
2. **Insufficient test-feedback coverage** — Although the system explicitly generates and executes tests, the available evidence shows that this feedback loop can still fail to expose or resolve the cases required by the benchmark. *Evidence:* `django__django-12774`, `sympy__sympy-17630`.
3. **Candidate convergence or selection failure** — Repeated sampling and test/model-based selection improve coverage but do not guarantee that a correct candidate is generated or selected, producing unresolved final outcomes. *Evidence:* `django__django-12774`, `sympy__sympy-17630`.

Top three why-records (verbatim).

1. **prompt-06** (Minimal-patch differential verification; 0.8150, $\Delta +0.0900$, proposed_not_applied): Baseline-versus-candidate differential verification gives comparable evidence for every viable patch and selects the smallest behaviorally effective change, directly addressing candidate selection failure. It is the strongest option overall, though minimality can under-correct issues whose requirements are broader than the focused reproducer.
2. **prompt-12** (Adversarial candidate discrimination; 0.8025, $\Delta +0.0775$, applied): Active counterexample search across boundaries, invalid inputs, repeated calls, and ordering is a strong discriminator against superficially passing patches and improves convergence. The strict no-approval rule can overtest or block on ambiguous evidence, keeping execution below the strongest compatibility-focused variant.
3. **struct-18** (Evidence-first candidate handoff protocol; 0.7975, $\Delta +0.0725$, proposed_not_applied): The ordered evidence-based handoff gives selection a concrete basis for comparing behavior coverage, reproduction, test results, regressions, and scope. It directly addresses convergence failure and improves integrity and execution, though it depends on upstream candidates actually producing the evidence.

Table 38: All 60 scored versions of SWE-Rizzo.

id	phase	technique	score	Δ	decision
struct-01	struct	Add an explicit semantic validation checklist to the edit-format ...	0.7625	+0.0375	applied
struct-02	struct	Increase the reasoning budget for relevance and edit planning	0.7475	+0.0225	applied
struct-03	struct	Allow longer-running test-feedback iterations	0.7450	+0.0200	applied
struct-04	struct	Expose a larger coherent repository context to relevance analysis	0.7000	-0.0250	applied
struct-05	struct	Reduce relevance-stage concurrency to improve deterministic feedb...	0.7075	-0.0175	applied
struct-06	struct	Require self-critique before accepting a sampled candidate	0.7800	+0.0550	applied
struct-07	struct	Add requirement-to-patch traceability before editing	0.7550	+0.0300	applied
struct-08	struct	Require adversarial and negative regression tests	0.7525	+0.0275	applied
struct-09	struct	Make final candidate output include an explicit hidden-test risk ...	0.7525	+0.0275	applied
struct-10	struct	Add a structured candidate handoff rubric for downstream selection	0.7450	+0.0200	applied
struct-11	struct	Normalize relevance-parser input before classification	0.7125	-0.0125	applied
struct-12	struct	Strengthen the final prompt with a pre-submission behavioral diff	0.7725	+0.0475	applied
struct-13	struct	Requirement-to-implementation traceability gate	0.7725	+0.0475	applied
struct-14	struct	Mandatory negative and regression test design	0.7775	+0.0525	proposed
struct-15	struct	Resource-stable relevance fanout	0.7475	+0.0225	applied
struct-16	struct	Increase context-decision reasoning budget	0.7175	-0.0075	proposed
struct-17	struct	Preserve complete model output before relevance parsing	0.7650	+0.0400	applied
struct-18	struct	Evidence-first candidate handoff protocol	0.7975	+0.0725	proposed
struct-19	struct	Add an explicit acceptance-contract gate to the editing prompt	0.7725	+0.0475	applied
struct-20	struct	Require adversarial and negative-case validation before approval	0.7725	+0.0475	applied
struct-21	struct	Use requirement-weighted candidate comparison rather than aggrega...	0.7725	+0.0475	applied
struct-22	struct	Add a caller-and-invariant review step to final patch approval	0.7900	+0.0650	applied
struct-23	struct	Increase relevance-stage timeout to prevent silent loss of feedba...	0.7350	+0.0100	applied
struct-24	struct	Reduce concurrent relevance pressure to improve completion reliab...	0.7225	-0.0025	applied
struct-25	struct	Add an explicit requirement-to-test coverage matrix to the agent'...	0.7525	+0.0275	applied
struct-26	struct	Replace aggregate test-count voting with requirement-by-requireme...	0.7500	+0.0250	applied
struct-27	struct	Add a mandatory caller/consumer impact audit with an explicit sto...	0.7750	+0.0500	applied
struct-28	struct	Use a staged relevance budget that limits concurrent context-sele...	0.7425	+0.0175	applied
struct-29	struct	Require test-failure triage before editing: classify failures as ...	0.7725	+0.0475	applied
struct-30	struct	Make final approval a two-pass process that separately verifies b...	0.7725	+0.0475	proposed
prompt-01	prompt	Requirement-to-test coverage matrix with executable acceptance ch...	0.7650	+0.0400	applied
prompt-02	prompt	Repository-wide contract and caller audit	0.7700	+0.0450	applied
prompt-03	prompt	Failure-to-requirement traceability after every test run	0.7950	+0.0700	applied
prompt-04	prompt	Adversarial edge-case and invariant review	0.7575	+0.0325	applied
prompt-05	prompt	Evidence-weighted candidate approval gate	0.7975	+0.0725	applied
prompt-06	prompt	Minimal-patch differential verification	0.8150	+0.0900	proposed

id	phase	technique	score	Δ	decision
prompt-07	prompt	Requirement decomposition with implementation tracing	0.7775	+0.0525	applied
prompt-08	prompt	Consumer and compatibility audit	0.7975	+0.0725	applied
prompt-09	prompt	Failure-directed test iteration	0.7925	+0.0675	applied
prompt-10	prompt	Coverage matrix with ambiguity blocking	0.7975	+0.0725	applied
prompt-11	prompt	Evidence-weighted candidate scoring	0.7225	-0.0025	applied
prompt-12	prompt	Adversarial candidate discrimination	0.8025	+0.0775	applied
prompt-13	prompt	Requirement-to-code traceability checklist	0.7725	+0.0475	applied
prompt-14	prompt	Caller and compatibility impact audit	0.7850	+0.0600	applied
prompt-15	prompt	Adversarial test matrix	0.7675	+0.0425	applied
prompt-16	prompt	Failure-driven iteration protocol	0.7900	+0.0650	applied
prompt-17	prompt	Pre-submission candidate self-review	0.7825	+0.0575	proposed
prompt-18	prompt	Evidence-weighted candidate selection rubric	0.7125	-0.0125	proposed
prompt-19	prompt	Requirement-to-contract matrix with explicit invariants	0.7650	+0.0400	applied
prompt-20	prompt	Adversarial verification planning	0.7550	+0.0300	applied
prompt-21	prompt	Call-graph and consumer-impact tracing	0.7675	+0.0425	applied
prompt-22	prompt	Failure taxonomy and feedback-directed iteration	0.7825	+0.0575	applied
prompt-23	prompt	Candidate self-adjudication against explicit acceptance gates	0.7950	+0.0700	applied
prompt-24	prompt	Patch-diff audit with independent alternative hypothesis	0.7750	+0.0500	proposed
prompt-25	prompt	Contract-first decomposition with regression obligations	0.7725	+0.0475	applied
prompt-26	prompt	Executable acceptance-gate design	0.7650	+0.0400	applied
prompt-27	prompt	Failure-to-requirement triage loop	0.7875	+0.0625	applied
prompt-28	prompt	Evidence-backed requirement matrix	0.7800	+0.0550	applied
prompt-29	prompt	Adversarial overfitting check	0.7575	+0.0325	applied
prompt-30	prompt	Deterministic candidate scorecard and rejection rule	0.7825	+0.0575	proposed

D.14 Agentless-1.5 + Claude-3.5 Sonnet (20241022)

Table 39: Record for Agentless-1.5 (rank 14).

Submission	20241202_agentless-1.5_claude-3.5-sonnet-20241022 (verified)
Official resolve rate	50.80% (254 instances)
Repository	OpenAutoCoder/Agentless @ 5ce588b9f
Recorded trajectories readable	0 (archive assets are credentialed)
Prompt-surface files found	4
Runnability	docker required
Baseline dims (TA/RS/EA/OI)	0.62 / 0.55 / 0.51 / 0.68
Trajectory	0.5900 $\rightarrow$ 0.6650 $\rightarrow$ 0.6650 (all PROXY)
Gate verdict	accept: mean not below reference; 4W/0L paired (sign p=0.125)
MDE ₈₀ at $n=4$	0.297; shipped $\Delta = +0.0750$ (below MDE: yes)
Paraphrase noise floor	+0.0000 (exceeds_noise_floor)
Blinded A/B ($k=5$)	patched 5 / baseline 0 / tie 0

Judge rationale (baseline). The two resolved probes show that the default Agentless-1.5 system can align with an issue and execute a valid repair, but the two unresolved probes demonstrate substantial failures in localization, reasoning, or patch validation. The dimensions are therefore kept near—but not above—the system’s reported 50.8% verified resolve rate, with execution accuracy weakest and output integrity somewhat higher because the pipeline generally produces structured patch artifacts even when they are incorrect.

Diagnosed failure modes.

1. **Localization and repair failure on repository-specific behavior** — The system can complete some fixes, but its unresolved outcomes on `pytest` and `seaborn` indicate recurring difficulty mapping issue descriptions to the correct implementation and producing a behaviorally correct patch. *Evidence:* `pytest-dev__pytest-6197`, `mwaskom__seaborn-3069`.
2. **Insufficient regression-test or reproduction-test guidance** — The explicit test-selection and reproduction-generation stages do not reliably prevent incorrect or incomplete fixes, particularly for the unresolved `pytest` and `seaborn` issues. *Evidence:* `pytest-dev__pytest-6197`, `mwaskom__seaborn-3069`.
3. **Brittle end-to-end patch validation across projects** — The pipeline’s validation can succeed on selected instances but does not generalize reliably across the benchmark’s varied repositories, as shown by two unresolved probes despite the successful `astropy` and `django` cases. *Evidence:* `pytest-dev__pytest-6197`, `mwaskom__seaborn-3069`, `astropy__astropy-13453`, `django__django-15104`.

Top three why-records (verbatim).

1. **struct-05** (Constrain generated reproductions to repository APIs and test convent...; 0.6650, $\Delta +0.0750$, applied): Requiring repository-present APIs, fixtures, and conventions most directly addresses repository-specific localization and repair failures, especially on the weak probes; the strict constraint can still block legitimate internal or less obvious entry points.
2. **prompt-16** (Executable behavior-contract reproduction; 0.6625, $\Delta +0.0725$, applied): The executable behavior contract forces explicit trigger conditions, pre-fix behavior, intended behavior, and boundaries before implementation. This is the strongest candidate for reducing reproduction ambiguity and improving project-specific reasoning, though it adds process overhead and still does not require independent end-to-end validation.
3. **prompt-18** (Validation gate with independent evidence; 0.6550, $\Delta +0.0650$, proposed_not_applied): The two-stage unmodified-versus-patched gate and exact-command reporting materially improve validation reliability and output integrity, directly addressing brittle end-to-end checks. It does less to

improve initial localization or repair reasoning, and requiring clean baseline execution can be impractical in repositories with unrelated failures.

Table 40: All 36 scored versions of Agentless-1.5.

id	phase	technique	score	Δ	decision
struct-01	struct	Correct multi-message token accounting and reserve prompt headroom	0.6025	+0.0125	applied
struct-02	struct	Turn regression-test selection into risk-aware, repository-specif...	0.6200	+0.0300	applied
struct-03	struct	Normalize and validate selected-test output before it reaches exe...	0.6325	+0.0425	applied
struct-04	struct	Require reproduction cases to assert the reported contract instea...	0.6500	+0.0600	applied
struct-05	struct	Constrain generated reproductions to repository APIs and test con...	0.6650	+0.0750	applied
struct-06	struct	Make the reproduction protocol machine-verifiable with stable pro...	0.6300	+0.0400	proposed
struct-07	struct		0.5900	+0.0000	proposed
struct-08	struct		0.5875	-0.0025	proposed
struct-09	struct		0.5850	-0.0050	proposed
struct-10	struct		0.5875	-0.0025	proposed
struct-11	struct		0.5850	-0.0050	proposed
struct-12	struct		0.5750	-0.0150	proposed
struct-13	struct		0.5900	+0.0000	proposed
struct-14	struct		0.5875	-0.0025	proposed
struct-15	struct		0.5875	-0.0025	proposed
struct-16	struct		0.5850	-0.0050	proposed
struct-17	struct		0.5850	-0.0050	proposed
struct-18	struct		0.5700	-0.0200	proposed
prompt-01	prompt	Behavior-contract decomposition	0.6275	+0.0375	applied
prompt-02	prompt	Repository-evidence gate	0.6475	+0.0575	applied
prompt-03	prompt	Affected-test classification	0.6325	+0.0425	applied
prompt-04	prompt	Call-path and boundary coverage	0.6425	+0.0525	proposed
prompt-05	prompt	Minimal failing test with control case	0.6450	+0.0550	proposed
prompt-06	prompt	Assertion robustness and implementation independence	0.6425	+0.0525	proposed
prompt-07	prompt	Evidence-traceable regression-test selection	0.6200	+0.0300	applied
prompt-08	prompt	Boundary and negative-control test selection	0.5900	+0.0000	proposed
prompt-09	prompt	Executable-test feasibility gate	0.6250	+0.0350	proposed
prompt-10	prompt	Repository-evidence-backed reproduction design	0.6400	+0.0500	applied
prompt-11	prompt	Intended-oracle and anti-overfitting constraints	0.6450	+0.0550	applied
prompt-12	prompt	Preflight execution and failure-signature specification	0.6375	+0.0475	applied
prompt-13	prompt	Evidence-linked regression test selection	0.6300	+0.0400	applied
prompt-14	prompt	Fail-before-fix test prioritization	0.6275	+0.0375	applied
prompt-15	prompt	Negative-evidence filtering for test selection	0.6350	+0.0450	applied
prompt-16	prompt	Executable behavior-contract reproduction	0.6625	+0.0725	applied
prompt-17	prompt	Boundary-pair and metamorphic reproduction design	0.6375	+0.0475	proposed
prompt-18	prompt	Validation gate with independent evidence	0.6550	+0.0650	proposed

D.15 Composio SWE-Kit (2024-10-25)

Table 41: Record for Composio SWE-Kit (rank 15).

Submission	20241025_composio_swekit (verified)
Official resolve rate	48.60% (243 instances)
Repository	ComposioHQ/composio @ 13cba53b1d
Recorded trajectories readable	0 (archive assets are credentialed)
Prompt-surface files found	20
Runnability	cli entry point plausible; execution still requires SWE-bench task in...
Baseline dims (TA/RS/EA/OI)	0.66 / 0.61 / 0.56 / 0.68
Trajectory	0.6275 $\rightarrow$ 0.6650 $\rightarrow$ 0.6850 (all PROXY)
Gate verdict	accept: mean not below reference; 4W/0L paired (sign p=0.125)
MDE ₈₀ at $n=4$	0.297; shipped $\Delta = +0.0575$ (below MDE: yes)
Paraphrase noise floor	+0.0075 (exceeds_noise_floor)
Blinded A/B ($k=5$)	patched 5 / baseline 0 / tie 0

Judge rationale (baseline). The system demonstrates meaningful task understanding and can produce successful changes, as shown by two resolved probes, but the 48.6% official resolve rate and two unresolved fixed probes require materially sub-perfect scores. Execution accuracy is weakest because the failures are end-to-end benchmark failures, while the available repository surfaces show substantial tooling and test infrastructure but do not establish consistently reliable completion.

Diagnosed failure modes.

1. **Incomplete end-to-end task execution** — The system resolves only two of the four fixed probes, with both failures representing inability to deliver an accepted repository change. *Evidence:* `django__django-14351`, `django__django-14534`.
2. **Inconsistent performance across similar framework tasks** — Performance is not reliably transferable even within Django: one Django probe is resolved while two others remain unresolved. *Evidence:* `django__django-13410`, `django__django-14351`, `django__django-14534`.
3. **Limited robustness across repository and task variations** — The resolved outcomes do not generalize across the heterogeneous benchmark workload, consistent with the system’s reported 48.6% overall resolve rate. *Evidence:* `astropy__astropy-14096`, `django__django-14351`, `django__django-14534`.

Top three why-records (verbatim).

1. **prompt-01** (Explicit completion contract; 0.6850, $\Delta +0.0575$, applied): Directly addresses incomplete end-to-end execution with a concrete implementation, testing, revision, and diff-verification contract. It should improve the weak Django outcomes, though it does not add much framework-specific guidance.
2. **prompt-02** (Test-first repair loop; 0.6750, $\Delta +0.0475$, applied): A focused reproduce-edit-test loop is useful and improves validation discipline, but the requirement to reproduce behavior before editing can be brittle when tests are absent, expensive, or difficult to isolate.
3. **prompt-04** (Repository pattern triangulation; 0.6750, $\Delta +0.0475$, applied): Triangulating the issue, implementation, neighboring commands, and tests is a stronger contract-discovery method than relying on a single source. It improves consistency and interface preservation, but the fragment is somewhat cluttered and still lacks an explicit completion loop.

Table 42: All 24 scored versions of Composio SWE-Kit.

id	phase	technique	score	Δ	decision
struct-01	struct	Use a unified configuration resolver with injected-config and env...	0.6550	+0.0275	applied
struct-02	struct	Validate required values after normalization and fail before cons...	0.6650	+0.0375	applied
struct-03	struct	Make request identifiers externally controllable so retries prese...	0.6350	+0.0075	applied
struct-04	struct	Expand transient-failure classification to cover status-code and ...	0.6425	+0.0150	applied
struct-05	struct	Prefer explicit runtime model configuration over a hard-coded inj...	0.6650	+0.0375	applied
struct-06	struct	Sanitize injected runtime configuration before reading model and ...	0.6100	-0.0175	applied
struct-07	struct	Bounded transport-level retry wrapper	0.6550	+0.0275	applied
struct-08	struct	Default network timeout enforcement	0.6425	+0.0150	applied
struct-09	struct	Expanded transient-failure classifier	0.6475	+0.0200	applied
struct-10	struct	Canonicalize injected runtime configuration	0.6500	+0.0225	applied
struct-11	struct	Explicit high-capability SWE model fallback	0.6600	+0.0325	applied
struct-12	struct	Reject ambiguous blank configuration values consistently	0.6475	+0.0200	applied
prompt-01	prompt	Explicit completion contract	0.6850	+0.0575	applied
prompt-02	prompt	Test-first repair loop	0.6750	+0.0475	applied
prompt-03	prompt	Framework-specific invariant extraction	0.6625	+0.0350	proposed
prompt-04	prompt	Repository pattern triangulation	0.6750	+0.0475	applied
prompt-05	prompt	Cross-repository reconnaissance checklist	0.6125	-0.0150	proposed
prompt-06	prompt	Diff-and-cleanliness gate	0.6425	+0.0150	proposed
prompt-07	prompt	Acceptance-gated execution loop	0.6650	+0.0375	applied
prompt-08	prompt	Test-first reproduction and repair loop	0.6225	-0.0050	applied
prompt-09	prompt	Contract triangulation matrix	0.6575	+0.0300	applied
prompt-10	prompt	Repository-adaptive workflow discovery	0.6600	+0.0325	proposed
prompt-11	prompt	Patch-integrity finalization gate	0.6675	+0.0400	proposed
prompt-12	prompt	Boundary-case and variant coverage	0.6675	+0.0400	proposed

D.16 DARS Agent

Table 43: Record for DARS Agent (rank 16).

Submission	20250205_dars_agent_claude_3.5_sonnet_deepseek_r1 (lite)
Official resolve rate	47.00% (141 instances)
Repository	darsagent/DARS-Agent @ eab35168a9
Recorded trajectories readable	0 (archive assets are credentialed)
Prompt-surface files found	8
Runnability	cli entry point plausible; execution still requires SWE-bench task in...
Baseline dims (TA/RS/EA/OI)	0.65 / 0.52 / 0.47 / 0.68
Trajectory	0.5800 $\rightarrow$ 0.6275 $\rightarrow$ 0.6450 (all PROXY)
Gate verdict	accept: mean not below reference; 4W/0L paired (sign p=0.125)
MDE ₈₀ at $n=4$	0.297; shipped $\Delta = +0.0650$ (below MDE: yes)
Paraphrase noise floor	+0.0000 (exceeds_noise_floor)
Blinded A/B ($k=5$)	patched 5 / baseline 0 / tie 0

Judge rationale (baseline). The system officially resolves 47.0% of lite instances, and the fixed probes show a mixed result: both Django tasks resolved while both SymPy tasks remained unresolved. Its explicit command protocol and DARS trajectory design support moderate target alignment and output integrity, but the unresolved probes and overall failure rate require substantially lower reasoning and execution scores.

Diagnosed failure modes.

1. **Failure to complete repository changes on difficult instances** — Both SymPy probes were unresolved, indicating that the agent frequently fails to carry an issue through to a correct completed patch. *Evidence:* `sympy__sympy-19007`, `sympy__sympy-16988`.
2. **Insufficient issue-specific reasoning or generalization** — The unresolved SymPy outcomes suggest weaker analysis of domain-specific behavior and edge cases than is needed for reliable SWE-bench performance. *Evidence:* `sympy__sympy-19007`, `sympy__sympy-16988`.
3. **Inadequate validation of proposed fixes** — The two unresolved outcomes are consistent with trajectories that do not reliably validate and refine changes until the task is actually solved. *Evidence:* `sympy__sympy-19007`, `sympy__sympy-16988`.

Top three why-records (verbatim).

1. **prompt-06** (Acceptance-criteria checklist and final audit; 0.6450, $\Delta +0.0650$, applied): The final checklist and diff audit most clearly improve output integrity, omission detection, and avoiding accidental edits. It is largely a finishing control rather than a mechanism for discovering the correct difficult fix, so reasoning and hard-instance gains remain limited.
2. **prompt-04** (Semantic oracle construction; 0.6350, $\Delta +0.0550$, applied): Constructing an explicit semantic oracle is the strongest improvement for contract-sensitive and symbolic bugs, producing the best reasoning score. It is narrower than the validation and completion techniques, and can overemphasize pre-edit analysis, so execution falls below baseline.
3. **prompt-01** (Definition-of-done gate; 0.6325, $\Delta +0.0525$, applied): The definition-of-done gate directly addresses premature termination and should improve completion, regression coverage, and final repository state. Its broad acceptance language adds useful execution discipline but only modestly improves issue-specific reasoning.

Table 44: All 24 scored versions of DARS Agent.

id	phase	technique	score	Δ	decision
struct-01	struct	Introduce an explicit implementation lifecycle with planning, loc...	0.6200	+0.0400	applied
struct-02	struct	Require repository archaeology and analogous-implementation analy...	0.6100	+0.0300	proposed
struct-03	struct	Add an explicit edge-case and regression-test design step to the ...	0.5900	+0.0100	proposed
struct-04	struct	Turn test failures into a bounded diagnose-revise-rerun loop inst...	0.6225	+0.0425	proposed
struct-05	struct	Make search_repo return-oriented by requiring definitions, refere...	0.6025	+0.0225	applied
struct-06	struct	Enforce a final acceptance checklist and diff review in the outpu...	0.6275	+0.0475	applied
struct-07	struct	Add an explicit issue-to-acceptance-criteria decomposition gate b...	0.6050	+0.0250	applied
struct-08	struct	Require an implementation-and-test trace before patch generation	0.6225	+0.0425	applied
struct-09	struct	Introduce a mandatory patch-integrity checkpoint immediately afte...	0.6150	+0.0350	applied
struct-10	struct	Replace single-stage testing with a focused-to-broad validation l...	0.6050	+0.0250	applied
struct-11	struct	Add an explicit failure-driven repair loop with a completion crit...	0.6175	+0.0375	applied
struct-12	struct	Lower the default sampling temperature for more deterministic mul...	0.5700	-0.0100	applied
prompt-01	prompt	Definition-of-done gate	0.6325	+0.0525	applied
prompt-02	prompt	Incremental validation ladder	0.6275	+0.0475	applied
prompt-03	prompt	Behavior-boundary and counterexample analysis	0.6275	+0.0475	applied
prompt-04	prompt	Semantic oracle construction	0.6350	+0.0550	applied
prompt-05	prompt	Test-failure diagnosis loop	0.6250	+0.0450	applied
prompt-06	prompt	Acceptance-criteria checklist and final audit	0.6450	+0.0650	applied
prompt-07	prompt	Enforce a hypothesis-to-patch completion loop	0.6125	+0.0325	proposed
prompt-08	prompt	Require an explicit behavioral contract and edge-case matrix	0.6025	+0.0225	proposed
prompt-09	prompt	Add a mandatory minimal-reproduction checkpoint	0.6100	+0.0300	applied
prompt-10	prompt	Use differential and regression-oriented validation	0.6200	+0.0400	proposed
prompt-11	prompt	Force repository-guided call-graph tracing	0.6150	+0.0350	applied
prompt-12	prompt	Add a final definition-of-done checklist	0.6075	+0.0275	applied

D.17 CodeShellAgent + Gemini 2.0 Flash (Experimental)

Table 45: Record for CodeShellAgent (rank 17).

Submission	20250118_codeshellagent_gemini_2.0_flash_experimental (verified)
Official resolve rate	44.20% (221 instances)
Repository	WisdomShell/codeshell @ 09d1adc88c
Recorded trajectories readable	0 (archive assets are credentialed)
Prompt-surface files found	0
Runnability	gpu/local-weights required
Baseline dims (TA/RS/EA/OI)	0.56 / 0.47 / 0.44 / 0.52
Trajectory	0.4975 $\rightarrow$ 0.5425 $\rightarrow$ 0.5525 (all PROXY)
Gate verdict	accept: mean not below reference; 4W/0L paired (sign p=0.125)
MDE ₈₀ at $n=4$	0.297; shipped $\Delta = +0.0550$ (below MDE: yes)
Paraphrase noise floor	+0.0000 (exceeds _noise_floor)
Blinded A/B ($k=5$)	— (no applied patch to compare)

Judge rationale (baseline). The archived outcomes show only two of four fixed probes resolved, consistent with the system’s broader 44.2% verified-split resolve rate. This supports moderate alignment and reasoning scores, but execution accuracy must remain the weakest dimension because more than half of instances fail, while output integrity is only moderately supported because the repository provides no prompt or patch-quality evidence.

Diagnosed failure modes.

1. **Inconsistent end-to-end task resolution** — The system resolves some repository tasks but fails to produce a successful outcome on other instances. *Evidence:* `django__django-15916`, `sphinx-doc__sphinx-9673`.
2. **Weak generalization across repositories and issue types** — Success on SymPy and one Django instance does not carry over reliably to another Django task or a Sphinx task. *Evidence:* `django__django-11880`, `django__django-15916`, `sphinx-doc__sphinx-9673`.
3. **Insufficient patch verification or correction** — The unresolved outcomes indicate that execution and final validation are not consistently completed, despite demonstrated successes on two probes. *Evidence:* `django__django-15916`, `sphinx-doc__sphinx-9673`.

Top three why-records (verbatim).

1. **prompt-10** (Failure-driven validation loop; 0.5525, $\Delta +0.0550$, proposed_not_applied): The explicit failure-classification and rerun loop is the strongest direct response to insufficient verification, improving correction of initially incomplete patches. Its gains are bounded because it can misclassify environment failures or overrun the available execution budget.
2. **prompt-07** (Explicit end-to-end completion contract; 0.5450, $\Delta +0.0475$, proposed_not_applied): The explicit completion contract directly addresses premature stopping and modestly improves implementation follow-through and final-state reporting. It does not substantially improve repository-specific diagnosis, so gains on the hard probes remain limited.
3. **struct-11** (Use a structured per-instance state machine with explicit states for ...; 0.5425, $\Delta +0.0450$, proposed_not_applied): Explicit phase transitions reduce premature termination and improve end-to-end consistency, but a rigid state machine can add overhead and still lacks concrete repository-specific repair logic.

Table 46: All 24 scored versions of CodeShellAgent.

id	phase	technique	score	Δ	decision
struct-01	struct	Add a mandatory test-first execution loop: generate or identify a...	0.5025	+0.0050	proposed
struct-02	struct	Implement a bounded repair loop that parses test failures, maps e...	0.5350	+0.0375	proposed
struct-03	struct	Introduce repository-aware setup and test discovery before reason...	0.5175	+0.0200	proposed
struct-04	struct	Add an explicit issue-to-code localization phase that extracts be...	0.5175	+0.0200	proposed
struct-05	struct	Add a patch-integrity gate that checks the final diff for non-emp...	0.5225	+0.0250	proposed
struct-06	struct	Persist a structured task state across exploration, edit, test, a...	0.5225	+0.0250	proposed
struct-07	struct	Add a mandatory diagnosis-to-test planning phase that converts th...	0.5350	+0.0375	proposed
struct-08	struct	Implement repository-adaptive test discovery: inspect project met...	0.5275	+0.0300	proposed
struct-09	struct	Add an iterative patch-repair loop that captures test failures, c...	0.5400	+0.0425	proposed
struct-10	struct	Introduce a clean-worktree validation gate that reruns the focuse...	0.5350	+0.0375	proposed
struct-11	struct	Use a structured per-instance state machine with explicit states ...	0.5425	+0.0450	proposed
struct-12	struct	Add a final-diff contract checker that verifies the diff is non-e...	0.4975	+0.0000	proposed
prompt-01	prompt	Explicit end-to-end completion protocol	0.5375	+0.0400	proposed
prompt-02	prompt	Repository reconnaissance checklist	0.4975	+0.0000	proposed
prompt-03	prompt	Issue-to-acceptance-criteria translation	0.5300	+0.0325	proposed
prompt-04	prompt	Failure-driven repair loop	0.5200	+0.0225	proposed
prompt-05	prompt	Independent regression validation	0.5250	+0.0275	proposed
prompt-06	prompt	Patch-audit and final-output gate	0.5075	+0.0100	proposed
prompt-07	prompt	Explicit end-to-end completion contract	0.5450	+0.0475	proposed
prompt-08	prompt	Repository reconnaissance before editing	0.5250	+0.0275	proposed
prompt-09	prompt	Test-first hypothesis and acceptance criteria	0.5425	+0.0450	proposed
prompt-10	prompt	Failure-driven validation loop	0.5525	+0.0550	proposed
prompt-11	prompt	Issue-to-code traceability checklist	0.5400	+0.0425	proposed
prompt-12	prompt	Strict finalization and patch-integrity gate	0.5100	+0.0125	proposed

D.18 CodeFuse-CGM

Table 47: Record for CodeFuse-CGM (rank 18).

Submission	20250310_codefuse-cgm (lite)
Official resolve rate	44.00% (132 instances)
Repository	codefuse-ai/CodeFuse-CGM @ 2c12754ade
Recorded trajectories readable	0 (archive assets are credentialed)
Prompt-surface files found	6
Runnability	cli entry point plausible; execution still requires SWE-bench task in...
Baseline dims (TA/RS/EA/OI)	0.68 / 0.56 / 0.49 / 0.64
Trajectory	0.5925 $\rightarrow$ 0.6425 $\rightarrow$ 0.6425 (all PROXY)
Gate verdict	accept: mean not below reference; 4W/0L paired (sign p=0.125)
MDE ₈₀ at $n=4$	0.297; shipped $\Delta = +0.0500$ (below MDE: yes)
Paraphrase noise floor	+0.0025 (exceeds_noise_floor)
Blinded A/B ($k=5$)	patched 4 / baseline 1 / tie 0

Judge rationale (baseline). The system is clearly aligned with SWE-bench repository-level repair and resolved both recorded Django probes, but its official 44.0% resolve rate and two unresolved probes show that successful reasoning does not generalize reliably. The strongest evidence of weakness is end-to-end execution accuracy, while the structured reranker prompts support only moderate output integrity rather than near-perfect reliability.

Diagnosed failure modes.

1. **Repository localization or retrieval misses** — The graph-based retrieval and reranking pipeline does not consistently identify the files needed for the issue, as indicated by unresolved Django and SymPy probes. *Evidence:* `django__django-15738`, `sympy__sympy-16503`.
2. **Issue-specific diagnosis is insufficiently reliable** — The system’s general repository-level prompting and model reasoning do not reliably translate issue descriptions into a correct implementation diagnosis on harder instances. *Evidence:* `django__django-15738`, `sympy__sympy-16503`.
3. **End-to-end patch execution or validation failure** — Despite an architecture aimed at repository-level code changes, the default system fails to produce an accepted solution on two of the four recorded probes, reflecting weak execution reliability. *Evidence:* `django__django-15738`, `sympy__sympy-16503`.

Top three why-records (verbatim).

1. **struct-10** (Counterexample and regression-contract reasoning; 0.6425, $\Delta +0.0500$, applied): Regression-contract and counterexample reasoning improves reproduction, invariant definition, test selection, and the completeness of the final plan. It raises execution and output integrity more than localization, while still depending on the initial diagnosis.
2. **struct-09** (Forced causal diagnosis before file ranking; 0.6400, $\Delta +0.0475$, applied): The explicit symptom-to-cause chain substantially improves issue-specific diagnosis and gives patch execution a clearer basis. It slightly improves localization and reporting, but causal hypotheses can still be wrong when repository evidence is ambiguous.
3. **prompt-08** (Contradiction-aware retrieval; 0.6325, $\Delta +0.0400$, applied): Competing localization hypotheses reduce premature commitment and are the strongest retrieval improvement, especially on ambiguous or difficult repositories, though the extra ranking step adds limited execution benefit.

Table 48: All 24 scored versions of CodeFuse-CGM.

id	phase	technique	score	Δ	decision
struct-01	struct	Require evidence-grounded file localization with explicit issue-t...	0.6175	+0.0250	applied
struct-02	struct	Make reranking inspect symbols, call relationships, and test owne...	0.6300	+0.0375	applied
struct-03	struct	Require ranked file roles and distinguish implementation files fr...	0.6150	+0.0225	applied
struct-04	struct	Enforce a machine-checkable reranker response contract with exact...	0.6175	+0.0250	applied
struct-05	struct	Increase the reranker generation budget to avoid truncated analys...	0.6050	+0.0125	applied
struct-06	struct	Retry transient inference failures and empty generations before r...	0.6100	+0.0175	applied
struct-07	struct	Graph-grounded evidence scoring in reranking	0.6250	+0.0325	applied
struct-08	struct	Query decomposition with coverage-oriented file selection	0.6150	+0.0225	applied
struct-09	struct	Forced causal diagnosis before file ranking	0.6400	+0.0475	applied
struct-10	struct	Counterexample and regression-contract reasoning	0.6425	+0.0500	applied
struct-11	struct	Deterministic decoding for reproducible patch planning	0.6075	+0.0150	applied
struct-12	struct	Larger response budget for complete patch and validation plans	0.5750	-0.0175	proposed
prompt-01	prompt	Issue-to-symbol decomposition with graph-anchored retrieval	0.6175	+0.0250	applied
prompt-02	prompt	Evidence coverage matrix with negative-evidence penalties	0.6050	+0.0125	applied
prompt-03	prompt	Competing causal hypotheses with falsification	0.6250	+0.0325	applied
prompt-04	prompt	Contract-and-regression triangulation	0.6300	+0.0375	applied
prompt-05	prompt	Patch-oriented file set and validation dependency ranking	0.6150	+0.0225	applied
prompt-06	prompt	Strict machine-readable handoff with uncertainty and validation f...	0.6100	+0.0175	proposed
prompt-07	prompt	Symbol-first graph triangulation	0.6175	+0.0250	applied
prompt-08	prompt	Contradiction-aware retrieval	0.6325	+0.0400	applied
prompt-09	prompt	Mechanism-and-invariant diagnosis	0.6250	+0.0325	applied
prompt-10	prompt	Counterexample-driven test tracing	0.6275	+0.0350	applied
prompt-11	prompt	Implementation-ready file-role ranking	0.6325	+0.0400	applied
prompt-12	prompt	Evidence-gated handoff	0.5900	-0.0025	applied

D.19 Agentless Lite + O3 Mini (20250214)

Table 49: Record for Agentless Lite (rank 19).

Submission	20250214_agentless_lite_o3_mini (verified)
Official resolve rate	42.40% (212 instances)
Repository	sorendunn/Agentless-Lite @ 01900cec17
Recorded trajectories readable	0 (archive assets are credentialed)
Prompt-surface files found	1
Runnability	cli entry point plausible; execution still requires SWE-bench task in...
Baseline dims (TA/RS/EA/OI)	0.66 / 0.55 / 0.48 / 0.62
Trajectory	0.5775 $\rightarrow$ 0.6300 $\rightarrow$ 0.6375 (all PROXY)
Gate verdict	accept: mean not below reference; 4W/0L paired (sign p=0.125)
MDE ₈₀ at $n=4$	0.297; shipped $\Delta = +0.0600$ (below MDE: yes)
Paraphrase noise floor	+0.0075 (exceeds_noise_floor)
Blinded A/B ($k=5$)	patched 5 / baseline 0 / tie 0

Judge rationale (baseline). The system resolves two of the four fixed probes, but its officially resolved rate is only 42.4%, so its dimensions are kept well below near-perfect performance. The successful Django and SymPy outcomes show useful task alignment and output production, while both unresolved scikit-learn outcomes directly support lower reasoning and execution scores.

Diagnosed failure modes.

1. **repository-specific localization and diagnosis failures** — Both unresolved scikit-learn probes indicate recurring difficulty identifying the relevant code paths and accurately diagnosing the issue in unfamiliar repository contexts. *Evidence:* `scikit-learn__scikit-learn-25747`, `scikit-learn__scikit-learn-12973`.
2. **incomplete or incorrect patch execution** — The two unresolved outcomes are consistent with the RAG-only system failing to translate retrieved context into a correct, applicable repository change. *Evidence:* `scikit-learn__scikit-learn-25747`, `scikit-learn__scikit-learn-12973`.
3. **insufficient validation of proposed fixes** — The unresolved scikit-learn results suggest weak confirmation that edits satisfy the task’s behavioral requirements, especially where repository-specific edge cases matter. *Evidence:* `scikit-learn__scikit-learn-25747`, `scikit-learn__scikit-learn-12973`.

Top three why-records (verbatim).

1. **prompt-06** (Post-patch proof obligation; 0.6375, $\Delta +0.0600$, proposed_not_applied): The proof obligation links the diagnosis to tests before editing and requires post-patch control-flow coverage, neighboring-call-site checks, and regression validation. It addresses the diagnosed validation gap while also catching some incomplete patches, though it cannot fully substitute for repository-specific localization.
2. **prompt-08** (Concrete execution-path tracing; 0.6375, $\Delta +0.0600$, applied): Concrete execution-path tracing is a stronger diagnostic intervention than a static repository map because it ties the requirement to an actual branch, transformation, and divergence. This improves target alignment and reasoning most, with a small downstream execution benefit, but it still does not impose a complete patch or validation protocol.
3. **struct-12** (Behavioral acceptance-checklist gate; 0.6300, $\Delta +0.0525$, proposed_not_applied): The acceptance checklist covers the trigger, edge cases, caller regressions, and repository tests, directly addressing the diagnosed validation failure and improving confidence in the final claim. It depends on a sound issue contract and may still miss hidden repository-specific paths, so target alignment improves only slightly.

Table 50: All 24 scored versions of Agentless Lite.

id	phase	technique	score	Δ	decision
struct-01	struct	Add a mandatory repository-traversal protocol to the editing-tool...	0.6075	+0.0300	applied
struct-02	struct	Make the tool prompt require an explicit root-cause and contract ...	0.6175	+0.0400	proposed
struct-03	struct	Require the agent to complete a minimal-but-complete change across...	0.6050	+0.0275	proposed
struct-04	struct	Add a transactional editing rule that requires checking whether e...	0.5675	-0.0100	proposed
struct-05	struct	Require post-edit behavioral validation against the issue's accep...	0.6225	+0.0450	proposed
struct-06	struct	Add adversarial validation requirements covering boundary values,...	0.5925	+0.0150	proposed
struct-07	struct	Repository boundary mapping gate	0.6050	+0.0275	applied
struct-08	struct	Semantic call-chain and test-trace requirement	0.6225	+0.0450	applied
struct-09	struct	Issue-contract extraction before patching	0.6150	+0.0375	applied
struct-10	struct	Patch precondition and minimality gate	0.6050	+0.0275	applied
struct-11	struct	Post-edit diff and source reinspection	0.6200	+0.0425	applied
struct-12	struct	Behavioral acceptance-checklist gate	0.6300	+0.0525	proposed
prompt-01	prompt	Evidence-backed repository map	0.6025	+0.0250	applied
prompt-02	prompt	Competing-hypothesis diagnosis	0.6050	+0.0275	applied
prompt-03	prompt	Minimal-diff patch contract	0.6025	+0.0250	applied
prompt-04	prompt	Patch completeness checklist	0.6150	+0.0375	applied
prompt-05	prompt	Acceptance-case matrix	0.6100	+0.0325	applied
prompt-06	prompt	Post-patch proof obligation	0.6375	+0.0600	proposed
prompt-07	prompt	Evidence-first repository localization	0.6175	+0.0400	applied
prompt-08	prompt	Concrete execution-path tracing	0.6375	+0.0600	applied
prompt-09	prompt	Requirement-to-edit checklist	0.6075	+0.0300	applied
prompt-10	prompt	Patch applicability and diff discipline	0.6150	+0.0375	applied
prompt-11	prompt	Behavioral validation matrix	0.6200	+0.0425	applied
prompt-12	prompt	Post-patch invariant and regression audit	0.6275	+0.0500	applied

D.20 SWE-Exp

Table 51: Record for SWE-Exp (rank 20).

Submission	20250806_SWE-Exp_DeepSeek-V3 (verified)
Official resolve rate	42.00% (210 instances)
Repository	YerbaPage/SWE-Exp @ 6b5c92ed0a
Recorded trajectories readable	0 (archive assets are credentialed)
Prompt-surface files found	20
Runnability	cli entry point plausible; execution still requires SWE-bench task in...
Baseline dims (TA/RS/EA/OI)	0.58 / 0.50 / 0.43 / 0.56
Trajectory	0.5175 $\rightarrow$ 0.5700 $\rightarrow$ 0.5875 (all PROXY)
Gate verdict	accept: mean not below reference; 4W/0L paired (sign p=0.125)
MDE ₈₀ at $n=4$	0.297; shipped $\Delta = +0.0700$ (below MDE: yes)
Paraphrase noise floor	+0.0050 (exceeds_noise_floor)
Blinded A/B ($k=5$)	patched 5 / baseline 0 / tie 0 (post-repair)

Judge rationale (baseline). The archive reports a 42.0% verified resolve rate, so the default system warrants middling rather than near-perfect scores, with execution accuracy lowest. Two probes resolved and two remained unresolved, showing meaningful capability but substantial reliability and generalization failures.

Diagnosed failure modes.

1. **End-to-end repair failure** — The system did not produce a verified resolution for these benchmark instances, consistent with execution accuracy being the principal weakness. *Evidence:* `django__django-14122`, `django__django-14351`.
2. **Unreliable transfer across issue types and repositories** — Despite an experience-learning and retrieval design, the fixed probes show successful handling on some tasks but failure on two Django tasks, indicating inconsistent generalization. *Evidence:* `django__django-13109`, `sympy__sympy-17655`, `django__django-14122`, `django__django-14351`.
3. **Insufficiently robust completion under difficult cases** — The unresolved outcomes indicate that the default trajectory, reasoning, or validation process can terminate without a correct patch even when the system is aimed at issue resolution. *Evidence:* `django__django-14122`, `django__django-14351`.

Top three why-records (verbatim).

1. **prompt-11** (Behavioral invariant decomposition; 0.5875, $\Delta +0.0700$, proposed_not_applied): Explicit current-versus-intended behavior, distinguishing inputs, invariants, and regression risks gives the strongest improvement in target alignment and reasoning. It can improve difficult repairs, though decomposition alone cannot ensure the patch is correctly executed.
2. **prompt-07** (Test-first repair contract; 0.5725, $\Delta +0.0550$, proposed_not_applied): The test-first contract improves requirement extraction, localization, and verification, directly reducing incomplete repairs. Gains are modest because it does not itself provide stronger recovery for difficult or misleading issues.
3. **prompt-12** (Mandatory completion and verification checklist; 0.5725, $\Delta +0.0550$, proposed_not_applied): The completion checklist directly improves verification, production-path correctness, and final-diff integrity, making it valuable for end-to-end completion. It is less effective at discovering the right behavioral hypothesis, so reasoning and target gains remain limited.

Table 52: All 24 scored versions of SWE-Exp.

id	phase	technique	score	Δ	decision
struct-01	struct	Repository-root-first import resolution	0.5550	+0.0375	applied
struct-02	struct	Absolute environment-file discovery	0.5275	+0.0100	applied
struct-03	struct	Retry-floor wrapper for transient completion failures	0.5375	+0.0200	applied
struct-04	struct	Deterministic experience-selection sampling	0.5075	-0.0100	applied
struct-05	struct	Bounded parallel retrieval executor	0.5350	+0.0175	applied
struct-06	struct	Fallback JSON extraction around repaired model output	0.2375	-0.2800	applied
struct-07	struct	Use a real JSONDecoder scan for malformed LLM responses instead o...	0.5700	+0.0525	applied
struct-08	struct	Make repository-local imports take precedence over installed pack...	0.5500	+0.0325	applied
struct-09	struct	Use an environment-controlled, bounded worker count with a minimu...	0.5375	+0.0200	proposed
struct-10	struct	Ensure the dotenv file explicitly overrides stale process-level s...	0.5175	+0.0000	applied
struct-11	struct	Seed PyTorch immediately after its import, including CUDA generat...	0.5375	+0.0200	proposed
struct-12	struct	Use a conservative default of one retrieval worker while retainin...	0.5350	+0.0175	applied
prompt-01	prompt	Hypothesis-first repair planning with explicit acceptance criteria	0.5525	+0.0350	proposed
prompt-02	prompt	Repository-local evidence and regression-test prioritization	0.5525	+0.0350	proposed
prompt-03	prompt	Structured alternative-hypothesis elimination	0.5300	+0.0125	proposed
prompt-04	prompt	Mandatory layered validation and failure-driven iteration	0.5600	+0.0425	proposed
prompt-05	prompt	Repository-specific retrieval gate with negative guidance	0.5250	+0.0075	proposed
prompt-06	prompt	Completion contract with patch-state and output integrity checks	0.5500	+0.0325	proposed
prompt-07	prompt	Test-first repair contract	0.5725	+0.0550	proposed
prompt-08	prompt	Repository-convention adaptation	0.5475	+0.0300	proposed
prompt-09	prompt	Failure-recovery trajectory	0.5625	+0.0450	proposed
prompt-10	prompt	Retrieval evidence gate	0.5375	+0.0200	proposed
prompt-11	prompt	Behavioral invariant decomposition	0.5875	+0.0700	proposed
prompt-12	prompt	Mandatory completion and verification checklist	0.5725	+0.0550	proposed

D.21 SWE-RL (Llama3-SWE-RL-70B + Agentless Mini) (20250226)

Table 53: Record for SWE-RL (rank 21).

Submission	20250226_swer1_llama3_70b (verified)
Official resolve rate	41.20% (206 instances)
Repository	facebookresearch/swe-rl @ 5aa10d67f1
Recorded trajectories readable	0 (archive assets are credentialed)
Prompt-surface files found	1
Runnability	gpu/local-weights required
Baseline dims (TA/RS/EA/OI)	0.59 / 0.55 / 0.46 / 0.62
Trajectory	0.5550 $\rightarrow$ 0.6075 $\rightarrow$ 0.6075 (all PROXY)
Gate verdict	accept: mean not below reference; 4W/0L paired (sign p=0.125)
MDE ₈₀ at $n=4$	0.297; shipped $\Delta = +0.0525$ (below MDE: yes)
Paraphrase noise floor	+0.0000 (exceeds_noise_floor)
Blinded A/B ($k=5$)	patched 5 / baseline 0 / tie 0

Judge rationale (baseline). The archive shows only two of four fixed probes resolved, consistent with the system’s broader 41.2% verified resolve rate rather than near-perfect performance. The two successful probes support moderately strong capability, but the Django and Sphinx failures indicate substantial weaknesses in alignment and, most clearly, executable patch correctness.

Diagnosed failure modes.

1. **Incorrect or incomplete code repair** — The system fails to produce a patch that resolves the issue on some repositories, indicating weak end-to-end repair accuracy. *Evidence:* `django__django-17084`, `sphinx-doc__sphinx-9281`.
2. **Insufficient issue-to-code alignment** — Unresolved instances suggest that localization or interpretation of the requested behavioral change is not consistently reliable across projects. *Evidence:* `django__django-17084`, `sphinx-doc__sphinx-9281`.
3. **Fragile generated-edit pipeline** — The repair path depends on extracting and splitting structured Python edit commands, so failures in producing applicable edits can prevent otherwise plausible reasoning from becoming a valid repository patch. *Evidence:* `django__django-17084`, `sphinx-doc__sphinx-9281`.

Top three why-records (verbatim).

1. **struct-11** (Recognize common Markdown fence language labels and optional indentat...; 0.6075, $\Delta +0.0525$, proposed_not_applied): This most directly broadens fallback recognition to realistic python, python3, and py fences with indentation and optional line breaks. It improves issue-to-code alignment and recovery coverage, though the permissive pattern can still accept malformed fence layouts.
2. **prompt-12** (Patch self-review against acceptance criteria; 0.6025, $\Delta +0.0475$, proposed_not_applied): The structured final review catches unmapped requirements, caller regressions, exception mistakes, and missing imports, producing broad but moderate gains; it is weaker than an actual test loop because it relies on the agent’s own inspection.
3. **prompt-04** (Test-first validation and repair loop; 0.6000, $\Delta +0.0450$, proposed_not_applied): A reproduce-fix-targeted-test loop is the strongest direct defense against incorrect repairs and improves reasoning through observed feedback. It can be costly or brittle when tests are hard to run, so alignment and output integrity remain near baseline.

Table 54: All 24 scored versions of SWE-RL.

id	phase	technique	score	Δ	decision
struct-01	struct	Increase repair-context budget so issue state-ments, localization ...	0.5750	+0.0200	applied
struct-02	struct	Add a resilient extraction wrapper that retries parsing fenced Py...	0.5875	+0.0325	applied
struct-03	struct	Guard the multifile application loop against unknown paths and em...	0.5675	+0.0125	proposed
struct-04	struct	Normalize model-produced path separators before matching generate...	0.5775	+0.0225	applied
struct-05	struct	Initialize each edited file from its repository content so a part...	0.5975	+0.0425	applied
struct-06	struct	Discard parsed command groups that are empty before entering file...	0.5700	+0.0150	applied
struct-07	struct	Correct the escaped-whitespace fence matcher so fallback extracti...	0.5925	+0.0375	applied
struct-08	struct	Merge primary-parser results with fenced-block recovery instead o...	0.5875	+0.0325	applied
struct-09	struct	Deduplicate recovered edit commands while preserving their first-...	0.5725	+0.0175	applied
struct-10	struct	Normalize non-string model responses before extraction	0.5775	+0.0225	applied
struct-11	struct	Recognize common Markdown fence language labels and optional inde...	0.6075	+0.0525	proposed
struct-12	struct	Dedent recovered fenced blocks before sending them to the structu...	0.4925	-0.0625	applied
prompt-01	prompt	Issue-contract and acceptance-criteria prompting	0.5800	+0.0250	proposed
prompt-02	prompt	Evidence-based localization with competing hypotheses	0.5925	+0.0375	proposed
prompt-03	prompt	Minimal semantic patch with regression-focused reasoning	0.5925	+0.0375	proposed
prompt-04	prompt	Test-first validation and repair loop	0.6000	+0.0450	proposed
prompt-05	prompt	Canonical structured-edit output contract	0.5800	+0.0250	proposed
prompt-06	prompt	Self-checking edit serialization	0.5825	+0.0275	proposed
prompt-07	prompt	Evidence-first localization gate	0.5775	+0.0225	proposed
prompt-08	prompt	Behavioral contract and edge-case checklist	0.5875	+0.0325	proposed
prompt-09	prompt	Single deterministic edit protocol	0.5800	+0.0250	proposed
prompt-10	prompt	Test-driven repair loop	0.5975	+0.0425	proposed
prompt-11	prompt	Repository-convention triangulation	0.5925	+0.0375	proposed
prompt-12	prompt	Patch self-review against acceptance criteria	0.6025	+0.0475	proposed

D.22 OrcaLoca + Agentless-1.5 + Claude-3.5 Sonnet (20241022)

Table 55: Record for OrcaLoca (rank 22).

Submission	20250113_OrcaLoca (lite)
Official resolve rate	41.00% (123 instances)
Repository	fishmingyu/OrcarLLM @ 341de75336
Recorded trajectories readable	0 (archive assets are credentialed)
Prompt-surface files found	4
Runnability	docker required
Baseline dims (TA/RS/EA/OI)	0.63 / 0.56 / 0.42 / 0.84
Trajectory	0.6125 $\rightarrow$ 0.6600 $\rightarrow$ 0.6775 (all PROXY)
Gate verdict	accept: mean not below reference; 4W/0L paired (sign p=0.125)
MDE ₈₀ at $n=4$	0.297; shipped $\Delta = +0.0650$ (below MDE: yes)
Paraphrase noise floor	+0.0075 (exceeds_noise_floor)
Blinded A/B ($k=5$)	patched 5 / baseline 0 / tie 0 (post-repair)

Judge rationale (baseline). The system resolved 41.0% of the SWE-bench Lite instances, so its target alignment and reasoning are only moderate and its execution accuracy is substantially weaker, despite a relatively strong structured-output interface. The probes show two successful cases but two unresolved SymPy cases, supporting good potential on some tasks but inconsistent localization-to-resolution performance.

Diagnosed failure modes.

1. **Failure to convert localization into a correct end-to-end resolution** — The system can produce structured search and bug-location outputs, but unresolved outcomes indicate that localization or subsequent execution did not reliably yield a correct solution. *Evidence:* `sympy__sympy-18698`, `sympy__sympy-14396`.
2. **Insufficient robustness on difficult SymPy issues** — Both recorded SymPy probes were unresolved, suggesting weak generalization or reasoning robustness on technically complex repository issues. *Evidence:* `sympy__sympy-18698`, `sympy__sympy-14396`.
3. **Inconsistent issue-specific prioritization and context selection** — The framework relies on LLM-guided action decomposition, relevance scoring, and context pruning, yet the mixed probe outcomes show that these mechanisms do not consistently identify the actionable code path. *Evidence:* `sympy__sympy-18621`, `sympy__sympy-18698`, `sympy__sympy-14396`.

Top three why-records (verbatim).

1. **prompt-18** (Evidence-preserving final decision protocol; 0.6775, $\Delta +0.0650$, proposed_not_applied): The evidence table and rejection of weakly supported candidates most directly improve context selection, traceability, and final decision quality. It slightly outperforms the baseline on alignment and integrity, but its strict acceptance protocol can reject workable repairs or add overhead, limiting execution gains and leaving the deepest SymPy reasoning cases below the semantic-tracing version.
2. **prompt-04** (Executable causal-chain scoring; 0.6700, $\Delta +0.0575$, proposed_not_applied): The executable causal-chain requirement best addresses the diagnosed localization-to-resolution gap by connecting reproduction, invariant, call path, minimal patch, and validation. It has the strongest end-to-end gains, but its strict chain can reject otherwise valid fixes when repository tests or environmental observations are imperfect.
3. **prompt-23** (Post-edit behavioral verification and regression expansion; 0.6675, $\Delta +0.0550$, applied): Diff inspection plus reproducer, regression, nearby-case, and existing-test verification most directly closes the baseline resolution gap. It substantially improves execution and output integrity, while only indirectly helping difficult diagnosis; weak hypotheses can still survive the verification set.

Table 56: All 60 scored versions of OrcaLoca.

id	phase	technique	score	Δ	decision
struct-01	struct	Add an explicit issue-modeling state to every search step	0.6175	+0.0050	applied
struct-02	struct	Require evidence-backed bug locations with confidence and causal ...	0.6425	+0.0300	applied
struct-03	struct	Replace raw relevance scoring with a repairability-aware rubric	0.6475	+0.0350	applied
struct-04	struct	Make search planning explicit about novelty and coverage	0.6300	+0.0175	proposed
struct-05	struct	Add a mandatory reproduce-patch-retest gate to the agent contract	0.6575	+0.0450	applied
struct-06	struct	Use a calibrated scoring rubric with competing hypotheses and abs...	0.6325	+0.0200	applied
struct-07	struct	Add an explicit causal-chain and repair-contract requirement to c...	0.6425	+0.0300	applied
struct-08	struct	Require executable reproduction obligations and semantic invarian...	0.6500	+0.0375	applied
struct-09	struct	Turn observation feedback into evidence-delta guidance for iterat...	0.6450	+0.0325	applied
struct-10	struct	Add a mandatory patch-and-validation plan to bug-location output	0.6575	+0.0450	applied
struct-11	struct	Require explicit rejection evidence for competing locations	0.6175	+0.0050	applied
struct-12	struct	Inject symbolic-semantics and edge-case checks into code scoring	0.6375	+0.0250	proposed
struct-13	struct	Add an explicit causal-to-patch execution contract to code scoring	0.6500	+0.0375	proposed
struct-14	struct	Expand the step state into a typed failure contract and patch obl...	0.6475	+0.0350	applied
struct-15	struct	Turn observation feedback into hypothesis falsification and infor...	0.6475	+0.0350	applied
struct-16	struct	Require multi-layer and invariant-based reproduction obligations	0.6500	+0.0375	applied
struct-17	struct	Make code scores conditional on execution-path and test consisten...	0.6025	-0.0100	applied
struct-18	struct	Attach causal roles and evidence requirements to every proposed l...	0.6425	+0.0300	proposed
struct-19	struct	Causal-chain scorer with executable-oracle gating	0.6475	+0.0350	applied
struct-20	struct	Explicit patch-contract schema in the planning exemplar	0.6500	+0.0375	applied
struct-21	struct	Execution-gate schema with staged validation checkpoints	0.6525	+0.0400	applied
struct-22	struct	Evidence-backed hypothesis ledger in the issue model	0.6500	+0.0375	applied
struct-23	struct	Distance-aware context contract centered on the causal slice	0.6500	+0.0375	applied
struct-24	struct	SymPy semantic-boundary regression matrix	0.6400	+0.0275	applied
struct-25	struct	Require a recorded baseline before accepting any localization hyp...	0.6400	+0.0275	applied
struct-26	struct	Add mandatory causal-chain validation before patch application	0.6425	+0.0300	applied
struct-27	struct	Turn the execution gate into a falsification-and-recovery loop	0.6550	+0.0425	applied
struct-28	struct	Require boundary and neighboring-behavior evidence in the regress...	0.6425	+0.0300	applied
struct-29	struct	Make the patch contract reject speculative multi-site edits	0.6175	+0.0050	applied
struct-30	struct	Strengthen the resolution acceptance rule with an exact oracle an...	0.6600	+0.0475	applied
prompt-01	prompt	Evidence-locked execution gate	0.6400	+0.0275	applied
prompt-02	prompt	Competing-hypothesis recovery loop	0.6325	+0.0200	applied
prompt-03	prompt	Invariant-first patch contract	0.6550	+0.0425	applied

id	phase	technique	score	Δ	decision
prompt-04	prompt	Executable causal-chain scoring	0.6700	+0.0575	proposed
prompt-05	prompt	Layered context packet	0.6475	+0.0350	applied
prompt-06	prompt	SymPy-specific semantic adversarial checks	0.6475	+0.0350	applied
prompt-07	prompt	Post-localization execution gate	0.6375	+0.0250	proposed
prompt-08	prompt	First-project-owned-frame prioritization	0.6200	+0.0075	proposed
prompt-09	prompt	Counterfactual causal-chain scoring	0.6475	+0.0350	proposed
prompt-10	prompt	Mandatory post-patch oracle sequence	0.6375	+0.0250	applied
prompt-11	prompt	Evidence-first localization packet	0.6425	+0.0300	applied
prompt-12	prompt	Executable falsification before edit commitment	0.6575	+0.0450	applied
prompt-13	prompt	Reproducer-first execution gate	0.6425	+0.0300	applied
prompt-14	prompt	Semantic transformation and invariant tracing	0.6450	+0.0325	proposed
prompt-15	prompt	Competing-hypothesis falsification	0.6375	+0.0250	applied
prompt-16	prompt	Patch-to-oracle closure checklist	0.6575	+0.0450	applied
prompt-17	prompt	Nearby-case and boundary preservation	0.6650	+0.0525	applied
prompt-18	prompt	Evidence-preserving final decision protocol	0.6775	+0.0650	proposed
prompt-19	prompt	Reproduction-first execution gate with explicit fallback	0.6425	+0.0300	applied
prompt-20	prompt	Invariant-and-counterexample reasoning for difficult cases	0.6350	+0.0225	applied
prompt-21	prompt	Causal-path localization rather than symbol relevance ranking	0.6375	+0.0250	applied
prompt-22	prompt	Discriminating-probe requirement with explicit next action	0.6250	+0.0125	proposed
prompt-23	prompt	Post-edit behavioral verification and regression expansion	0.6675	+0.0550	applied
prompt-24	prompt	Fact-hypothesis separation with an actionable uncertainty ledger	0.6375	+0.0250	applied
prompt-25	prompt	Reproduction-first execution gate with explicit patch-validation ...	0.6600	+0.0475	applied
prompt-26	prompt	SymPy-specific semantic counterexample protocol	0.6500	+0.0375	applied
prompt-27	prompt	Evidence-backed candidate comparison table	0.6475	+0.0350	applied
prompt-28	prompt	Falsification-before-edit instruction	0.6650	+0.0525	applied
prompt-29	prompt	Assumption-and-evaluation-state checklist for SymPy edits	0.6550	+0.0425	applied
prompt-30	prompt	Causal-path and context-boundary selection	0.6225	+0.0100	applied

D.23 Patched.Codes Patchwork

Table 57: Record for Patched.Codes Patchwork (rank 23).

Submission	20250104_patched_codes_claude-3.5-sonnet-20241022 (lite)
Official resolve rate	37.00% (111 instances)
Repository	patched-codes/patchwork @ 21948cbec4
Recorded trajectories readable	0 (archive assets are credentialed)
Prompt-surface files found	20
Runnability	gpu/local-weights required
Baseline dims (TA/RS/EA/OI)	0.62 / 0.57 / 0.46 / 0.54
Trajectory	0.5475 $\rightarrow$ 0.5925 $\rightarrow$ 0.6175 (all PROXY)
Gate verdict	accept: mean not below reference; 4W/0L paired (sign p=0.125)
MDE ₈₀ at $n=4$	0.297; shipped $\Delta = +0.0700$ (below MDE: yes)
Paraphrase noise floor	+0.0075 (exceeds_noise_floor)
Blinded A/B ($k=5$)	patched 5 / baseline 0 / tie 0

Judge rationale (baseline). The official 37.0% resolve rate indicates broad unreliability, and the fixed probes show only the two SymPy tasks resolved while both Django tasks failed. The prompt, planning, and agentic-tool code provides substantial capability surface, but the observed Django failures warrant materially lower execution and reasoning scores rather than near-perfect ratings.

Diagnosed failure modes.

1. **Repository- and issue-type generalization failure** — The system resolved both sampled SymPy instances but failed both sampled Django instances, indicating weak transfer across repository conventions and issue domains. *Evidence:* `django__django-11001`, `django__django-11910`.
2. **Unreliable multi-step patch execution** — Although the repository exposes planning, agentic roles, and code-edit tools, both Django probes remained unresolved, suggesting that planning does not reliably become a correct, applied, and validated patch. *Evidence:* `django__django-11001`, `django__django-11910`.
3. **Insufficient completion or validation of final changes** — The unresolved outcomes are consistent with failures to finish the requested task or verify the resulting behavior, especially on the Django probes despite the system’s structured prompt and tool surfaces. *Evidence:* `django__django-11001`, `django__django-11910`.

Top three why-records (verbatim).

1. **prompt-23** (Evidence-based validation barrier; 0.6175, $\Delta +0.0700$, proposed_not_applied): The validation barrier most directly addresses premature completion, improving execution accuracy and output integrity through diff inspection and focused regression testing; it also helps expose wrong assumptions in low-performing Django cases.
2. **prompt-30** (Final completion gate; 0.6075, $\Delta +0.0600$, proposed_not_applied): The final completion gate checks behavior, full-diff integrity, and intended file changes, directly addressing incomplete or poorly validated submissions. It provides balanced gains across execution and output reliability, though it is less specific than the mandatory edit loop about how patches are produced.
3. **prompt-21** (Explicit inspect-edit-verify execution loop; 0.6050, $\Delta +0.0575$, proposed_not_applied): The closed inspect-edit-verify loop materially improves multi-step patch reliability and final state awareness, while producing moderate secondary gains in reasoning and target alignment.

Table 58: All 60 scored versions of Patched.Codes Patchwork.

id	phase	technique	score	Δ	decision
struct-01	struct	Inject a repository-aware execution contract into every agent con...	0.5900	+0.0425	applied
struct-02	struct	Prepend an explicit reconnaissance step to every generated plan	0.5675	+0.0200	applied
struct-03	struct	Make plan advancement idempotent at the end of the plan	0.5550	+0.0075	applied
struct-04	struct	Expose an explicit plan-completion predicate to the planning agent	0.5775	+0.0300	applied
struct-05	struct	Require completion claims to carry validation evidence	0.5650	+0.0175	applied
struct-06	struct	Reject malformed structured execution payloads at the result boun...	0.5200	-0.0275	applied
struct-07	struct	Repository-aware execution contract	0.5750	+0.0275	proposed
struct-08	struct	Issue-reproduction reconnaissance step	0.5850	+0.0375	applied
struct-09	struct	Explicit edit-test-completion protocol	0.5925	+0.0450	proposed
struct-10	struct	Reject unverified successful completions	0.5525	+0.0050	applied
struct-11	struct	Validation evidence on plan completion	0.5625	+0.0150	applied
struct-12	struct	Non-empty structured execution output	0.5525	+0.0050	applied
struct-13	struct	Repository-adaptive reconnaissance contract	0.5800	+0.0325	applied
struct-14	struct	Transactional edit-and-verify agent protocol	0.5900	+0.0425	proposed
struct-15	struct	Evidence-gated plan completion	0.5500	+0.0025	applied
struct-16	struct	Strict execution-result serialization	0.5600	+0.0125	applied
struct-17	struct	Step-level validation evidence propagation	0.5625	+0.0150	applied
struct-18	struct	Prompt-template format normalization	0.5675	+0.0200	applied
struct-19	struct	Enforce evidence at step-result construction time	0.5750	+0.0275	applied
struct-20	struct	Require comprehensive plan-level validation markers	0.5500	+0.0025	applied
struct-21	struct	Reject completed execution results without an execution message	0.5600	+0.0125	applied
struct-22	struct	Sanitize the tool registry before agent execution	0.5550	+0.0075	applied
struct-23	struct	Reinitialize the system instruction when conversation state is ma...	0.5450	-0.0025	applied
struct-24	struct	Normalize prompt identifiers before template lookup	0.5675	+0.0200	applied
struct-25	struct	Repository-aware bootstrap instruction	0.5850	+0.0375	applied
struct-26	struct	Deterministic tool-schema ordering	0.5600	+0.0125	applied
struct-27	struct	Evidence-bearing completion gate	0.5750	+0.0275	applied
struct-28	struct	Credibility threshold for planned steps	0.5325	-0.0150	applied
struct-29	struct	Portable prompt-template path resolution	0.5625	+0.0150	applied
struct-30	struct	Normalized prompt identifier handling	0.5625	+0.0150	applied
prompt-01	prompt	Repository reconnaissance checklist	0.5600	+0.0125	proposed
prompt-02	prompt	Issue-contract decomposition	0.5600	+0.0125	proposed
prompt-03	prompt	Atomic edit-and-observe loop	0.5800	+0.0325	proposed
prompt-04	prompt	Patch-state checkpointing	0.5800	+0.0325	proposed
prompt-05	prompt	Regression-test-first validation	0.5800	+0.0325	proposed
prompt-06	prompt	Evidence-gated completion	0.5925	+0.0450	proposed
prompt-07	prompt	Repository fingerprinting before diagnosis	0.5775	+0.0300	proposed
prompt-08	prompt	Issue contract and invariant extraction	0.5950	+0.0475	proposed
prompt-09	prompt	Checkpointed edit execution loop	0.5725	+0.0250	proposed
prompt-10	prompt	Minimal patch with state verification	0.5325	-0.0150	proposed
prompt-11	prompt	Layered validation and regression proof	0.5850	+0.0375	proposed
prompt-12	prompt	Final completion gate	0.5825	+0.0350	proposed
prompt-13	prompt	Repository fingerprinting and convention checklist	0.5750	+0.0275	proposed
prompt-14	prompt	Issue-archetype triangulation	0.5925	+0.0450	proposed
prompt-15	prompt	Transactional tool-execution protocol	0.5750	+0.0275	proposed
prompt-16	prompt	Incremental checkpointing and diff verification	0.5950	+0.0475	proposed
prompt-17	prompt	Focused regression-test mandate	0.5475	+0.0000	proposed
prompt-18	prompt	Completion gate with final-state evidence	0.5775	+0.0300	proposed
prompt-19	prompt	Repository-adaptive investigation gate	0.5925	+0.0450	proposed

id	phase	technique	score	Δ	decision
prompt-20	prompt	Framework-specific convention checklist	0.5750	+0.0275	proposed
prompt-21	prompt	Explicit inspect-edit-verify execution loop	0.6050	+0.0575	proposed
prompt-22	prompt	Minimal-diff and tool-state confirmation	0.5725	+0.0250	proposed
prompt-23	prompt	Evidence-based validation barrier	0.6175	+0.0700	proposed
prompt-24	prompt	Final acceptance checklist	0.5925	+0.0450	proposed
prompt-25	prompt	Framework-aware repository reconnaissance	0.5675	+0.0200	proposed
prompt-26	prompt	Acceptance-contract extraction	0.5950	+0.0475	proposed
prompt-27	prompt	Mandatory edit-loop execution	0.5950	+0.0475	proposed
prompt-28	prompt	Tool-state confirmation	0.5425	-0.0050	proposed
prompt-29	prompt	Evidence-gated focused testing	0.5700	+0.0225	proposed
prompt-30	prompt	Final completion gate	0.6075	+0.0600	proposed

D.24 SWE-Fixer (Qwen2.5-7b retriever + Qwen2.5-72b editor)

Table 59: Record for SWE-Fixer (rank 24).

Submission	20250306_SWE-Fixer_Qwen2.5-7b-retriever_Qwen2.5-72b-editor (verified)
Official resolve rate	32.80% (164 instances)
Repository	InternLM/SWE-Fixer @ 7871693672
Recorded trajectories readable	0 (archive assets are credentialed)
Prompt-surface files found	1
Runnability	cli entry point plausible; execution still requires SWE-bench task in...
Baseline dims (TA/RS/EA/OI)	0.52 / 0.46 / 0.36 / 0.58
Trajectory	0.4800 $\rightarrow$ 0.5500 $\rightarrow$ 0.5500 (all PROXY)
Gate verdict	accept: mean not below reference; 4W/0L paired (sign p=0.125)
MDE ₈₀ at $n=4$	0.297; shipped $\Delta = +0.0700$ (below MDE: yes)
Paraphrase noise floor	+0.0025 (exceeds_noise_floor)
Blinded A/B ($k=5$)	patched 5 / baseline 0 / tie 0

Judge rationale (baseline). The two resolved probes show that SWE-Fixer can successfully align with and solve some issues, but the two unresolved probes in this fixed set and its 32.8% official resolve rate require substantially lower scores, especially for execution accuracy. The exposed client surface mainly delegates generation to a selected chat model, and the available evidence does not demonstrate reliable repository-level validation or patch correctness.

Diagnosed failure modes.

1. **Incorrect or incomplete repository modification** — The system was unresolved on both probes below, indicating a recurring failure to turn its analysis into a patch that satisfies the benchmark tests. *Evidence:* `sphinx-doc__sphinx-9230`, `django__django-13279`.
2. **Insufficient issue-specific reasoning** — The unresolved outcomes on both a Sphinx issue and a Django issue suggest difficulty reliably grounding the solution in the repository-specific requirements rather than producing a generally plausible response. *Evidence:* `sphinx-doc__sphinx-9230`, `django__django-13279`.
3. **Insufficient validation before submission** — Because the system resolves only 32.8% of verified instances overall and failed these two recorded probes, it appears to lack reliable test-driven verification or correction of its generated changes. *Evidence:* `sphinx-doc__sphinx-9230`, `django__django-13279`.

Top three why-records (verbatim).

1. **struct-10** (Add a self-review gate for patch completeness before allowing output; 0.5500, $\Delta +0.0700$, proposed_not_applied): A completeness review specifically targets omitted call sites, edge cases, imports, and patches that only alter tests or documentation. This most directly improves patch integrity and execution accuracy, but the self-review is still model-generated and can cause overly conservative rejection of legitimate narrow fixes.
2. **struct-17** (Introduce an adversarial reviewer pass before emitting the JSON; 0.5400, $\Delta +0.0600$, proposed_not_applied): The adversarial pass checks several concrete failure modes, including stale paths, API assumptions, missing files, and non-matching edits. This materially improves patch completeness and output integrity, though it is still self-review and does not provide the same confidence as running tests.
3. **struct-08** (Inject an issue-to-code localization workflow into the editor request; 0.5375, $\Delta +0.0575$, applied): The localization workflow directly addresses issue-specific reasoning by requiring behavioral-contract inference, code-path tracing, and relevant-test identification. It should help most on the two weak probes, though prompting alone cannot ensure the model actually edits every required location or validates the patch.

Table 60: All 36 scored versions of SWE-Fixer.

id	phase	technique	score	Δ	decision
struct-01	struct	Select the strongest available editor model instead of relying on...	0.4950	+0.0150	applied
struct-02	struct	Add an explicit patch-output contract to the user request	0.5200	+0.0400	applied
struct-03	struct	Inject a repository-grounding checklist into the system instructi...	0.5175	+0.0375	applied
struct-04	struct	Add a second-pass self-review that returns a corrected final patch	0.5250	+0.0450	applied
struct-05	struct	Retry transient model/API failures before exposing a failed gener...	0.4850	+0.0050	applied
struct-06	struct	Force deterministic decoding for repository edits	0.5000	+0.0200	applied
struct-07	struct	Select the editor by capability instead of falling back to the fi...	0.5175	+0.0375	applied
struct-08	struct	Inject an issue-to-code localization workflow into the editor req...	0.5375	+0.0575	applied
struct-09	struct	Require test-oriented acceptance criteria and regression preserva...	0.5200	+0.0400	proposed
struct-10	struct	Add a self-review gate for patch completeness before allowing out...	0.5500	+0.0700	proposed
struct-11	struct	Make inference deterministic even when a nonzero temperature is a...	0.4825	+0.0025	applied
struct-12	struct	Use a longer request timeout so the editor can finish repository-...	0.4850	+0.0050	applied
struct-13	struct	Preserve the caller-selected decoding temperature	0.4450	-0.0350	applied
struct-14	struct	Force an explicit repository-tracing phase before patch synthesis	0.5150	+0.0350	applied
struct-15	struct	Require a test-derived acceptance matrix	0.5225	+0.0425	proposed
struct-16	struct	Add a mandatory implementation-and-test execution loop	0.5375	+0.0575	proposed
struct-17	struct	Introduce an adversarial reviewer pass before emitting the JSON	0.5400	+0.0600	proposed
struct-18	struct	Make model selection capability-aware instead of substring-based	0.5225	+0.0425	applied
prompt-01	prompt	Issue-to-requirement decomposition	0.5175	+0.0375	proposed
prompt-02	prompt	Repository-grounded change planning	0.5275	+0.0475	proposed
prompt-03	prompt	Test-first patch loop	0.5325	+0.0525	proposed
prompt-04	prompt	Edge-case and regression contract	0.5000	+0.0200	proposed
prompt-05	prompt	Competing-hypotheses diagnosis	0.4975	+0.0175	proposed
prompt-06	prompt	Submission completeness and diff audit	0.5225	+0.0425	proposed
prompt-07	prompt	Issue-to-code contract before editing	0.5225	+0.0425	proposed
prompt-08	prompt	Evidence-first repository investigation	0.5325	+0.0525	proposed
prompt-09	prompt	Test-driven patch loop	0.5200	+0.0400	proposed
prompt-10	prompt	Hidden-case and compatibility checklist	0.4825	+0.0025	proposed
prompt-11	prompt	Post-edit diff and scope audit	0.5150	+0.0350	proposed
prompt-12	prompt	Patch-only completion contract	0.5300	+0.0500	proposed
prompt-13	prompt	Issue-to-acceptance-criteria decomposition	0.5150	+0.0350	proposed
prompt-14	prompt	Evidence-gated diagnosis	0.5025	+0.0225	proposed
prompt-15	prompt	Requirement-to-diff mapping	0.5250	+0.0450	proposed
prompt-16	prompt	Focused test-driven repair loop	0.5225	+0.0425	proposed
prompt-17	prompt	Final diff and repository-state audit	0.5125	+0.0325	proposed
prompt-18	prompt	Adversarial hidden-test simulation	0.5025	+0.0225	proposed

D.25 CodeShellTester + GPT 4o (2024-05-13)

Table 61: Record for CodeShellTester (rank 25).

Submission	20241111_codeshelltester_gpt4o (lite)
Official resolve rate	31.33% (94 instances)
Repository	WisdomShell/codeshell @ 09d1adc88c
Recorded trajectories readable	0 (archive assets are credentialed)
Prompt-surface files found	0
Runnability	gpu/local-weights required
Baseline dims (TA/RS/EA/OI)	0.43 / 0.37 / 0.32 / 0.45
Trajectory	0.3925 → 0.4525 → 0.4525 (all PROXY)
Gate verdict	accept: mean not below reference; 4W/0L paired (sign p=0.125)
MDE ₈₀ at $n=4$	0.297; shipped $\Delta = +0.0600$ (below MDE: yes)
Paraphrase noise floor	+0.0075 (exceeds_noise_floor)
Blinded A/B ($k=5$)	— (no applied patch to compare)

Judge rationale (baseline). The system resolved both Django probes but failed the xarray and SymPy probes, while its overall official resolve rate is only 31.33%, so its capabilities should be treated as inconsistent rather than near-perfect. The strongest evidence is for a substantial execution-accuracy gap, with corresponding weaknesses in reasoning robustness and reliable final output despite some successful completions.

Diagnosed failure modes.

1. **Failure to produce a correct repository patch** — The system did not resolve the xarray and SymPy instances, indicating repeated inability to translate its analysis into a patch that satisfies the repository tests. *Evidence:* `pydata__xarray-4094`, `sympy__sympy-17139`.
2. **Weak generalization across project ecosystems** — Performance was successful on the two Django probes but failed on both non-Django probes, suggesting limited robustness when adapting to different codebases and conventions. *Evidence:* `django__django-11620`, `django__django-11583`, `pydata__xarray-4094`, `sympy__sympy-17139`.
3. **Insufficient debugging and validation** — The unresolved outcomes on two probes provide evidence that diagnosis, test-driven iteration, or final validation was not reliable enough for consistent SWE-bench completion. *Evidence:* `pydata__xarray-4094`, `sympy__sympy-17139`.

Top three why-records (verbatim).

1. **struct-05** (Require semantic diff validation before submission, including checkin...; 0.4525, $\Delta +0.0600$, proposed_not_applied): Semantic diff checks most directly improve output integrity by detecting empty, unrelated, or API-breaking patches. They are post-edit safeguards rather than a solution to diagnosis or ecosystem differences, and the repository exposes no submission validator to execute them.
2. **prompt-10** (Test-first failure reproduction loop; 0.4525, $\Delta +0.0600$, proposed_not_applied): The mandatory reproduce-edit-rerun loop directly addresses the diagnosed debugging and validation failure, producing the strongest execution gains, especially on the weak xarray and SymPy probes. It may add friction when tests are difficult to isolate, and it does not substantially improve repository interpretation by itself.
3. **prompt-12** (Final patch audit with executable acceptance criteria; 0.4525, $\Delta +0.0600$, proposed_not_applied): The final diff and acceptance-criteria gate most directly improves patch translation and reporting integrity, reducing cases where analysis is not reflected in the submitted change. It provides moderate execution gains through required tests, but is weaker than test-first reproduction for discovering the actual defect.

Table 62: All 24 scored versions of CodeShellTester.

id	phase	technique	score	Δ	decision
struct-01	struct	Enforce a structured patch-planning phase that requires the agent...	0.4350	+0.0425	proposed
struct-02	struct	Add an ecosystem-aware repository reconnaissance phase that detec...	0.4375	+0.0450	proposed
struct-03	struct	Implement an automatic failing-test localization loop that runs t...	0.4425	+0.0500	proposed
struct-04	struct	Add bounded iterative repair with separate diagnosis, edit, test,...	0.4500	+0.0575	proposed
struct-05	struct	Require semantic diff validation before submission, including che...	0.4525	+0.0600	proposed
struct-06	struct	Create a test-selection fallback hierarchy using issue keywords, ...	0.4300	+0.0375	proposed
struct-07	struct	Add a repository reconnaissance phase that identifies the project...	0.4150	+0.0225	proposed
struct-08	struct	Introduce an explicit failing-test localization stage: run the sm...	0.4200	+0.0275	proposed
struct-09	struct	Add an iterative repair loop with bounded retries. After every ed...	0.4250	+0.0325	proposed
struct-10	struct	Implement ecosystem-specific strategy selection using repository ...	0.4225	+0.0300	proposed
struct-11	struct	Add a final validation gate that runs the directly implicated tes...	0.4300	+0.0375	proposed
struct-12	struct	Add patch-integrity and recovery handling: capture the pre-edit s...	0.3900	-0.0025	proposed
prompt-01	prompt	Repository-first reconnaissance and convention mapping	0.4225	+0.0300	proposed
prompt-02	prompt	Issue-to-code contract extraction	0.4325	+0.0400	proposed
prompt-03	prompt	Minimal-diff implementation policy	0.4075	+0.0150	proposed
prompt-04	prompt	Test-guided hypothesis and iteration loop	0.4375	+0.0450	proposed
prompt-05	prompt	Cross-ecosystem semantic tracing	0.4350	+0.0425	proposed
prompt-06	prompt	Pre-submission verification and patch-integrity checklist	0.4350	+0.0425	proposed
prompt-07	prompt	Repository-convention reconnaissance before editing	0.4275	+0.0350	proposed
prompt-08	prompt	Issue-to-code contract extraction	0.4450	+0.0525	proposed
prompt-09	prompt	Minimal patch and compatibility guardrails	0.4050	+0.0125	proposed
prompt-10	prompt	Test-first failure reproduction loop	0.4525	+0.0600	proposed
prompt-11	prompt	Differential and edge-case validation checklist	0.4450	+0.0525	proposed
prompt-12	prompt	Final patch audit with executable acceptance criteria	0.4525	+0.0600	proposed

D.26 Aegis - o3-mini_1.0

Table 63: Record for Aegis - o3-mini_1.0 (rank 26).

Submission	20250207_aegis_o3mini (lite)
Official resolve rate	30.33% (91 instances)
Repository	evandiewald/aegis @ cd81da38f4
Recorded trajectories readable	0 (archive assets are credentialed)
Prompt-surface files found	2
Runnability	cli entry point plausible; execution still requires SWE-bench task in...
Baseline dims (TA/RS/EA/OI)	0.48 / 0.43 / 0.37 / 0.46
Trajectory	0.4350 $\rightarrow$ 0.4825 $\rightarrow$ 0.4975 (all PROXY)
Gate verdict	accept: mean not below reference; 4W/0L paired (sign p=0.125)
MDE ₈₀ at $n=4$	0.297; shipped $\Delta = +0.0625$ (below MDE: yes)
Paraphrase noise floor	+0.0050 (exceeds_noise_floor)
Blinded A/B ($k=5$)	patched 5 / baseline 0 / tie 0

Judge rationale (baseline). The system resolved both Django 11848 and Django 11583 but failed on Django 12286 and SymPy 12236 in the fixed probes, consistent with a capable but unreliable agent rather than near-uniform success. Its LangGraph, editor, search, and repository-environment surfaces support reasonable task alignment, but the recorded 30.33% overall resolve rate warrants low-to-moderate scores, especially for execution accuracy.

Diagnosed failure modes.

1. **Failure to produce a correct end-to-end patch** — The system did not resolve the Django 12286 or SymPy 12236 probes, indicating that issue understanding did not reliably translate into executable, test-passing changes. *Evidence:* `django__django-12286`, `sympy__sympy-12236`.
2. **Insufficient issue-specific reasoning on harder tasks** — The unresolved outcomes suggest weaknesses in diagnosing or planning solutions for at least some repository-specific issues despite the available browsing and editing workflow. *Evidence:* `django__django-12286`, `sympy__sympy-12236`.
3. **Unreliable completion and validation of results** — The agent architecture exposes tools and result-saving infrastructure, but the unresolved probes show that the final submitted state was not consistently validated into a successful resolution. *Evidence:* `django__django-12286`, `sympy__sympy-12236`.

Top three why-records (verbatim).

1. **prompt-06** (Explicit completion gate; 0.4975, $\Delta +0.0625$, applied): The explicit final-state gate directly addresses premature completion by requiring reproducible commands, results, and an evidenced environment block. It improves completion and reporting more than diagnosis; strict gating can also leave some environment-blocked tasks unresolved rather than producing a potentially usable patch.
2. **prompt-12** (Completion and diff verification gate; 0.4925, $\Delta +0.0575$, applied): This most directly addresses unreliable completion: it separates product defects from harness failures, verifies the actual files and full diff, and requires focused plus regression testing. It therefore has the strongest output-integrity and execution gains, though it contributes less to issue-specific reasoning and may occasionally misclassify environment failures without a stronger diagnosis framework.
3. **prompt-01** (Root-cause-to-patch contract; 0.4875, $\Delta +0.0525$, applied): The diagnosis-to-patch contract directly constrains scope and expected behavior, improving alignment and reducing incomplete or unrelated edits. It provides less help with discovering the correct root cause, so gains on the harder probes remain modest.

Table 64: All 24 scored versions of Aegis - o3-mini_1.0.

id	phase	technique	score	Δ	decision
struct-01	struct	Add a mandatory inspect-plan-edit-test-repair loop to the agent's...	0.4825	+0.0475	applied
struct-02	struct	Require an explicit issue-to-code diagnosis before any modificati...	0.4725	+0.0375	proposed
struct-03	struct	Make repository exploration trace both definitions and callers	0.4575	+0.0225	proposed
struct-04	struct	Add targeted-test selection and failure interpretation guidance t...	0.4675	+0.0325	proposed
struct-05	struct	Require a minimal-diff and regression-test contract for every edit	0.4675	+0.0325	proposed
struct-06	struct	Add an explicit completion gate requiring evidence from the final...	0.4625	+0.0275	proposed
struct-07	struct	Require a pre-edit hypothesis and acceptance-criteria artifact	0.4625	+0.0275	applied
struct-08	struct	Make targeted test selection and execution an explicit post-edit ...	0.4625	+0.0275	applied
struct-09	struct	Add bounded failure-driven repair iterations with failure classif...	0.4700	+0.0350	applied
struct-10	struct	Require an independent diff-review pass before declaring completi...	0.4725	+0.0375	applied
struct-11	struct	Make completion contingent on reproducible validation evidence	0.4375	+0.0025	applied
struct-12	struct	Force explicit separation of repository failures from harness fai...	0.4575	+0.0225	applied
prompt-01	prompt	Root-cause-to-patch contract	0.4875	+0.0525	applied
prompt-02	prompt	Repository-pattern triangulation	0.4725	+0.0375	applied
prompt-03	prompt	Smallest-test-first repair loop	0.4675	+0.0325	applied
prompt-04	prompt	Failure triage with actionable retry	0.4775	+0.0425	applied
prompt-05	prompt	Independent diff and compatibility audit	0.4875	+0.0525	applied
prompt-06	prompt	Explicit completion gate	0.4975	+0.0625	applied
prompt-07	prompt	Test-driven execution loop	0.4725	+0.0375	applied
prompt-08	prompt	Evidence-backed repository diagnosis	0.4675	+0.0325	applied
prompt-09	prompt	Acceptance-criteria to test mapping	0.4800	+0.0450	applied
prompt-10	prompt	Explicit regression-boundary analysis	0.4750	+0.0400	applied
prompt-11	prompt	Pre-edit consistency gate	0.4750	+0.0400	applied
prompt-12	prompt	Completion and diff verification gate	0.4925	+0.0575	applied

D.27 Agentless + RepoGraph + GPT-4o

Table 65: Record for Agentless (rank 27).

Submission	20240808_RepoGraph_gpt4o (lite)
Official resolve rate	29.67% (89 instances)
Repository	ozyyshr/RepoGraph @ 6c3977d878
Recorded trajectories readable	0 (archive assets are credentialed)
Prompt-surface files found	17
Runnability	cli entry point plausible; execution still requires SWE-bench task in...
Baseline dims (TA/RS/EA/OI)	0.46 / 0.39 / 0.33 / 0.58
Trajectory	0.4400 $\rightarrow$ 0.4800 $\rightarrow$ 0.5025 (all PROXY)
Gate verdict	accept: mean not below reference; 4W/0L paired (sign p=0.125)
MDE ₈₀ at $n=4$	0.297; shipped $\Delta = +0.0625$ (below MDE: yes)
Paraphrase noise floor	+0.0025 (exceeds_noise_floor)
Blinded A/B ($k=5$)	patched 5 / baseline 0 / tie 0 (post-repair)

Judge rationale (baseline). The repository shows substantial support for context retrieval, structured editing, parsing, syntax checks, and patch postprocessing, so target alignment and output integrity are stronger than semantic execution. However, the recorded 29.67% overall resolve rate and two unresolved fixed probes require low-to-moderate scores for reasoning and especially execution accuracy; the two resolved probes indicate useful but inconsistent capability rather than dependable repair performance.

Diagnosed failure modes.

1. **Incorrect or incomplete patch implementation** — The system can generate a plausible repair but fails to produce a patch that fully satisfies the issue, as shown by unresolved outcomes. *Evidence:* `sphinx-doc__sphinx-10451`, `django__django-14016`.
2. **Insufficient repository-level localization and reasoning** — Despite RepoGraph-supported context retrieval and relevant-file prompting, the agent does not reliably identify and reason through the complete set of affected code paths. *Evidence:* `sphinx-doc__sphinx-10451`, `django__django-14016`.
3. **Weak validation against behavioral requirements** — The pipeline includes syntax, lint, and diff postprocessing, but these mechanical checks do not ensure semantic correctness, reflected in the unresolved probes. *Evidence:* `sphinx-doc__sphinx-10451`, `django__django-14016`.

Top three why-records (verbatim).

1. **prompt-11** (Behavioral oracle and scenario validation; 0.5025, $\Delta +0.0625$, proposed_not_applied): Concrete behavioral scenarios and explicit oracles most directly address weak validation, improving implementation completeness and regression resistance; gains remain moderate because validation quality depends on correctly understanding the repository and issue semantics.
2. **prompt-06** (Adversarial patch review; 0.5000, $\Delta +0.0600$, proposed_not_applied): Adversarial counterexamples and pre/post behavioral comparison directly target weak validation while also exposing incomplete implementations. It provides the strongest execution and output gains, although it can overfocus on anticipated hidden cases without fully solving repository localization.
3. **prompt-12** (Patch-level regression and invariant audit; 0.4950, $\Delta +0.0550$, proposed_not_applied): A post-edit regression and invariant audit improves patch reliability and output integrity, but it is weaker than scenario-driven validation for discovering missing behavioral cases and does not materially improve initial localization.

Table 66: All 24 scored versions of Agentless.

id	phase	technique	score	Δ	decision
struct-01	struct	Dependency-closure localization prompt	0.4675	+0.0275	applied
struct-02	struct	Requirements-and-edge-case decomposition	0.4750	+0.0350	proposed
struct-03	struct	Test-driven repair instruction	0.4700	+0.0300	proposed
struct-04	struct	Patch completeness and edit-scope guardrails	0.4725	+0.0325	proposed
struct-05	struct	Correct message-level token accounting	0.4525	+0.0125	applied
struct-06	struct	Deterministic repair-generation default	0.4550	+0.0150	proposed
struct-07	struct	Dependency-closure repair specification	0.4800	+0.0400	applied
struct-08	struct	Behavioral acceptance-contract prompting	0.4750	+0.0350	applied
struct-09	struct	Deterministic repair generation	0.4450	+0.0050	proposed
struct-10	struct	Token-budget-aware context accounting	0.4575	+0.0175	applied
struct-11	struct	Bounded API retry and timeout policy	0.4475	+0.0075	applied
struct-12	struct	Multi-candidate repair generation	0.4600	+0.0200	proposed
prompt-01	prompt	Acceptance-contract decomposition	0.4675	+0.0275	applied
prompt-02	prompt	Repository graph closure before editing	0.4900	+0.0500	proposed
prompt-03	prompt	Executable behavioral oracle design	0.4875	+0.0475	proposed
prompt-04	prompt	Edge-case and regression matrix	0.4750	+0.0350	proposed
prompt-05	prompt	Call-chain and state-flow tracing	0.4825	+0.0425	proposed
prompt-06	prompt	Adversarial patch review	0.5000	+0.0600	proposed
prompt-07	prompt	Acceptance-contract decomposition with trace-ability	0.4800	+0.0400	applied
prompt-08	prompt	Multi-file impact and compatibility guardrails	0.4850	+0.0450	proposed
prompt-09	prompt	Graph-guided bidirectional localization	0.4775	+0.0375	proposed
prompt-10	prompt	Competing localization hypotheses	0.4625	+0.0225	proposed
prompt-11	prompt	Behavioral oracle and scenario validation	0.5025	+0.0625	proposed
prompt-12	prompt	Patch-level regression and invariant audit	0.4950	+0.0550	proposed

D.28 CodeR + GPT 4 (1106)

Table 67: Record for CodeR (rank 28).

Submission	20240604_CodeR (lite)
Official resolve rate	28.33% (85 instances)
Repository	NL2Code/CodeR @ d63468344b
Recorded trajectories readable	0 (archive assets are credentialed)
Prompt-surface files found	5
Runnability	no source code in linked repo
Baseline dims (TA/RS/EA/OI)	0.48 / 0.43 / 0.38 / 0.56
Trajectory	0.4625 $\rightarrow$ 0.5025 $\rightarrow$ 0.5075 (all PROXY)
Gate verdict	accept: mean not below reference; 4W/0L paired (sign p=0.125)
MDE ₈₀ at $n=4$	0.297; shipped $\Delta = +0.0450$ (below MDE: yes)
Paraphrase noise floor	+0.0050 (exceeds_noise_floor)
Blinded A/B ($k=5$)	patched 5 / baseline 0 / tie 0

Judge rationale (baseline). CodeR has a structured multi-agent workflow with explicit localization, reproduction, and verification roles, and it demonstrably resolves the two resolved probes. However, its official 28.33% SWE-bench lite rate means failures dominate overall, so the dimensions remain moderate rather than high, with execution accuracy weakest because the unresolved probes show that the workflow often does not produce a correct tested patch.

Diagnosed failure modes.

1. **Incorrect or incomplete fault localization** — The system’s localization stage is prompt-driven and the unresolved Astropy and Django probes indicate that it often fails to identify the correct edit locations or relevant call paths. *Evidence:* `astropy__astropy-14995`, `django__django-11848`.
2. **Patch implementation and test-execution failures** — Despite dedicated reproducer and verifier roles, the low overall 28.33% resolve rate and both unresolved probes indicate frequent failures converting understanding into a correct patch that passes the relevant tests. *Evidence:* `astropy__astropy-14995`, `django__django-11848`.
3. **Insufficient validation or generalization** — The one-submission setting and unresolved probes suggest that verification does not reliably catch incomplete fixes, regressions, or behavior outside the reproduced case. *Evidence:* `astropy__astropy-14995`, `django__django-11848`.

Top three why-records (verbatim).

1. **prompt-02** (Issue-to-code trace with localization stopping criteria; 0.5075, $\Delta +0.0450$, applied): Tracing from the public entry point to the observed failure with an explicit stopping criterion is a more concrete localization improvement than candidate ranking alone. It modestly raises target alignment and reasoning, but provides little direct protection against implementation and validation failures.
2. **prompt-06** (Post-fix regression matrix with independent confirmation; 0.5075, $\Delta +0.0450$, applied): Independent confirmation across the reproducer, regression tests, and neighboring cases best addresses insufficient validation and generalization. The indentation warning also reduces edit-format failures. Gains are strongest on probes requiring regression confidence, though the matrix adds time and can expose unrelated test instability.
3. **prompt-11** (Behavioral boundary matrix; 0.5050, $\Delta +0.0425$, applied): The behavioral contract and boundary matrix improve validation, regression awareness, and output integrity, with moderate reasoning gains. The approach depends on constructing reliable reproducers and does not itself ensure accurate patch execution, so execution remains only modestly above baseline.

Table 68: All 24 scored versions of CodeR.

id	phase	technique	score	Δ	decision
struct-01	struct	Evidence-first localization pipeline	0.4600	-0.0025	proposed
struct-02	struct	Competing localization hypotheses with elimination	0.4650	+0.0025	proposed
struct-03	struct	Patch-plan contract and pre-edit invariant checks	0.4800	+0.0175	proposed
struct-04	struct	Transactional edit-and-test loop	0.4500	-0.0125	proposed
struct-05	struct	Layered validation matrix	0.4650	+0.0025	proposed
struct-06	struct	Independent semantic review and regression oracle	0.4825	+0.0200	proposed
struct-07	struct	Evidence-backed two-pass fault localization	0.4850	+0.0225	proposed
struct-08	struct	Competing localization hypotheses with falsification	0.4825	+0.0200	proposed
struct-09	struct	Plan-then-edit patch workflow with immediate reproducer feedback	0.4900	+0.0275	proposed
struct-10	struct	Tool-mediated edit and command verification loop	0.4975	+0.0350	proposed
struct-11	struct	Layered validation matrix	0.4825	+0.0200	proposed
struct-12	struct	Behavioral contract and regression-oracle review	0.5025	+0.0400	proposed
prompt-01	prompt	Evidence-backed fault localization with explicit candidate ranking	0.4900	+0.0275	applied
prompt-02	prompt	Issue-to-code trace with localization stopping criteria	0.5075	+0.0450	applied
prompt-03	prompt	Failing-test-first patch loop with minimal-diff discipline	0.4975	+0.0350	applied
prompt-04	prompt	Command-result checkpointing and recovery protocol	0.4850	+0.0225	applied
prompt-05	prompt	Behavioral contract and boundary-case reproduction	0.5000	+0.0375	applied
prompt-06	prompt	Post-fix regression matrix with independent confirmation	0.5075	+0.0450	applied
prompt-07	prompt	Evidence-ranked candidate elimination	0.4925	+0.0300	applied
prompt-08	prompt	Execution-path proof before handoff	0.5025	+0.0400	applied
prompt-09	prompt	Failing-before-patching transaction	0.4950	+0.0325	applied
prompt-10	prompt	Failure-classification repair loop	0.5025	+0.0400	applied
prompt-11	prompt	Behavioral boundary matrix	0.5050	+0.0425	applied
prompt-12	prompt	Three-gate submission checklist	0.5025	+0.0400	applied

D.29 Aider + GPT 4o & Claude 3 Opus

Table 69: Record for Aider (rank 29).

Submission	20240523_aider (lite)
Official resolve rate	26.33% (79 instances)
Repository	paul-gauthier/aider @ 5dc9490bb3
Recorded trajectories readable	0 (archive assets are credentialed)
Prompt-surface files found	13
Runnability	cli entry point plausible; execution still requires SWE-bench task in...
Baseline dims (TA/RS/EA/OI)	0.47 / 0.40 / 0.34 / 0.45
Trajectory	0.4150 $\rightarrow$ 0.4450 $\rightarrow$ 0.4650 (all PROXY)
Gate verdict	accept: mean not below reference; 4W/0L paired (sign p=0.125)
MDE ₈₀ at $n=4$	0.297; shipped $\Delta = +0.0500$ (below MDE: yes)
Paraphrase noise floor	+0.0050 (exceeds_noise_floor)
Blinded A/B ($k=5$)	patched 5 / baseline 0 / tie 0

Judge rationale (baseline). The recorded overall resolution rate is only 26.33%, and two of the four fixed probes were unresolved, so all dimensions are kept well below 1.0, with execution accuracy lowest. The two resolved probes demonstrate useful capability on some repository tasks, but the unresolved SymPy and Pytest cases indicate substantial weaknesses in deep reasoning, edge-case compliance, and test-driven correction.

Diagnosed failure modes.

1. **Complex semantic reasoning failure** — The system did not reach a resolved patch on the SymPy instance, indicating difficulty tracing and implementing nontrivial behavioral changes across a complex codebase. *Evidence: `sympy__sympy-16792`.*
2. **Failure to satisfy behavioral edge cases** — The unresolved Pytest instance suggests insufficient attention to exact expected behavior and regression-test requirements. *Evidence: `pytest-dev__pytest-5413`.*
3. **Insufficient iterative validation and recovery** — Both unresolved probes show that the system can fail to recover from an initially incomplete or incorrect edit through testing and follow-up corrections. *Evidence: `sympy__sympy-16792`, `pytest-dev__pytest-5413`.*

Top three why-records (verbatim).

1. **prompt-12** (Add a completion gate with diff, test, and scope checks; 0.4650, $\Delta +0.0500$, proposed_not_applied): The completion gate materially improves diff discipline, test coverage, focused validation, and broader-suite verification. It is among the strongest changes for execution and output integrity and helps all probes, but it occurs late and cannot fully repair flawed initial semantic reasoning.
2. **prompt-09** (Mandate a failing-test-to-fix-to-regression-test loop; 0.4550, $\Delta +0.0400$, proposed_not_applied): The mandated test-fix-regression loop directly addresses insufficient recovery and validation. It improves execution and output integrity most strongly and raises both difficult probes, though it does not substantially improve initial semantic targeting.
3. **prompt-11** (Require evidence-based navigation before selecting an edit location; 0.4525, $\Delta +0.0375$, proposed_not_applied): Evidence-based navigation is a strong improvement over repository-name guessing: it requires callers, analogues, tests, control flow, and API confirmation before editing. This most directly improves reasoning on complex cases, while validation and edge handling remain incomplete.

Table 70: All 24 scored versions of Aider.

id	phase	technique	score	Δ	decision
struct-01	struct	Make retry backoff configurable and substantially longer for diff...	0.4225	+0.0075	applied
struct-02	struct	Increase the default model request deadline while preserving oper...	0.4300	+0.0150	applied
struct-03	struct	Disable lazy context loading for the primary GPT-4o configuration	0.4400	+0.0250	applied
struct-04	struct	Use unified diffs for the primary GPT-4o edit channel	0.3900	-0.0250	proposed
struct-05	struct	Use the strong model for weak-model summarization and auxiliary d...	0.4375	+0.0225	applied
struct-06	struct	Disable lazy loading for the dated GPT-4o fallback configuration ...	0.4200	+0.0050	applied
struct-07	struct	Use unified-diff editing for more explicit, line-anchored patches	0.4275	+0.0125	applied
struct-08	struct	Disable lazy context loading for the dated GPT-4o profile	0.4375	+0.0225	proposed
struct-09	struct	Use the full model rather than the mini model for auxiliary model...	0.4250	+0.0100	proposed
struct-10	struct	Prefer Claude Opus as the default reasoning model	0.4450	+0.0300	applied
struct-11	struct	Allow longer recovery intervals for transient model and test-loop...	0.4225	+0.0075	applied
struct-12	struct	Increase the per-request timeout for long repository-analysis and...	0.4175	+0.0025	applied
prompt-01	prompt	Require an explicit semantic-tracing plan before editing	0.4375	+0.0225	proposed
prompt-02	prompt	Add a hypothesis-driven investigation loop	0.4450	+0.0300	proposed
prompt-03	prompt	Turn the issue into an explicit behavior matrix	0.4325	+0.0175	proposed
prompt-04	prompt	Require tests to encode the observable contract	0.4225	+0.0075	proposed
prompt-05	prompt	Install a test-failure recovery protocol	0.4450	+0.0300	proposed
prompt-06	prompt	Add a final patch-and-repository integrity gate	0.4375	+0.0225	proposed
prompt-07	prompt	Require an explicit repository-tracing plan before editing	0.4375	+0.0225	proposed
prompt-08	prompt	Turn the request into an explicit behavioral contract	0.4400	+0.0250	proposed
prompt-09	prompt	Mandate a failing-test-to-fix-to-regression-test loop	0.4550	+0.0400	proposed
prompt-10	prompt	Use an edge-case matrix and preserve existing semantics	0.4425	+0.0275	proposed
prompt-11	prompt	Require evidence-based navigation before selecting an edit locati...	0.4525	+0.0375	proposed
prompt-12	prompt	Add a completion gate with diff, test, and scope checks	0.4650	+0.0500	proposed

D.30 RAG + Claude 3 Opus

Table 71: Record for RAG (rank 30).

Submission	20240402_rag_claude3opus (verified)
Official resolve rate	7.00% (35 instances)
Repository	SWE-bench/SWE-bench @ cd37836ffe
Recorded trajectories readable	0 (archive assets are credentialed)
Prompt-surface files found	1
Runnability	docker required
Baseline dims (TA/RS/EA/OI)	0.24 / 0.20 / 0.08 / 0.32
Trajectory	0.2100 → 0.2600 → 0.2725 (all PROXY)
Gate verdict	accept: mean not below reference; 4W/0L paired (sign p=0.125)
MDE ₈₀ at $n=4$	0.297; shipped $\Delta = +0.0625$ (below MDE: yes)
Paraphrase noise floor	+0.0025 (exceeds_noise_floor)
Blinded A/B ($k=5$)	patched 5 / baseline 0 / tie 0

Judge rationale (baseline). The two resolved probes show that RAG plus Claude 3 Opus can occasionally produce a correct result, but the official verified resolve rate is only 35 instances (7.0%), so execution accuracy and overall TEI must remain low. The two unresolved probes, including one from Django despite another Django success, provide direct evidence of substantial instance-level and cross-issue inconsistency.

Diagnosed failure modes.

1. **Low instance-level solve reliability** — The system resolves only 35 verified instances, a 7.0% rate, indicating that successful issue understanding does not consistently become a correct repository change. *Evidence:* `django__django-16116`, `pydata__xarray-3677`.
2. **Weak cross-instance generalization** — It resolves one Django instance but fails another, showing that success on a repository or framework does not reliably transfer to a different issue in the same codebase. *Evidence:* `django__django-10914`, `django__django-16116`.
3. **Inconsistent cross-repository execution** — The system succeeds on the SymPy probe but fails on the Xarray probe, indicating unreliable adaptation to repository-specific code, tests, and conventions. *Evidence:* `sympy__sympy-21847`, `pydata__xarray-3677`.

Top three why-records (verbatim).

1. **prompt-12** (Mandatory patch verification loop; 0.2725, $\Delta +0.0625$, proposed_not_applied): The verification loop combines falsifiable diagnosis, full-diff requirement review, focused testing, and failure interpretation, giving the broadest practical improvement across reliability and generalization. Gains remain modest because it depends on the preceding analysis being accurate and may impose substantial process overhead.
2. **struct-10** (Impose an iterative validation-and-repair loop, requiring focused tes...; 0.2600, $\Delta +0.0500$, proposed_not_applied): The focused-test, failure-inspection, repair, and regression sequence most directly improves solve reliability and catches implementation mistakes before completion. It is the strongest candidate, though requiring a regression test can add friction when the repository has unusual test setup or the reported failure is difficult to reproduce.
3. **prompt-06** (Patch-integrity and regression audit; 0.2575, $\Delta +0.0475$, proposed_not_applied): A minimal-root-cause plan followed by patch and regression auditing most directly reduces accidental edits, missed regressions, and incomplete final reporting. It gives the strongest execution and output-integrity improvement, though its extra audit steps provide only modest additional target and reasoning gains.

Table 72: All 24 scored versions of RAG.

id	phase	technique	score	Δ	decision
struct-01	struct	Add explicit task-boundary markers to the Anthropic prompt protoc...	0.2150	+0.0050	applied
struct-02	struct	Use a bounded, longer retry policy for transient model/API failur...	0.2200	+0.0100	applied
struct-03	struct	Reserve context for repository evidence and generated patches	0.2300	+0.0200	applied
struct-04	struct	Validate and normalize extracted patches before writing predictio...	0.2375	+0.0275	applied
struct-05	struct	Make dataset loading tolerant of repository-specific dataset scri...	0.2175	+0.0075	applied
struct-06	struct	Increase Anthropic client-level retry coverage independently of p...	0.2175	+0.0075	proposed
struct-07	struct	Add a mandatory evidence-to-change workflow before editing: ident...	0.2375	+0.0275	applied
struct-08	struct	Require repository-local evidence and explicitly prohibit framewo...	0.2175	+0.0075	proposed
struct-09	struct	Add a repository-adaptation phase that discovers the project's te...	0.2400	+0.0300	proposed
struct-10	struct	Impose an iterative validation-and-repair loop, requiring focused...	0.2600	+0.0500	proposed
struct-11	struct	Replace the unconstrained final response delimiter with a strict ...	0.2275	+0.0175	applied
struct-12	struct	Increase transient API retry budget and cap backoff at a longer i...	0.2225	+0.0125	applied
prompt-01	prompt	Acceptance-criteria traceability matrix	0.2375	+0.0275	applied
prompt-02	prompt	Evidence-gated diagnosis and edit plan	0.2450	+0.0350	proposed
prompt-03	prompt	Repository-local convention anchoring	0.2325	+0.0225	proposed
prompt-04	prompt	Call-graph and boundary tracing	0.2300	+0.0200	proposed
prompt-05	prompt	Repository-adaptive test discovery	0.2325	+0.0225	proposed
prompt-06	prompt	Patch-integrity and regression audit	0.2575	+0.0475	proposed
prompt-07	prompt	Acceptance-contract decomposition	0.2425	+0.0325	applied
prompt-08	prompt	Repository evidence trace	0.2400	+0.0300	applied
prompt-09	prompt	Competing-hypothesis debugging	0.2275	+0.0175	applied
prompt-10	prompt	Edit-to-requirement coverage gate	0.2575	+0.0475	proposed
prompt-11	prompt	Repository-native test planning	0.2425	+0.0325	proposed
prompt-12	prompt	Mandatory patch verification loop	0.2725	+0.0625	proposed